

# Writing R Extensions

---

Version 2.15.0 (2012-03-30)

R Development Core Team

---

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions, except that this permission notice may be stated in a translation approved by the R Development Core Team.

Copyright © 1999–2012 R Development Core Team

ISBN 3-900051-11-9

# Table of Contents

|  |          |
|--|----------|
| <b>Acknowledgements</b>                        | <b>1</b> |
| <b>1 Creating R packages</b>                   | <b>2</b> |
| 1.1 Package structure                          | 3        |
| 1.1.1 The ‘DESCRIPTION’ file                   | 4        |
| 1.1.2 The ‘INDEX’ file                         | 9        |
| 1.1.3 Package subdirectories                   | 10       |
| 1.1.4 Package bundles                          | 13       |
| 1.1.5 Data in packages                         | 13       |
| 1.1.6 Non-R scripts in packages                | 14       |
| 1.2 Configure and cleanup                      | 15       |
| 1.2.1 Using ‘Makevars’                         | 18       |
| 1.2.1.1 OpenMP support                         | 21       |
| 1.2.1.2 Using pthreads                         | 22       |
| 1.2.1.3 Compiling in sub-directories           | 22       |
| 1.2.2 Configure example                        | 23       |
| 1.2.3 Using F95 code                           | 25       |
| 1.3 Checking and building packages             | 26       |
| 1.3.1 Checking packages                        | 26       |
| 1.3.2 Building package tarballs                | 29       |
| 1.3.3 Building binary packages                 | 31       |
| 1.4 Writing package vignettes                  | 32       |
| 1.4.1 Encodings and vignettes                  | 34       |
| 1.5 Submitting a package to CRAN               | 34       |
| 1.5.1 PDF size                                 | 36       |
| 1.5.2 Package timing                           | 37       |
| 1.5.3 Windows external software                | 38       |
| 1.6 Package namespaces                         | 38       |
| 1.6.1 Specifying imports and exports           | 39       |
| 1.6.2 Registering S3 methods                   | 40       |
| 1.6.3 Load hooks                               | 40       |
| 1.6.4 An example                               | 43       |
| 1.6.5 Summary – converting an existing package | 43       |
| 1.6.6 Namespaces with S4 classes and methods   | 44       |
| 1.7 Writing portable packages                  | 45       |
| 1.7.1 Encoding issues                          | 47       |
| 1.7.2 Binary distribution                      | 48       |
| 1.8 Diagnostic messages                        | 48       |
| 1.9 Internationalization                       | 50       |
| 1.9.1 C-level messages                         | 50       |
| 1.9.2 R messages                               | 51       |
| 1.9.3 Installing translations                  | 51       |
| 1.9.4 Makefile support                         | 51       |

|          |  |           |
|----------|--|-----------|
| 1.10     | CITATION files.....                        | 52        |
| 1.11     | Package types.....                         | 53        |
| 1.11.1   | Frontend.....                              | 53        |
| 1.11.2   | Translation .....                          | 53        |
| 1.12     | Services.....                              | 54        |
| <b>2</b> | <b>Writing R documentation files .....</b> | <b>55</b> |
| 2.1      | Rd format .....                            | 55        |
| 2.1.1    | Documenting functions.....                 | 56        |
| 2.1.2    | Documenting data sets.....                 | 61        |
| 2.1.3    | Documenting S4 classes and methods.....    | 62        |
| 2.1.4    | Documenting packages .....                 | 63        |
| 2.2      | Sectioning .....                           | 63        |
| 2.3      | Marking text.....                          | 63        |
| 2.4      | Lists and tables.....                      | 65        |
| 2.5      | Cross-references.....                      | 66        |
| 2.6      | Mathematics.....                           | 67        |
| 2.7      | Figures .....                              | 68        |
| 2.8      | Insertions.....                            | 68        |
| 2.9      | Indices.....                               | 69        |
| 2.10     | Platform-specific documentation .....      | 69        |
| 2.11     | Conditional text .....                     | 70        |
| 2.12     | Dynamic pages .....                        | 70        |
| 2.13     | User-defined macros.....                   | 71        |
| 2.14     | Encoding .....                             | 72        |
| 2.15     | Processing Rd format.....                  | 73        |
| 2.16     | Editing Rd files.....                      | 73        |
| <b>3</b> | <b>Tidying and profiling R code.....</b>   | <b>74</b> |
| 3.1      | Tidying R code .....                       | 74        |
| 3.2      | Profiling R code for speed.....            | 74        |
| 3.3      | Profiling R code for memory use.....       | 76        |
| 3.3.1    | Memory statistics from <b>Rprof</b> .....  | 77        |
| 3.3.2    | Tracking memory allocations .....          | 77        |
| 3.3.3    | Tracing copies of an object .....          | 77        |
| 3.4      | Profiling compiled code .....              | 78        |
| 3.4.1    | Linux.....                                 | 78        |
| 3.4.1.1  | sprof.....                                 | 78        |
| 3.4.1.2  | oprofile.....                              | 79        |
| 3.4.2    | Solaris .....                              | 80        |
| 3.4.3    | Mac OS X .....                             | 80        |

|          |   |           |
|----------|---|-----------|
| <b>4</b> | <b>Debugging</b>  | <b>81</b> |
| 4.1      | Browsing  | 81        |
| 4.2      | Debugging R code  | 82        |
| 4.3      | Using gctorture and valgrind                                      | 86        |
| 4.3.1    | Using gctorture   | 86        |
| 4.3.2    | Using valgrind  | 87        |
| 4.4      | Debugging compiled code   | 89        |
| 4.4.1    | Finding entry points in dynamically loaded code                   | 90        |
| 4.4.2    | Inspecting R objects when debugging                               | 90        |
| <b>5</b> | <b>System and foreign language interfaces</b>                     | <b>93</b> |
| 5.1      | Operating system access   | 93        |
| 5.2      | Interface functions <code>.C</code> and <code>.Fortran</code>     | 93        |
| 5.3      | <code>dyn.load</code> and <code>dyn.unload</code>                 | 95        |
| 5.4      | Registering native routines                                       | 97        |
| 5.5      | Creating shared objects   | 99        |
| 5.6      | Interfacing C++ code  | 101       |
| 5.7      | Fortran I/O   | 103       |
| 5.8      | Linking to other packages   | 103       |
| 5.8.1    | Unix-alikes   | 103       |
| 5.8.2    | Windows   | 104       |
| 5.9      | Handling R objects in C   | 105       |
| 5.9.1    | Handling the effects of garbage collection                        | 106       |
| 5.9.2    | Allocating storage  | 108       |
| 5.9.3    | Details of R types  | 108       |
| 5.9.4    | Attributes  | 109       |
| 5.9.5    | Classes   | 112       |
| 5.9.6    | Handling lists  | 112       |
| 5.9.7    | Handling character data   | 113       |
| 5.9.8    | Finding and setting variables                                     | 113       |
| 5.9.9    | Some convenience functions  | 114       |
| 5.9.9.1  | Semi-internal convenience functions                               | 114       |
| 5.9.10   | Named objects and copying   | 115       |
| 5.10     | Interface functions <code>.Call</code> and <code>.External</code> | 116       |
| 5.10.1   | Calling <code>.Call</code>  | 116       |
| 5.10.2   | Calling <code>.External</code>                                    | 117       |
| 5.10.3   | Missing and special values  | 119       |
| 5.11     | Evaluating R expressions from C                                   | 119       |
| 5.11.1   | Zero-finding  | 121       |
| 5.11.2   | Calculating numerical derivatives                                 | 122       |
| 5.12     | Parsing R code from C   | 125       |
| 5.13     | External pointers and weak references                             | 126       |
| 5.13.1   | An example  | 128       |
| 5.14     | Vector accessor functions   | 129       |
| 5.15     | Character encoding issues   | 129       |

|          |   |            |
|----------|---|------------|
| <b>6</b> | <b>The R API: entry points for C code .....</b>   | <b>130</b> |
| 6.1      | Memory allocation .....                           | 130        |
| 6.1.1    | Transient storage allocation .....                | 131        |
| 6.1.2    | User-controlled memory .....                      | 131        |
| 6.2      | Error handling .....                              | 132        |
| 6.2.1    | Error handling from FORTRAN .....                 | 132        |
| 6.3      | Random number generation .....                    | 132        |
| 6.4      | Missing and IEEE special values .....             | 133        |
| 6.5      | Printing .....                                    | 133        |
| 6.5.1    | Printing from FORTRAN .....                       | 134        |
| 6.6      | Calling C from FORTRAN and vice versa .....       | 134        |
| 6.7      | Numerical analysis subroutines .....              | 135        |
| 6.7.1    | Distribution functions .....                      | 135        |
| 6.7.2    | Mathematical functions .....                      | 136        |
| 6.7.3    | Numerical Utilities .....                         | 137        |
| 6.7.4    | Mathematical constants .....                      | 138        |
| 6.8      | Optimization .....                                | 139        |
| 6.9      | Integration .....                                 | 140        |
| 6.10     | Utility functions .....                           | 141        |
| 6.11     | Re-encoding .....                                 | 143        |
| 6.12     | Allowing interrupts .....                         | 143        |
| 6.13     | Platform and version information .....            | 144        |
| 6.14     | Inlining C functions .....                        | 144        |
| 6.15     | Controlling visibility .....                      | 145        |
| 6.16     | Using these functions in your own C code .....    | 145        |
| 6.17     | Organization of header files .....                | 146        |
| <b>7</b> | <b>Generic functions and methods .....</b>        | <b>148</b> |
| 7.1      | Adding new generics .....                         | 149        |
| <b>8</b> | <b>Linking GUIs and other front-ends to R ...</b> | <b>150</b> |
| 8.1      | Embedding R under Unix-alikes .....               | 150        |
| 8.1.1    | Compiling against the R library .....             | 152        |
| 8.1.2    | Setting R callbacks .....                         | 152        |
| 8.1.3    | Registering symbols .....                         | 155        |
| 8.1.4    | Meshing event loops .....                         | 155        |
| 8.1.5    | Threading issues .....                            | 156        |
| 8.2      | Embedding R under Windows .....                   | 157        |
| 8.2.1    | Using (D)COM .....                                | 157        |
| 8.2.2    | Calling R.dll directly .....                      | 157        |
| 8.2.3    | Finding R_HOME .....                              | 160        |
|          | <b>Function and variable index .....</b>          | <b>162</b> |
|          | <b>Concept index .....</b>                        | <b>165</b> |

## Acknowledgements

The contributions of Saikat DebRoy (who wrote the first draft of a guide to using `.Call` and `.External`) and of Adrian Trapletti (who provided information on the C++ interface) are gratefully acknowledged.

# 1 Creating R packages

Packages provide a mechanism for loading optional code, data and documentation as needed. The R distribution itself includes about 30 packages.

In the following, we assume that you know the `library()` command, including its `lib.loc` argument, and we also assume basic knowledge of the R CMD INSTALL utility. Otherwise, please look at R’s help pages on

```
?library  
?INSTALL
```

before reading on.

A computing environment including a number of tools is assumed; the “R Installation and Administration” manual describes what is needed. Under a Unix-alike most of the tools are likely to be present by default, but Microsoft Windows may require careful setup.

Once a source package is created, it must be installed by the command R CMD INSTALL. See [Section “Add-on-packages” in R Installation and Administration](#).

Other types of extensions are supported (but rare): See [Section 1.11 \[Package types\]](#), [page 53](#).

Some notes on terminology complete this introduction. These will help with the reading of this manual, and also in describing concepts accurately when asking for help.

A *package* is a directory of files which extend R, either a *source package* (the master files of a package), or a tarball containing the files of a source package, or an *installed* package, the result of running R CMD INSTALL on a source package. On some platforms there are also *binary packages*, a zip file or tarball containing the files of an installed package which can be unpacked rather than installing from sources.

A package is **not**<sup>1</sup> a *library*. The latter is used in two senses in R documentation. The first is a directory into which packages are installed, e.g. `‘/usr/lib/R/library’`: in that sense it is sometimes referred to as a *library directory* or *library tree* (since the library is a directory which contains packages as directories, which themselves contain directories). The second sense is that used by the operating system, as a shared library or static library or (especially on Windows) a DLL, where the second L stands for ‘library’. Installed packages may contain compiled code in what is known on most Unix-alikes as a *shared object* and on Windows as a DLL (and used to be called a *shared library* on some Unix-alikes). The concept of a *shared library* (*dynamic library* on Mac OS X) as a collection of compiled code to which a package might link is also used, especially for R itself on some platforms.

There are a number of well-defined operations on source packages. The most common is *installation* which takes a source package and installs it in a library using R CMD INSTALL or `install.packages`. Source packages can be *built*, a distinct concept. This involves taking a source directory and creating a tarball ready for distribution, including cleaning it up and creating PDF documentation from any *vignettes* it may contain. Source packages (and most often tarballs) can be *checked*, when a test installation is done and tested (including running its examples); also, the contents of the package are tested in various ways for consistency and portability.

---

<sup>1</sup> although this is common mis-usage. It seems to stem from S, whose analogues of R’s packages were officially known as *library sections* and later as *chapters*, but almost always referred to as *libraries*.



*Compilation* is not a correct term for a package. Installing a source package which contains C, C++ or Fortran code will involve compiling that code. As from R 2.13.0 there is also the possibility of ('byte') compiling the R code in a package (using the facilities of package **compiler**): at some future time this might be done routinely when *compiling* a package may come to mean compiling its R code.

It used to be unambiguous to talk about *loading* an installed package using `library()`, but since the advent of package name spaces this has been less clear: people now often talk about *loading* the package's namespace and then *attaching* the package so it becomes visible on the search path. Function `library` performs both steps, but a package's namespace can be loaded without the package being attached (for example by calls like `splines::ns`).

The option of *lazy loading* of code or data is mentioned at several points. This is part of the installation, always selected for R code (since R 2.14.0) but optional for data. When used the R objects of the package are created at installation time, and stored in a database in the 'R' directory of the installed package, being loaded into the session at first use. This makes the R session run faster and use less (virtual) memory. (For technical details, see [Section "Lazy loading" in R Internals.](#))

## 1.1 Package structure

The sources of a R package consists of a subdirectory containing a file 'DESCRIPTION' and the subdirectories 'R', 'data', 'demo', 'exec', 'inst', 'man', 'po', 'src', and 'tests' (some of which can be missing, but which should not be empty). The package subdirectory may also contain files 'INDEX', 'NAMESPACE', 'configure', 'cleanup', 'LICENSE', 'LICENCE' and 'NEWS'. Other files such as 'INSTALL' (for non-standard installation instructions), 'README' or 'ChangeLog' will be ignored by R, but may be useful to end users.

The 'DESCRIPTION' and 'INDEX' files are described in the subsections below. The 'NAMESPACE' file is described in the section on [Section 1.6 \[Package namespaces\], page 38](#).

The optional files 'configure' and 'cleanup' are (Bourne shell) script files which are, respectively, executed before and (provided that option '--clean' was given) after installation on Unix-alikes, see [Section 1.2 \[Configure and cleanup\], page 15](#). The analogues on Windows are 'configure.win' and 'cleanup.win'.

The optional file 'LICENSE'/'LICENCE' contains a copy of the license to the package. Whereas you should feel free to include a license file in your *source* distribution, please do not arrange to *install* yet another copy of the GNU 'COPYING' or 'COPYING.LIB' files but refer to the copies on <http://www.r-project.org/Licenses/> and included in the R distribution (in directory 'share/licenses'). Since files named 'LICENSE' or 'LICENCE' *will* be installed, do not use these names for standard licence files.

For the conventions for files 'NEWS' and 'ChangeLog' in the GNU project see <http://www.gnu.org/prep/standards/standards.html#Documentation>.

The package subdirectory should be given the same name as the package. Because some file systems (e.g., those on Windows and by default on Mac OS X) are not case-sensitive, to maintain portability it is strongly recommended that case distinctions not be used to distinguish different packages. For example, if you have a package named 'foo', do not also create a package named 'Foo'.

To ensure that file names are valid across file systems and supported operating system platforms, the ASCII control characters as well as the characters '"', '\*', ':', '/', '<', '>', '?',

‘\’, and ‘|’ are not allowed in file names. In addition, files with names ‘con’, ‘prn’, ‘aux’, ‘clock\$', ‘nul’, ‘com1’ to ‘com9’, and ‘lpt1’ to ‘lpt9’ after conversion to lower case and stripping possible “extensions” (e.g., ‘lpt5.foo.bar’), are disallowed. Also, file names in the same directory must not differ only by case (see the previous paragraph). In addition, the basenames of ‘.Rd’ files may be used in URLs and so must be ASCII and not contain %. For maximal portability filenames should only contain only ASCII characters not excluded already (that is A-Za-z0-9.\_!#\$%&+;=@~(){}'[] — we exclude space as many utilities do not accept spaces in file paths): non-English alphabetic characters cannot be guaranteed to be supported in all locales. It would be good practice to avoid the shell metacharacters (){}'[]\$.

A source package if possible should not contain binary executable files: they are not portable, and a security risk if they are of the appropriate architecture. R CMD check will warn about them<sup>2</sup> unless they are listed (one filepath per line) in a file ‘BinaryFiles’ at the top level of the package. Note that CRAN will no longer accept submissions containing binary files even if they are listed.

The R function `package.skeleton` can help to create the structure for a new package: see its help page for details.

### 1.1.1 The ‘DESCRIPTION’ file

The ‘DESCRIPTION’ file contains basic information about the package in the following format:

```
Package: pkgname
Version: 0.5-1
Date: 2004-01-01
Title: My First Collection of Functions
Authors@R: c(person("Joe", "Developer", role = c("aut", "cre"),
                  email = "Joe.Developer@some.domain.net"),
             person("Pat", "Developer", role = "aut"),
             person("A.", "User", role = "ctb",
                  email = "A.User@whereever.net"))
Author: Joe Developer and Pat Developer, with contributions from A. User
Maintainer: Joe Developer <Joe.Developer@some.domain.net>
Depends: R (>= 1.8.0), nlme
Suggests: MASS
Description: A short (one paragraph) description of what
             the package does and why it may be useful.
License: GPL (>= 2)
URL: http://www.r-project.org, http://www.another.url
BugReports: http://pkgname.bugtracker.url
```

The format is that of a ‘Debian Control File’ (see the help for ‘read.dcf’ and <http://www.debian.org/doc/debian-policy/ch-controlfields.html>: R does not require encoding in UTF-8). Continuation lines (for example, for descriptions longer than one line) start with a space or tab. The ‘Package’, ‘Version’, ‘License’, ‘Description’, ‘Title’, ‘Author’, and ‘Maintainer’ fields are mandatory, all other fields are optional. For R 2.14.0 or later, ‘Author’ and ‘Maintainer’ can be auto-generated from ‘Authors@R’, and should be omitted if the latter is provided (and the package depends on R (>= 2.14): see below for details).

<sup>2</sup> false positives are possible, but only a handful have been seen so far.

For maximal portability, the ‘DESCRIPTION’ file should be written entirely in ASCII — if this is not possible it must contain an ‘Encoding’ field (see below).

The mandatory ‘Package’ field gives the name of the package. This should contain only letters, numbers and dot, have at least two characters and start with a letter and not end in a dot. (Translation packages are allowed names of the form ‘Translation-11’.)

The mandatory ‘Version’ field gives the version of the package. This is a sequence of at least *two* (and usually three) non-negative integers separated by single ‘.’ or ‘-’ characters. The canonical form is as shown in the example, and a version such as ‘0.01’ or ‘0.01.0’ will be handled as if it were ‘0.1-0’.

The mandatory ‘License’ field should specify the license of the package in a standardized form. Alternatives are indicated *via* vertical bars. Individual specifications must be one of

- One of the “standard” short specifications

GPL-2 GPL-3 LGPL-2 LGPL-2.1 LGPL-3 AGPL-3 Artistic-1.0 Artistic-2.0

as made available *via* <http://www.r-project.org/Licenses/> and contained in sub-directory ‘share/licenses’ of the R source or home directory.

- The names of abbreviations of free or open source software (FOSS, e.g., <http://en.wikipedia.org/wiki/FOSS>) licenses as contained in the license data base in file ‘share/licenses/license.db’ in the R source or home directory, possibly (for versioned licenses) followed by a version restriction of the form ‘(op v)’ with *op* one of the comparison operators ‘<’, ‘<=’, ‘>’, ‘>=’, ‘==’, or ‘!=’ and *v* a numeric version specification (strings of non-negative integers separated by ‘.’), possibly combined *via* ‘,’ (see below for an example). For versioned licenses, one can also specify the name followed by the version, or combine an existing abbreviation and the version with a ‘-’. Further free (see <http://www.fsf.org/licenses/license-list.html>) or open software (see <http://www.opensource.org/licenses/bsd-license.php>) licenses will be added to this data base if necessary.
- One of the strings ‘file LICENSE’ or ‘file LICENCE’ referring to a file named ‘LICENSE’ or ‘LICENCE’ in the package (source and installation) top-level directory.
- The string ‘Unlimited’, meaning that there are no restrictions on distribution or use other than those imposed by relevant laws (including copyright laws).

If a package license *extends* a base FOSS license (e.g., using GPL-3 or AGPL-3 with an attribution clause), the extension should be placed in file ‘LICENSE’ (or ‘LICENCE’), and the string ‘+ file LICENSE’ (or ‘+ file LICENCE’, respectively) should be appended to the corresponding individual license specification.

Examples for standardized specifications include

```
License: GPL-2
License: GPL (>= 2) | BSD
License: LGPL (>= 2.0, < 3) | Mozilla Public License
License: GPL-2 | file LICENCE
License: Artistic-1.0 | AGPL-3 + file LICENSE
```

Please note in particular that “Public domain” is not a valid license, since it is not recognized in some jurisdictions.

It is very important that you include this license information! Otherwise, it may not even be legally correct for others to distribute copies of the package. Do not use the ‘**License**’ field for copyright information: if needed, use a ‘**Copyright**’ field.

Please ensure that the license you choose also covers any dependencies (including system dependencies) of your package: it is particularly important that any restrictions on the use of such dependencies are evident to people reading your ‘**DESCRIPTION**’ file.

The mandatory ‘**Description**’ field should give a comprehensive description of what the package does. One can use several (complete) sentences, but only one paragraph.

The mandatory ‘**Title**’ field should give a short description of the package. Some package listings may truncate the title to 65 characters. It should be capitalized, not use any markup, not have any continuation lines, and not end in a period.

The mandatory ‘**Author**’ field describes who wrote *the package*. It is a plain text field intended for human readers, but not for automatic processing (such as extracting the email addresses of all listed contributors: for that use ‘**Authors@R**’). Note that all significant contributors must be included: if you wrote an R wrapper for the work of others included in the ‘**src**’ directory, you are not the sole (and maybe not even the main) author.

The mandatory ‘**Maintainer**’ field should give a *single* name with a *valid* (RFC 2822) email address in angle brackets (for sending bug reports etc.). It should not end in a period or comma. For a public package it should be a *person*, not a mailing list and not a corporate entity: do ensure that it is valid and will remain valid for the lifetime of the package.

Both ‘**Author**’ and ‘**Maintainer**’ fields can be omitted (as from R 2.14.0) if a suitable ‘**Authors@R**’ field is given. This field can be used to provide a refined, machine-readable description of the package “authors” (in particular specifying their precise *roles*), via suitable R code. The roles can include “**aut**” (author) for full authors, “**cre**” (creator) for the package maintainer, and “**ctb**” (contributor) for other contributors, among others. See `?person` for more information. Note that no role is assumed by default. Auto-generated package citation information takes advantage of this specification; in R 2.14.0 or later, the ‘**Author**’ and ‘**Maintainer**’ fields are auto-generated from it if needed when building or installing.

Several optional fields take *logical values*: these can be specified as ‘**yes**’, ‘**true**’, ‘**no**’ or ‘**false**’: capitalized values are also accepted.

The ‘**Date**’ field gives the release date of the current version of the package. It is strongly recommended to use the yyyy-mm-dd format conforming to the ISO 8601 standard.

The ‘**Depends**’ field gives a comma-separated list of package names which this package depends on. The package name may be optionally followed by a comment in parentheses. The comment should contain a comparison operator, whitespace and a valid version number. You can also use the special package name ‘**R**’ if your package depends on a certain version of R — e.g., if the package works only with R version 2.11.0 or later, include ‘**R (>= 2.11.0)**’ in the ‘**Depends**’ field. You can also require a certain SVN revision for R-devel or R-patched, e.g. ‘**R (>= 2.14.0), R (>= r56550)**’ requires a version later than R-devel of late July 2011 (including released versions of 2.14.0). Both `library` and the R package checking facilities use this field: hence it is an error to use improper syntax or misuse the ‘**Depends**’ field for comments on other software that might be needed. Other dependencies (external to the R system) should be listed in the ‘**SystemRequirements**’ field, possibly amplified in a separate ‘**README**’ file. The R `INSTALL` facilities check if the version of R used is recent enough for the

package being installed, and the list of packages which is specified will be attached (after checking version requirements) before the current package, both when `library` is called and when preparing for lazy-loading during installation.

A package (or ‘R’) can appear more than once in the ‘`Depends`’, but only the first occurrence was used in versions of R prior to 2.7.0: these are now very unlikely to be encountered.

It makes no sense to declare a dependence on R without a version specification, nor on the package **base**: this is an R package and **base** is always available.

The ‘`Imports`’ field lists packages whose namespaces are imported from (as specified in the ‘`NAMESPACE`’ file) but which do not need to be attached. Namespaces accessed by the ‘`::`’ and ‘`:::`’ operators must be listed here, or in ‘`Suggests`’ or ‘`Enhances`’ (see below). Ideally this field will include all the standard packages that are used, and it is important to include S4-using packages (as their class definitions can change and the ‘`DESCRIPTION`’ file is used to decide which packages to re-install when this happens). Packages declared in the ‘`Depends`’ field should not also be in the ‘`Imports`’ field. Version requirements can be specified, but will **not** be checked when the namespace is loaded (whereas they are checked by R CMD check).

The ‘`Suggests`’ field uses the same syntax as ‘`Depends`’ and lists packages that are not necessarily needed. This includes packages used only in examples, tests or vignettes (see [Section 1.4 \[Writing package vignettes\]](#), page 32), and packages loaded in the body of functions. E.g., suppose an example from package **foo** uses a dataset from package **bar**. Then it is not necessary to have **bar** use **foo** unless one wants to execute all the examples/tests/vignettes: it is useful to have **bar**, but not necessary. Version requirements can be specified, and will be used by R CMD check. Note that someone wanting to run the examples/tests/vignettes may not have a suggested package available (and it may not even be possible to install it for that platform), so it is helpful if the use of suggested packages is made conditional *via* `if(require(pkgname))`.

Finally, the ‘`Enhances`’ field lists packages “enhanced” by the package at hand, e.g., by providing methods for classes from these packages, or ways to handle objects from these packages (so several packages have ‘`Enhances: chron`’ because they can handle datetime objects from **chron** even though they prefer R’s native datetime functions). Version requirements can be specified, but are currently not used. Such packages cannot be required to check the package: any tests which use them must be conditional on the presence of the package. (If your tests use e.g. a dataset from another package it should be in ‘`Suggests`’ and not ‘`Enhances`’.)

The general rules are

- Packages whose namespace only is needed to load the package using `library(pkgname)` must be listed in the ‘`Imports`’ field and not in the ‘`Depends`’ field.
- Packages that need to be attached to successfully load the package using `library(pkgname)` must be listed in the ‘`Depends`’ field, only.
- All packages that are needed<sup>3</sup> to successfully run R CMD check on the package must be listed in one of ‘`Depends`’ or ‘`Suggests`’ or ‘`Imports`’. Packages used to run examples or

---

<sup>3</sup> This includes all packages directly called by `library` and `require` calls, as well as data obtained *via* `data(theirdata, package = "somepkg")` calls: R CMD check will warn about all of these. But there are subtler uses which it will not detect: e.g. if package A uses package B and makes use of functionality in



tests conditionally (e.g. *via* `if(require(pkgname))`) should be listed in ‘Suggests’ or ‘Enhances’. (This allows checkers to ensure that all the packages needed for a complete check are installed.)

In particular, large packages providing “only” data for examples or vignettes should be listed in ‘Suggests’ rather than ‘Depends’ in order to make lean installations possible.

Version dependencies in the ‘Depends’ field are used by `library` when it loads the package, and `install.packages` checks versions for the ‘Imports’ and (for `dependencies = TRUE`) ‘Suggests’ fields.

It is increasingly important that the information in these fields is complete and accurate: it is for example used to compute which packages depend on an updated package and which packages can safely be installed in parallel.

The ‘URL’ field may give a list of URLs separated by commas or whitespace, for example the homepage of the author or a page where additional material describing the software can be found. These URLs are converted to active hyperlinks in CRAN package listings.

The ‘BugReports’ field may contain a single URL to which bug reports about the package should be submitted. This URL will be used by `bug.reports` instead of sending an email to the maintainer.

Base and recommended packages (i.e., packages contained in the R source distribution or available from CRAN and recommended to be included in every binary distribution of R) have a ‘Priority’ field with value ‘base’ or ‘recommended’, respectively. These priorities must not be used by other packages.

An ‘Collate’ field can be used for controlling the collation order for the R code files in a package when these are processed for package installation. The default is to collate according to the ‘C’ locale. If present, the collate specification must list *all* R code files in the package (taking possible OS-specific subdirectories into account, see [Section 1.1.3 \[Package subdirectories\]](#), page 10) as a whitespace separated list of file paths relative to the ‘R’ subdirectory. Paths containing white space or quotes need to be quoted. An OS-specific collation field (‘Collate.unix’ or ‘Collate.windows’) will be used instead of ‘Collate’.

The ‘LazyData’ logical field controls whether the R datasets use lazy-loading. A ‘LazyLoad’ field was used in versions prior to 2.14.0, but now is ignored.

The ‘KeepSource’ logical field controls if the package code is sourced using `keep.source = TRUE` or `FALSE`: it might be needed exceptionally for a package designed to always be used with `keep.source = TRUE`.

The ‘ByteCompile’ logical field controls if the package code is byte-compiled on installation: the default is currently not to, so this may be useful for a package known to benefit particularly from byte-compilation (which can take quite a long time and increases the installed size of the package).

The ‘ZipData’ logical field used to control whether the automatic Windows build would zip up the data directory or not: set this to ‘no’ if your package will not work with a zipped data directory. (Setting to any other value is deprecated, and it is unused from R 2.13.0: but it might still be needed if the package can be installed under earlier versions of R.)

---

package B which uses package C which package B suggests or enhances, then package C needs to be in the ‘Suggests’ list for package A. Nor will undeclared uses in included files be reported, nor unconditional uses of packages listed under ‘Enhances’.

The ‘`BuildVignettes`’ logical field can be set to a false value to stop R CMD build from attempting to rebuild the vignettes, as well as preventing R CMD check from testing this. This should only be used exceptionally, for example if the PDFs need large figures which are not part of the package sources.

If the ‘`DESCRIPTION`’ file is not entirely in ASCII it should contain an ‘`Encoding`’ field specifying an encoding. This is used as the encoding of the ‘`DESCRIPTION`’ file itself and of the ‘`R`’ and ‘`NAMESPACE`’ files, and as the default encoding of ‘`.Rd`’ files. The examples are assumed to be in this encoding when running R CMD check, and it is used for the encoding of the CITATION file. Only encoding names `latin1`, `latin2` and `UTF-8` are known to be portable. (Do not specify an encoding unless one is actually needed: doing so makes the package *less* portable.)

The ‘`OS_type`’ field specifies the OS(es) for which the package is intended. If present, it should be one of `unix` or `windows`, and indicates that the package can only be installed on a platform with ‘`.Platform$OS.type`’ having that value.

The ‘`Type`’ field specifies the type of the package: see [Section 1.11 \[Package types\]](#), [page 53](#).

**Note:** There should be no ‘`Built`’ or ‘`Packaged`’ fields, as these are added by the package management tools.

One can add subject classifications for the content of the package using the fields ‘`Classification/ACM`’ (using the Computing Classification System of the Association for Computing Machinery, <http://www.acm.org/class/>), ‘`Classification/JEL`’ (the Journal of Economic Literature Classification System, [http://www.aeaweb.org/journal/jel\\_class\\_system.html](http://www.aeaweb.org/journal/jel_class_system.html)), or ‘`Classification/MSC`’ (the Mathematics Subject Classification of the American Mathematical Society, <http://www.ams.org/msc/>). The subject classifications should be comma-separated lists of the respective classification codes, e.g., ‘`Classification/ACM: G.4, H.2.8, I.5.1`’.

Finally, an ‘`Language`’ field can be used to indicate if the package documentation is not in English: this should be a comma-separated list of standard (not private use or grandfathered) IETF language tags as currently defined by RFC 5646 (<http://tools.ietf.org/html/rfc5646>, see also [http://en.wikipedia.org/wiki/IETF\\_language\\_tag](http://en.wikipedia.org/wiki/IETF_language_tag)), i.e., use language subtags which in essence are 2-letter ISO 639-1 ([http://en.wikipedia.org/wiki/ISO\\_639-1](http://en.wikipedia.org/wiki/ISO_639-1)) or 3-letter ISO 639-3 ([http://en.wikipedia.org/wiki/ISO\\_639-3](http://en.wikipedia.org/wiki/ISO_639-3)) language codes.

### 1.1.2 The ‘`INDEX`’ file

The optional file ‘`INDEX`’ contains a line for each sufficiently interesting object in the package, giving its name and a description (functions such as print methods not usually called explicitly might not be included). Normally this file is missing and the corresponding information is automatically generated from the documentation sources (using `tools::Rdindex()`) when installing from source.

Rather than editing this file, it is preferable to put customized information about the package into an overview man page (see [Section 2.1.4 \[Documenting packages\]](#), [page 63](#)) and/or a vignette (see [Section 1.4 \[Writing package vignettes\]](#), [page 32](#)).

### 1.1.3 Package subdirectories

The ‘R’ subdirectory contains R code files, only. The code files to be installed must start with an ASCII (lower or upper case) letter or digit and have one of the extensions<sup>4</sup> ‘.R’, ‘.S’, ‘.q’, ‘.r’, or ‘.s’. We recommend using ‘.R’, as this extension seems to be not used by any other software. It should be possible to read in the files using `source()`, so R objects must be created by assignments. Note that there need be no connection between the name of the file and the R objects created by it. Ideally, the R code files should only directly assign R objects and definitely should not call functions with side effects such as `require` and `options`. If computations are required to create objects these can use code ‘earlier’ in the package (see the ‘Collate’ field) plus functions in the ‘Depends’ packages provided that the objects created do not depend on those packages except *via* namespace imports.

Two exceptions are allowed: if the ‘R’ subdirectory contains a file ‘`sysdata.rda`’ (a saved image of R objects: please use suitable compression as suggested by `tools::resaveRdaFiles`) this will be lazy-loaded into the namespace/package environment – this is intended for system datasets that are not intended to be user-accessible *via* data. Also, files ending in ‘.in’ will be allowed in the ‘R’ directory to allow a ‘`configure`’ script to generate suitable files.

Only ASCII characters (and the control characters tab, formfeed, LF and CR) should be used in code files. Other characters are accepted in comments, but then the comments may not be readable in e.g. a UTF-8 locale. Non-ASCII characters in object names will normally<sup>5</sup> fail when the package is installed. Any byte will be allowed in a quoted character string but `\uxxxx` escapes should be used<sup>6</sup> for non-ASCII characters. However, non-ASCII character strings may not be usable in some locales and may display incorrectly in others.

Various R functions in a package can be used to initialize and clean up. See [Section 1.6.3 \[Load hooks\]](#), [page 40](#).<sup>7</sup>

The ‘man’ subdirectory should contain (only) documentation files for the objects in the package in *R documentation* (Rd) format. The documentation filenames must start with an ASCII (lower or upper case) letter or digit and have the extension ‘.Rd’ (the default) or ‘.rd’. Further, the names must be valid in ‘file://’ URLs, which means<sup>8</sup> they must be entirely ASCII and not contain ‘%’. See [Chapter 2 \[Writing R documentation files\]](#), [page 55](#), for more information. Note that all user-level objects in a package should be documented; if a package *pkg* contains user-level objects which are for “internal” use only,

<sup>4</sup> Extensions ‘.S’ and ‘.s’ arise from code originally written for S(-PLUS), but are commonly used for assembler code. Extension ‘.q’ was used for S, which at one time was tentatively called QPE.

<sup>5</sup> This is true for OSes which implement the ‘C’ locale: Windows’ idea of the ‘C’ locale uses the WinAnsi charset.

<sup>6</sup> the package needs to depend on R ( $\geq 2.10$ )

<sup>7</sup> Prior to R version 2.14.0, the `.First.lib` and `.Last.lib` functions handled these tasks in packages without namespaces. (In current R versions, *all* packages have namespaces.) To help with conversion of old packages, here is how they were handled: It was conventional to define these functions in a file called ‘`zzz.R`’. If `.First.lib` was defined in a package, it was called with arguments `libname` and `pkgname` after the package was loaded and attached. A common use was to call `library.dynam` inside `.First.lib` to load compiled code: another use was to call some functions with side effects. If `.Last.lib` existed in a package it was called (with argument the full path to the installed package) just before the package was detached.

<sup>8</sup> More precisely, they can contain the English alphanumeric characters and the symbols ‘\$ - \_ . + ! ’ ( ) , ; = &’.



it should provide a file `'pkg-internal.Rd'` which documents all such objects, and clearly states that these are not meant to be called by the user. See e.g. the sources for package **grid** in the R distribution for an example. Note that packages which use internal objects extensively should not export those objects from their namespace, when they do not need to be documented (see [Section 1.6 \[Package namespaces\]](#), page 38).

Having a `'man'` directory containing no documentation files may give an installation error.

The `'R'` and `'man'` subdirectories may contain OS-specific subdirectories named `'unix'` or `'windows'`.

The sources and headers for the compiled code are in `'src'`, plus optionally a file `'Makevars'` or `'Makefile'`. When a package is installed using R CMD INSTALL, `make` is used to control compilation and linking into a shared object for loading into R. There are default `make` variables and rules for this (determined when R is configured and recorded in `'R_HOME/etcR_ARCH/Makeconf'`), providing support for C, C++, FORTRAN 77, Fortran 9x<sup>9</sup>, Objective C and Objective C++<sup>10</sup> with associated extensions `'.c'`, `'.cc'` or `'.cpp'`, `'.f'`, `'.f90'` or `'.f95'`, `'.m'`, and `'.mm'` or `'.M'`, respectively. We recommend using `'.h'` for headers, also for C++<sup>11</sup> or Fortran 9x include files. (Use of extension `'.C'` for C++ is no longer supported.) Files in the `'src'` directory should not be hidden (start with a dot), and hidden files will under some versions of R be ignored.

It is not portable (and may not be possible at all) to mix all these languages in a single package, and we do not support using both C++ and Fortran 9x. Because R itself uses it, we know that C and FORTRAN 77 can be used together and mixing C and C++ seems to be widely successful.

If your code needs to depend on the platform there are certain defines which can be used in C or C++. On all Windows builds (even 64-bit ones) `'WIN32'` will be defined: on 64-bit Windows builds also `'WIN64'`, and on Mac OS X `'__APPLE__'` and `'__APPLE_CC__'` are defined.

The default rules can be tweaked by setting macros<sup>12</sup> in a file `'src/Makevars'` (see [Section 1.2.1 \[Using Makevars\]](#), page 18). Note that this mechanism should be general enough to eliminate the need for a package-specific `'src/Makefile'`. If such a file is to be distributed, considerable care is needed to make it general enough to work on all R platforms. If it has any targets at all, it should have an appropriate first target named `'all'` and a (possibly empty) target `'clean'` which removes all files generated by running `make` (to be used by `'R CMD INSTALL --clean'` and `'R CMD INSTALL --preclean'`). There are platform-specific file names on Windows: `'src/Makevars.win'` takes precedence over `'src/Makevars'` and `'src/Makefile.win'` must be used. Some `make` programs require makefiles to have a complete final line, including a newline.

A few packages use the `'src'` directory for purposes other than making a shared object (e.g. to create executables). Such packages should have files `'src/Makefile'` and `'src/Makefile.win'` (unless intended for only Unix-alikes or only Windows).

<sup>9</sup> Note that Ratfor is not supported. If you have Ratfor source code, you need to convert it to FORTRAN. Only FORTRAN-77 (which we write in upper case) is supported on all platforms, but most also support Fortran-95 (for which we use title case). If you want to ship Ratfor source files, please do so in a subdirectory of `'src'` and not in the main subdirectory.

<sup>10</sup> either or both of which may not be supported on particular platforms

<sup>11</sup> Using `' .hpp'`, although somewhat popular, is not guaranteed to be portable.

<sup>12</sup> the POSIX terminology, called 'make variables' by GNU make.

In very special cases packages may create binary files other than the shared objects/DLLs in the ‘src’ directory. Such files will not be installed in multi-arch setting since R CMD INSTALL `--libs-only` is used to merge multiple architectures and it only copies shared objects/DLLs. If a package wants to install other binaries (for example executable programs), it should provide an R script ‘src/install.libs.R’ which will be run as part of the installation in the `src` build directory *instead of* copying the shared objects/DLLs. The script is run in a separate R environment containing the following variables: `R_PACKAGE_NAME` (the name of the package), `R_PACKAGE_SOURCE` (the path to the source directory of the package), `R_PACKAGE_DIR` (the path of the target installation directory of the package), `R_ARCH` (the arch-dependent part of the path), `SHLIB_EXT` (the extension of shared objects) and `WINDOWS` (TRUE on Windows, FALSE elsewhere). Something close to the default behavior could be replicated with the following ‘src/install.libs.R’ file:

```
files <- Sys.glob(paste("*", SHLIB_EXT, sep=''))
libarch <- if (nzchar(R_ARCH)) paste('libs', R_ARCH, sep='') else 'libs'
dest <- file.path(R_PACKAGE_DIR, libarch)
dir.create(dest, recursive = TRUE, showWarnings = FALSE)
file.copy(files, dest, overwrite = TRUE)
```

The ‘data’ subdirectory is for data files: See [Section 1.1.5 \[Data in packages\]](#), page 13.

The ‘demo’ subdirectory is for R scripts (for running *via* `demo()`) that demonstrate some of the functionality of the package. Demos may be interactive and are not checked automatically, so if testing is desired use code in the ‘tests’ directory to achieve this. The script files must start with a (lower or upper case) letter and have one of the extensions ‘.R’ or ‘.r’. If present, the ‘demo’ subdirectory should also have a ‘00Index’ file with one line for each demo, giving its name and a description separated by white space. (Note that it is not possible to generate this index file automatically.)

The contents of the ‘inst’ subdirectory will be copied recursively to the installation directory. Subdirectories of ‘inst’ should not interfere with those used by R (currently, ‘R’, ‘data’, ‘demo’, ‘exec’, ‘libs’, ‘man’, ‘help’, ‘html’ and ‘Meta’, and earlier versions used ‘latex’, ‘R-ex’). The copying of the ‘inst’ happens after ‘src’ is built so its ‘Makefile’ can create files to be installed. Prior to R 2.12.2, the files were installed on POSIX platforms with the permissions in the package sources, so care should be taken to ensure these are not too restrictive: R CMD build will make suitable adjustments. To exclude files from being installed, one can specify a list of exclude patterns in file ‘.Rinstignore’ in the top-level source directory. These patterns should be Perl-like regular expressions (see the help for `regexp` in R for the precise details), one per line, to be matched<sup>13</sup> against the file and directory paths, e.g. ‘doc/.\*.png\$’ will exclude all PNG files in ‘inst/doc’ based on the (lower-case) extension.

Note that with the exceptions of ‘INDEX’, ‘LICENSE’/‘LICENCE’ and ‘NEWS’, information files at the top level of the package will *not* be installed and so not be known to users of Windows and Mac OS X compiled packages (and not seen by those who use R CMD INSTALL or `install.packages` on the tarball). So any information files you wish an end user to see should be included in ‘inst’. Note that if the named exceptions also occur in ‘inst’, the versions in ‘inst’ will be that seen in the installed package.

---

<sup>13</sup> case-insensitively on Windows.

One thing you might like to add to ‘inst’ is a ‘CITATION’ file for use by the `citation` function.

Subdirectory ‘tests’ is for additional package-specific test code, similar to the specific tests that come with the R distribution. Test code can either be provided directly in a ‘.R’ file, or *via* a ‘.Rin’ file containing code which in turn creates the corresponding ‘.R’ file (e.g., by collecting all function objects in the package and then calling them with the strangest arguments). The results of running a ‘.R’ file are written to a ‘.Rout’ file. If there is a corresponding<sup>14</sup> ‘.Rout.save’ file, these two are compared, with differences being reported but not causing an error. The directory ‘tests’ is copied to the check area, and the tests are run with the copy as the working directory and with `R_LIBS` set to ensure that the copy of the package installed during testing will be found by `library(pkg_name)`. Note that the package-specific tests are run in a vanilla R session without setting the random-number seed, so tests which use random numbers will need to set the seed to obtain reproducible results (and it can be helpful to do so in all cases, to avoid occasional failures when tests are run).

If ‘tests’ has a subdirectory ‘Examples’ containing a file `pkg-Ex.Rout.save`, this is compared to the output file for running the examples when the latter are checked.

Subdirectory ‘exec’ could contain additional executable scripts the package needs, typically scripts for interpreters such as the shell, Perl, or Tcl. This mechanism is currently used only by a very few packages, and still experimental. NB: only files (and not directories) under ‘exec’ are installed (and those with names starting with a dot are ignored), and they are all marked as executable (mode 755, moderated by ‘umask’) on POSIX platforms. Note too that this may not be suitable for executable *programs* since some platforms (including Mac OS X and Windows) support multiple architectures using the same installed package directory.

Subdirectory ‘po’ is used for files related to *localization*: see [Section 1.9 \[Internationalization\]](#), page 50.

### 1.1.4 Package bundles

Support for package bundles was removed in R 2.11.0.

### 1.1.5 Data in packages

The ‘data’ subdirectory is for data files, either to be made available *via* lazy-loading or for loading using `data()`. (The choice is made by the ‘LazyData’ field in the ‘DESCRIPTION’ file: the default is not to do so.) It should not be used for other data files needed by the package, and the convention has grown up to use directory ‘inst/extdata’ for such files.

Data files can have one of three types as indicated by their extension: plain R code (‘.R’ or ‘.r’), tables (‘.tab’, ‘.txt’, or ‘.csv’, see `?data` for the file formats, and note that ‘.csv’ is **not** the standard<sup>15</sup> CSV format), or `save()` images (‘.RData’ or ‘.rda’). The files should not be hidden (have names starting with a dot). Note that R code should be “self-sufficient” and not make use of extra functionality provided by the package, so that the data file can also be used without having to load the package.

<sup>14</sup> The best way to generate such a file is to copy the ‘.Rout’ from a successful run of `R CMD check`. If you want to generate it separately, do run R with options ‘--vanilla --slave’ and with environment variable `LANGUAGE=en` set to get messages in English.

<sup>15</sup> e.g. <http://tools.ietf.org/html/rfc4180>.

Images (extensions `‘.RData’` or `‘.rda’`) can contain references to the namespaces of packages that were used to create them. Preferably there should be no such references in data files, and in any case they should only be to packages listed in the `Depends` and `Imports` fields, as otherwise it may be impossible to install the package. To check for such references, load all the images into a vanilla R session, and look at the output of `loadedNamespaces()`.

If your data files are large and you are not using `‘LazyData’` you can speed up installation by providing a file `‘datalist’` in the `‘data’` subdirectory. This should have one line per topic that `data()` will find, in the format `‘foo’` if `data(foo)` provides `‘foo’`, or `‘foo: bar bah’` if `data(foo)` provides `‘bar’` and `‘bah’`. R CMD build will automatically add a `‘datalist’` file to `‘data’` directories of over 1Mb, using the function `tools::add_datalist`.

Tables (`‘.tab’`, `‘.txt’`, or `‘.csv’` files) can be compressed by `gzip`, `bzip2` or `xz`, optionally with additional extension `‘.gz’`, `‘.bz2’` or `‘.xz’`. However, such files can only be used with R 2.10.0 or later, and so the package should have an appropriate `‘Depends’` entry in its DESCRIPTION file.

If your package is to be distributed, do consider the resource implications of large datasets for your users: they can make packages very slow to download and use up unwelcome amounts of storage space, as well as taking many seconds to load. It is normally best to distribute large datasets as `‘.rda’` images prepared by `save(, compress = TRUE)` (the default): there is no excuse for distributing ASCII saves. Using `bzip2` or `xz` compression will usually reduce the size of both the package tarball and the installed package, in some cases by a factor of two or more. However, such compression can only be used with R 2.10.0 or later, and so the package should have an appropriate `‘Depends’` entry in its DESCRIPTION file.

Package **tools** has a couple of functions to help with data images: `checkRdaFiles` reports on the way the image was saved, and `resaveRdaFiles` will re-save with a different type of compression, including choosing the best type for that particular image.

Some packages using `‘LazyData’` will benefit from using a form of compression other than `gzip` in the installed lazy-loading database. This can be selected by the `‘--data-compress’` option to R CMD INSTALL or by using the `‘LazyDataCompression’` field in the `‘DESCRIPTION’` file. Useful values are `bzip2`, `xz` and the default, `gzip`. The only way to discover which is best is to try them all and look at the size of the `‘pkgname/data/Rdata.rdb’` file.

Lazy-loading is not supported for very large datasets (those which when serialized exceed 2GB).

### 1.1.6 Non-R scripts in packages

Code which needs to be compiled (C, C++, FORTRAN, Fortran 95 ...) is included in the `‘src’` subdirectory and discussed elsewhere in this document.

Subdirectory `‘exec’` could be used for scripts for interpreters such as the shell (e.g. **arulesSequences**), BUGS, Java, JavaScript, Matlab, Perl (**FEST**), php (**amap**), Python or Tcl, or even R. However, it seems more common to use the `‘inst’` directory, for example `‘AMA/inst/java’`, `‘WriteXLS/inst/Perl’`, `‘Amelia/inst/tklibs’`, `‘CGIwithR/inst/cgi-bin’`, `‘NMF/inst/matlab’` and `‘emdbook/inst/BUGS’`.

If your package requires one of these interpreters or an extension then this should be declared in the `‘SystemRequirements’` field of its `‘DESCRIPTION’` file. Windows users should be aware that the Tcl extensions `‘BWidget’` and `‘Tktable’` which are included with the R

for Windows installer *are* extensions and do need to be declared. ‘Tktable’ does ship as part of the Tcl/Tk provided on CRAN for Mac OS X, but you will need to tell your users how to make use of it:

```
> addTclPath('/usr/local/lib/Tktable2.9')
> tclRequire('Tktable')
<Tcl> 2.9
```

## 1.2 Configure and cleanup

Note that most of this section is specific to Unix-alikes: see the comments later on about the Windows port of R.

If your package needs some system-dependent configuration before installation you can include an executable (Bourne shell) script ‘**configure**’ in your package which (if present) is executed by R CMD INSTALL before any other action is performed. This can be a script created by the Autoconf mechanism, but may also be a script written by yourself. Use this to detect if any nonstandard libraries are present such that corresponding code in the package can be disabled at install time rather than giving error messages when the package is compiled or used. To summarize, the full power of Autoconf is available for your extension package (including variable substitution, searching for libraries, etc.).

Under a Unix-alike only, an executable (Bourne shell) script ‘**cleanup**’ is executed as the last thing by R CMD INSTALL if option ‘--clean’ was given, and by R CMD build when preparing the package for building from its source. It can be used to clean up the package source tree: in particular, it should remove all files created by **configure**.

As an example consider we want to use functionality provided by a (C or FORTRAN) library **foo**. Using Autoconf, we can create a configure script which checks for the library, sets variable **HAVE\_FOO** to **TRUE** if it was found and with **FALSE** otherwise, and then substitutes this value into output files (by replacing instances of ‘@HAVE\_FOO@’ in input files with the value of **HAVE\_FOO**). For example, if a function named **bar** is to be made available by linking against library **foo** (i.e., using ‘-lfoo’), one could use

```
AC_CHECK_LIB(foo, fun, [HAVE_FOO=TRUE], [HAVE_FOO=FALSE])
AC_SUBST(HAVE_FOO)
.....
AC_CONFIG_FILES([foo.R])
AC_OUTPUT
```

in ‘**configure.ac**’ (assuming Autoconf 2.50 or later).

The definition of the respective R function in ‘**foo.R.in**’ could be

```
foo <- function(x) {
  if(!@HAVE_FOO@)
    stop("Sorry, library 'foo' is not available")
  ...
}
```

From this file **configure** creates the actual R source file ‘**foo.R**’ looking like

```
foo <- function(x) {
  if(!FALSE)
    stop("Sorry, library 'foo' is not available")
  ...
}
```

if library `foo` was not found (with the desired functionality). In this case, the above R code effectively disables the function.

One could also use different file fragments for available and missing functionality, respectively.

You will very likely need to ensure that the same C compiler and compiler flags are used in the ‘`configure`’ tests as when compiling R or your package. Under a Unix-alike, you can achieve this by including the following fragment early in ‘`configure.ac`’

```
: ${R_HOME}='R RHOME'
if test -z "${R_HOME}"; then
  echo "could not determine R_HOME"
  exit 1
fi
CC='${R_HOME}/bin/R' CMD config CC
CFLAGS='${R_HOME}/bin/R' CMD config CFLAGS
CPPFLAGS='${R_HOME}/bin/R' CMD config CPPFLAGS
```

(Using ‘`${R_HOME}/bin/R`’ rather than just ‘`R`’ is necessary in order to use the correct version of R when running the script as part of R CMD INSTALL, and the quotes since ‘`${R_HOME}`’ might contain spaces.)

If your code does load checks then you may also need

```
LDFLAGS='${R_HOME}/bin/R' CMD config LDFLAGS
```

and packages written with C++ need to pick up the details for the C++ compiler and switch the current language to C++ by

```
AC_LANG(C++)
```

The latter is important, as for example C headers may not be available to C++ programs or may not be written to avoid C++ name-mangling.

You can use R CMD `config` for getting the value of the basic configuration variables, or the header and library flags necessary for linking against R, see *R CMD config --help* for details.

To check for an external BLAS library using the `ACX_BLAS` macro from the official Autoconf Macro Archive, one can simply do

```
F77='${R_HOME}/bin/R' CMD config F77
AC_PROG_F77
FLIBS='${R_HOME}/bin/R' CMD config FLIBS
ACX_BLAS([], AC_MSG_ERROR([could not find your BLAS library], 1))
```

Note that `FLIBS` as determined by R must be used to ensure that FORTRAN 77 code works on all R platforms. Calls to the Autoconf macro `AC_F77_LIBRARY_LDFLAGS`, which would overwrite `FLIBS`, must not be used (and hence e.g. removed from `ACX_BLAS`). (Recent versions of Autoconf in fact allow an already set `FLIBS` to override the test for the FORTRAN linker flags. Also, recent versions of R can detect external BLAS and LAPACK libraries.)

You should bear in mind that the `configure` script will not be used on Windows systems. If your package is to be made publicly available, please give enough information for a user on a non-Unix-alike platform to configure it manually, or provide a ‘`configure.win`’ script to be used on that platform. (Optionally, there can be a ‘`cleanup.win`’ script. Both should



be shell scripts to be executed by `ash`, which is a minimal version of Bourne-style `sh`.) When `'configure.win'` is run the environment variables `R_HOME` (which uses `/` as the file separator) and `R_ARCH` will be set. Use `R_ARCH` to decide if this is a 64-bit build (its value there is `'x64'`) and to install DLLs to the correct place (`${R_HOME}/libs${R_ARCH}`). Use `R_ARCH_BIN` to find the correct place under the `'bin'` directory, e.g. `${R_HOME}/bin${R_ARCH_BIN}/Rscript.exe`.

In some rare circumstances, the configuration and cleanup scripts need to know the location into which the package is being installed. An example of this is a package that uses C code and creates two shared object/DLLs. Usually, the object that is dynamically loaded by R is linked against the second, dependent, object. On some systems, we can add the location of this dependent object to the object that is dynamically loaded by R. This means that each user does not have to set the value of the `LD_LIBRARY_PATH` (or equivalent) environment variable, but that the secondary object is automatically resolved. Another example is when a package installs support files that are required at run time, and their location is substituted into an R data structure at installation time. (This happens with the Java Archive files in the Omegahat **SJava** package.) The names of the top-level library directory (i.e., specifiable *via* the `'-l'` argument) and the directory of the package itself are made available to the installation scripts *via* the two shell/environment variables `R_LIBRARY_DIR` and `R_PACKAGE_DIR`. Additionally, the name of the package (e.g. `'survival'` or `'MASS'`) being installed is available from the environment variable `R_PACKAGE_NAME`. (Currently the value of `R_PACKAGE_DIR` is always `${R_LIBRARY_DIR}/${R_PACKAGE_NAME}`, but this used not to be the case when versioned installs were allowed. Its main use is in `'configure.win'` scripts for the installation path of external software's DLLs.) Note that the value of `R_PACKAGE_DIR` may contain spaces and other shell-unfriendly characters, and so should be quoted in makefiles and configure scripts.

One of the more tricky tasks can be to find the headers and libraries of external software. One tool which is increasingly available on Unix-alikes (but not Mac OS X) to do this is `pkg-config`. The `'configure'` script will need to test for the presence of the command itself (see for example package **Cairo**), and if present it can be asked if the software is installed, of a suitable version and for compilation/linking flags by e.g.

```
$ pkg-config --exists 'QtCore >= 4.0.0' # check the status
$ pkg-config --modversion QtCore
4.7.1
$ pkg-config --cflags QtCore
-DQT_SHARED -I/usr/include/QtCore
$ pkg-config --libs QtCore
-lQtCore
```

Note that `pkg-config --libs` gives the information required to link against the default version of that library (usually the dynamic one), and `pkg-config --static` is needed if the static library is to be used.

Sometimes the name by which the software is known to `pkg-config` is not what one might expect (e.g. `'gtk+-2.0'` even for 2.22). To get a complete list use

```
pkg-config --list-all | sort
```

### 1.2.1 Using ‘Makevars’

Sometimes writing your own ‘configure’ script can be avoided by supplying a file ‘Makevars’: also one of the most common uses of a ‘configure’ script is to make ‘Makevars’ from ‘Makevars.in’.

A ‘Makevars’ file is a makefile and is used as one of several makefiles by R CMD SHLIB (which is called by R CMD INSTALL to compile code in the ‘src’ directory). It should be written if at all possible in a portable style, in particular (except for ‘Makevars.win’) without the use of GNU extensions.

The most common use of a ‘Makevars’ file is to set additional preprocessor options (for example include paths) for C/C++ files *via* PKG\_CPPFLAGS, and additional compiler flags by setting PKG\_CFLAGS, PKG\_CXXFLAGS, PKG\_FFLAGS or PKG\_FCFLAGS, for C, C++, FORTRAN or Fortran 9x respectively (see [Section 5.5 \[Creating shared objects\]](#), page 99).

**NB:** Include paths are preprocessor options, not compiler options, and **must** be set in PKG\_CPPFLAGS as otherwise platform-specific paths (e.g. ‘-I/usr/local/include’) will take precedence.

‘Makevars’ can also be used to set flags for the linker, for example ‘-L’ and ‘-l’ options, *via* PKG\_LIBS.

When writing a ‘Makevars’ file for a package you intend to distribute, take care to ensure that it is not specific to your compiler: flags such as ‘-O2 -Wall -pedantic’ are all specific to GCC.

There are some macros<sup>16</sup> which are set whilst configuring the building of R itself and are stored in ‘R\_HOME/etcR\_ARCH/Makeconf’. That makefile is included as a ‘Makefile’ *after* ‘Makevars[.win]’, and the macros it defines can be used in macro assignments and make command lines in the latter. These include

**FLIBS**      A macro containing the set of libraries need to link FORTRAN code. This may need to be included in PKG\_LIBS: it will normally be included automatically if the package contains FORTRAN source files.

**BLAS\_LIBS**      A macro containing the BLAS libraries used when building R. This may need to be included in PKG\_LIBS. Beware that if it is empty then the R executable will contain all the double-precision and double-complex BLAS routines, but no single-precision or complex routines. If BLAS\_LIBS is included, then FLIBS also needs to be<sup>17</sup>, as most BLAS libraries are written at least partially in FORTRAN.

**LAPACK\_LIBS**      A macro containing the LAPACK libraries (and paths where appropriate) used when building R. This may need to be included in PKG\_LIBS. It may point to a dynamic library `libRlapack` which contains all the double-precision LAPACK routines as well as those double-complex LAPACK and BLAS routines needed to build R, or it may point to an external LAPACK library, or may be empty if an external BLAS library also contains LAPACK.

<sup>16</sup> in POSIX parlance: GNU **make** calls these ‘make variables’.

<sup>17</sup> at least on Unix-alikes: the Windows build currently resolves such dependencies to a static FORTRAN library when ‘Rblas.dll’ is built.



[There is no guarantee that the LAPACK library will provide more than all the double-precision and double-complex routines, and some do not provide all the auxiliary routines.]

For portability, the macros `BLAS_LIBS` and `FLIBS` should always be included *after* `LAPACK_LIBS`.

#### `SAFE_FFLAGS`

A macro containing flags which are needed to circumvent over-optimization of FORTRAN code: it is typically `'-g -O2 -ffloat-store'` on 'ix86' platforms using `gfortran`. Note that this is **not** an additional flag to be used as part of `PKG_FFLAGS`, but a replacement for `FFLAGS`, and that it is intended for the FORTRAN-77 compiler 'F77' and not necessarily for the Fortran 90/95 compiler 'FC'. See the example later in this section.

Setting certain macros in 'Makevars' will prevent R CMD SHLIB setting them: in particular if 'Makevars' sets 'OBJECTS' it will not be set on the `make` command line. This can be useful in conjunction with implicit rules to allow other types of source code to be compiled and included in the shared object. It can also be used to control the set of files which are compiled, either by excluding some files in 'src' or including some files in subdirectories. For example

```
OBJECTS = 4dfp/endianio.o 4dfp/Getifh.o R4dfp-object.o
```

Note that 'Makevars' should not normally contain targets, as it is included before the default makefile and `make` will call the first target, intended to be `all` in the default makefile. If you really need to circumvent that, use a suitable (phony) target `all` before any actual targets in 'Makevars.[win]': for example package **fastICA** has

```
PKG_LIBS = @BLAS_LIBS@

SLAMC_FFLAGS=$(R_XTRA_FFLAGS) $(FPICFLAGS) $(SHLIB_FFLAGS) $(SAFE_FFLAGS)

all: $(SHLIB)

slamc.o: slamc.f
    $(F77) $(SLAMC_FFLAGS) -c -o slamc.o slamc.f
```

needed to ensure that the LAPACK routines find some constants without infinite looping. The Windows equivalent is

```
all: $(SHLIB)

slamc.o: slamc.f
    $(F77) $(SAFE_FFLAGS) -c -o slamc.o slamc.f
```

(since the other macros are all empty on that platform, and R's internal BLAS is not used). Note that the first target in 'Makevars' will be called, but for back-compatibility it is best named `all`.

If you want to create and then link to a library, say using code in a subdirectory, use something like

```
.PHONY: all mylibs

all: $(SHLIB)
$(SHLIB): mylibs

mylibs:
```

```
(cd subdir; make)
```

Be careful to create all the necessary dependencies, as there is a no guarantee that the dependencies of `all` will be run in a particular order (and some of the CRAN build machines use multiple CPUs and parallel makes).

Note that on Windows it is required that `'Makevars.win'` does create a DLL: this is needed as it is the only reliable way to ensure that building a DLL succeeded. If you want to use the `'src'` directory for some purpose other than building a DLL, use a `'Makefile.win'` file.

It is sometimes useful to have a target `'clean'` in `'Makevars'` or `'Makevars.win'`: this will be used by R CMD `build` to clean up (a copy of) the package sources. When it is run by `build` it will have fewer macros set, in particular not `$(SHLIB)`, nor `$(OBJECTS)` unless set in the file itself. It would also be possible to add tasks to the target `'shlib-clean'` which is run by R CMD `INSTALL` and R CMD `SHLIB` with options `'--clean'` and `'--preclean'`.

If you want to run R code in `'Makevars'`, e.g. to find configuration information, please do ensure that you use the correct copy of R or `Rscript`: there might not be one in the path at all, or it might be the wrong version or architecture. The correct way to do this is *via*

```
"$(R_HOME)/bin$(R_ARCH_BIN)/Rscript" filename
"$(R_HOME)/bin$(R_ARCH_BIN)/Rscript" -e 'R expression'
```

where `$(R_ARCH_BIN)` is only needed currently on Windows.

Environment or make variables can be used to select different macros for 32- and 64-bit code, for example (GNU make syntax, allowed on Windows)

```
ifeq "$(WIN)" "64"
PKG_LIBS = value for 64-bit Windows
else
PKG_LIBS = value for 32-bit Windows
endif
```

On Windows there is normally a choice between linking to an import library or directly to a DLL. Where possible, the latter is much more reliable: import libraries are tied to a specific toolchain, and in particular on 64-bit Windows two different conventions have been commonly used. So for example instead of

```
PKG_LIBS = -L$(XML_DIR)/lib -lxml2
```

one can use

```
PKG_LIBS = -L$(XML_DIR)/bin -lxml2
```

since on Windows `-lxxx` will look in turn for

```
libxxx.dll.a
xxx.dll.a
libxxx.a
xxx.lib
libxxx.dll
xxx.dll
```

where the first and second are conventionally import libraries, the third and fourth often static libraries (with `.lib` intended for Visual C++), but might be import libraries. See for example <http://sourceware.org/binutils/docs-2.20/ld/WIN32.html#WIN32>.

The fly in the ointment is that the DLL might not be named ‘libxxx.dll’, and in fact on 32-bit Windows there is a ‘libxml2.dll’ whereas on one build for 64-bit Windows the DLL is called ‘libxml2-2.dll’. Using import libraries can cover over these differences but can cause equal difficulties.

If static libraries are available they can save a lot of problems with run-time finding of DLLs, especially when binary packages are to be distributed and even more when these support both architectures. Where using DLLs is unavoidable we normally arrange (*via* ‘configure.win’) to ship them in the same directory as the package DLL.

### 1.2.1.1 OpenMP support

As from R 2.13.0 there is some support for packages which wish to use OpenMP<sup>18</sup>. The `make` macros

```
SHLIB_OPENMP_CFLAGS
SHLIB_OPENMP_CXXFLAGS
SHLIB_OPENMP_FCFLAGS
SHLIB_OPENMP_FFLAGS
```

are available for use in ‘src/Makevars’ or ‘src/Makevars.win’.<sup>19</sup> Include the appropriate macro in `PKG_CFLAGS`, `PKG_CPPFLAGS` and so on, and also in `PKG_LIBS`. C/C++ code that needs to be conditioned on the use of OpenMP can be used inside `#ifdef SUPPORT_OPENMP`, a macro defined in the header ‘Rconfig.h’ (see [Section 6.13 \[Platform and version information\]](#), page 144): however the use of OpenMP is most often indicated by ‘`#pragma`’ statements.

For example, a package with C code written for OpenMP should have in ‘src/Makevars’ the lines

```
PKG_CFLAGS = $(SHLIB_OPENMP_CFLAGS)
PKG_LIBS = $(SHLIB_OPENMP_CFLAGS)
```

There is nothing to say what version of OpenMP is supported: version 3.0 (May 2008) is supported by recent versions of the main platforms (but note that Mac OS X binaries are currently built for 10.5 using compilers which support version 2.5), but portable packages cannot assume that end users have recent versions (there are some years-old versions of Linux in use), and it may be safest to assume version 2.5.

OpenMP support is expected to be available on Windows in the toolchain used for R 2.15.0 and these macros will be set appropriately. It is not available on Windows prior to R 2.14.2.

The performance of OpenMP varies substantially between platforms. Both the Mac OS X and Windows implementations have substantial overheads and are only beneficial if quite substantial tasks are run in parallel.

Calling any of the R API from threaded code is ‘for experts only’: they will need to read the source code to determine if it is thread-safe.

<sup>18</sup> <http://www.openmp.org/>, <http://en.wikipedia.org/wiki/OpenMP>, <https://computing.llnl.gov/tutorials/openMP/>

<sup>19</sup> But OpenMP support is not available with the default toolchain used prior to R 2.14.2 on Windows.

### 1.2.1.2 Using pthreads

There is no direct support for the POSIX threads (more commonly known as **pthreads**): by the time we considered adding it several packages were using it unconditionally so it seems that nowadays it is universally available on POSIX operating systems (hence not Windows).

For reasonably recent versions of **gcc** the correct specification is

```
PKG_CPPFLAGS = -pthread
PKG_LIBS = -pthread
```

(and the plural version is also accepted on some systems/versions). For other platforms the specification is

```
PKG_CPPFLAGS = -D_REENTRANT
PKG_LIBS = -lpthread
```

(and note that the library name is singular). This is what ‘**-pthread**’ does on all known current platforms (although earlier version of OpenBSD used a different library name).

For a tutorial see <https://computing.llnl.gov/tutorials/pthreads/>.

POSIX threads are not normally used on Windows, which has its own native concepts of threads. However, there are two projects implementing **pthreads** on top of Windows, **pthreads-w32** and **winpthreads** (a recent part of the MinGW-w64 project). Both implement **libpthreads** as an import library for a DLL.

Whether Windows toolchains implement **pthreads** is up to the toolchain provider. One issue has been licenses: **pthreads-w32** is licensed under LGPL which requires source code to be made available. The toolchains used to compile R prior to version 2.14.2 do not contain **pthreads**, although in some cases **pthreads-w32** could be retro-fitted. As from R 2.14.2 a **make** variable **SHLIB\_PTHREAD\_FLAGS** is available: this should be included in both **PKG\_CPPFLAGS** (or the Fortran or F9x equivalents) and **PKG\_LIBS**.

The presence of a working **pthreads** implementation cannot be unambiguously determined without testing for yourself: however, that ‘**\_REENTRANT**’ is defined<sup>20</sup> in C/C++ code is a good indication.

See also the comments on thread-safety and performance under OpenMP: on all known R platforms OpenMP is implemented *via* **pthreads** and the known performance issues are in the latter.

### 1.2.1.3 Compiling in sub-directories

Package authors fairly often want to organize code in sub-directories of ‘**src**’, for example if they are including a separate piece of external software to which this is an R interface.

One simple way is simply to set **OBJECTS** to be all the objects that need to be compiled, including in sub-directories. For example, CRAN package **RSiena** has

```
SOURCES = $(wildcard data/*.cpp network/*.cpp utils/*.cpp model/*.cpp model/**/*.cpp mo
OBJECTS = siena07utilities.o siena07internals.o siena07setup.o siena07models.o $(SOURC
```

One problem with that approach is that unless GNU make extensions are used, the source files need to be listed and kept up-to-date. As in the following from CRAN package **lossDev**:

<sup>20</sup> some Windows toolchains have the typo ‘**\_REENTRANCE**’ instead.

```

OBJECTS.samplers = samplers/ExpandableArray.o samplers/Knots.o \
  samplers/RJumpSpline.o samplers/RJumpSplineFactory.o \
  samplers/RealSlicerOV.o samplers/SliceFactoryOV.o samplers/MNorm.o
OBJECTS.distributions = distributions/DSpline.o \
  distributions/DChisqrOV.o distributions/DTOV.o \
  distributions/DNormOV.o distributions/DUnifOV.o distributions/RScalarDist.o
OBJECTS.root = RJump.o

```

```

OBJECTS = $(OBJECTS.samplers) $(OBJECTS.distributions) $(OBJECTS.root)

```

Where the subdirectory is self-contained code with a suitable makefile, the best approach is something like

```

PKG_LIBS = -LCsdp/lib -lsdp $(LAPACK_LIBS) $(BLAS_LIBS) $(FLIBS)

$(SHLIB): Csdp/lib/libsdp.a

Csdp/lib/libsdp.a
    @(cd Csdp/lib && $(MAKE) libsdp.a \
      CC="$(CC)" CFLAGS="$(CFLAGS) $(CPICFLAGS)" AR="$(AR)" RANLIB="$(RANLIB)")

```

Note the quotes: the macros can contain spaces, e.g. `gcc -m64 -std=gnu99`. Several authors have forgotten about parallel makes: the static library in the subdirectory must be made before the shared library and so must depend on the latter. Others forget the need for position-independent code.

We really do not recommend using a ‘src/Makefile’ instead on ‘src/Makevars’, and as the example above shows, it is not necessary.

## 1.2.2 Configure example

It may be helpful to give an extended example of using a ‘configure’ script to create a ‘src/Makevars’ file: this is based on that in the **RODBC** package.

The ‘configure.ac’ file follows: ‘configure’ is created from this by running `autoconf` in the top-level package directory (containing ‘configure.ac’).

```

AC_INIT([RODBC], 1.1.8) dnl package name, version

dnl A user-specifiable option
odbc_mgr=""
AC_ARG_WITH([odbc-manager],
    AC_HELP_STRING([--with-odbc-manager=MGR],
        [specify the ODBC manager, e.g. odbc or iodbc]),
    [odbc_mgr=$withval])

if test "$odbc_mgr" = "odbc" ; then
    AC_PATH_PROGS(ODBC_CONFIG, odbc_config)
fi

dnl Select an optional include path, from a configure option
dnl or from an environment variable.
AC_ARG_WITH([odbc-include],
    AC_HELP_STRING([--with-odbc-include=INCLUDE_PATH],
        [the location of ODBC header files]),
    [odbc_include_path=$withval])

```

```

RODBC_CPPFLAGS="-I."
if test [ -n "$odbc_include_path" ] ; then
    RODBC_CPPFLAGS="-I. -I${odbc_include_path}"
else
    if test [ -n "${ODBC_INCLUDE}" ] ; then
        RODBC_CPPFLAGS="-I. -I${ODBC_INCLUDE}"
    fi
fi

dnl ditto for a library path
AC_ARG_WITH([odbc-lib],
             AC_HELP_STRING([--with-odbc-lib=LIB_PATH],
                             [the location of ODBC libraries]),
             [odbc_lib_path=$withval])
if test [ -n "$odbc_lib_path" ] ; then
    LIBS="-L$odbc_lib_path ${LIBS}"
else
    if test [ -n "${ODBC_LIBS}" ] ; then
        LIBS="-L${ODBC_LIBS} ${LIBS}"
    else
        if test -n "${ODBC_CONFIG}"; then
            odbc_lib_path='odbc_config --libs | sed s/-lodbc//'
            LIBS="${odbc_lib_path} ${LIBS}"
        fi
    fi
fi

dnl Now find the compiler and compiler flags to use
: ${R_HOME='R RHOME'}
if test -z "${R_HOME}"; then
    echo "could not determine R_HOME"
    exit 1
fi
CC='${R_HOME}/bin/R' CMD config CC
CPP='${R_HOME}/bin/R' CMD config CPP
CFLAGS='${R_HOME}/bin/R' CMD config CFLAGS
CPPFLAGS='${R_HOME}/bin/R' CMD config CPPFLAGS
AC_PROG_CC
AC_PROG_CPP

if test -n "${ODBC_CONFIG}"; then
    RODBC_CPPFLAGS='odbc_config --cflags'
fi
CPPFLAGS="${CPPFLAGS} ${RODBC_CPPFLAGS}"

dnl Check the headers can be found
AC_CHECK_HEADERS(sql.h sql.h)
if test "${ac_cv_header_sql_h}" = no ||
    test "${ac_cv_header_sql_h}" = no; then
    AC_MSG_ERROR("ODBC headers sql.h and sql.h not found")
fi

dnl search for a library containing an ODBC function
if test [ -n "${odbc_mgr}" ] ; then
    AC_SEARCH_LIBS(SQLTables, ${odbc_mgr}, ,
                  AC_MSG_ERROR("ODBC driver manager ${odbc_mgr} not found"))
else

```

```

        AC_SEARCH_LIBS(SQLTables, odbc odbc32 iodbc, ,
            AC_MSG_ERROR("no ODBC driver manager found"))
    fi

    dnl for 64-bit ODBC need SQL[U]LEN, and it is unclear where they are defined.
    AC_CHECK_TYPES([SQLLEN, SQLULEN], , , [# include <sql.h>])
    dnl for unixODBC header
    AC_CHECK_SIZEOF(long, 4)

    dnl substitute RODB_CPPFLAGS and LIBS
    AC_SUBST(RODBC_CPPFLAGS)
    AC_SUBST(LIBS)
    AC_CONFIG_HEADERS([src/config.h])
    dnl and do substitution in the src/Makevars.in and src/config.h
    AC_CONFIG_FILES([src/Makevars])
    AC_OUTPUT

```

where ‘src/Makevars.in’ would be simply

```

PKG_CPPFLAGS = @RODBC_CPPFLAGS@
PKG_LIBS = @LIBS@

```

A user can then be advised to specify the location of the ODBC driver manager files by options like (lines broken for easier reading)

```

R CMD INSTALL \
  --configure-args='--with-odbc-include=/opt/local/include \
  --with-odbc-lib=/opt/local/lib --with-odbc-manager=iodbc' \
  RODBC

```

or by setting the environment variables `ODBC_INCLUDE` and `ODBC_LIBS`.

### 1.2.3 Using F95 code

R assumes that source files with extension ‘.f’ are FORTRAN 77, and passes them to the compiler specified by ‘F77’. On most but not all platforms that compiler will accept Fortran 90/95 code: some platforms have a separate Fortran 90/95 compiler and a few (by now quite rare<sup>21</sup>) platforms have no Fortran 90/95 support.

This means that portable packages need to be written in correct FORTRAN 77, which will also be valid Fortran 95. See <http://developer.r-project.org/Portability.html> for reference resources. In particular, *free source form* F95 code is not portable.

On some systems an alternative F95 compiler is available: from the `gcc` family this might be `gfortran` or `g95`. Configuring R will try to find a compiler which (from its name) appears to be a Fortran 90/95 compiler, and set it in macro ‘FC’. Note that it does not check that such a compiler is fully (or even partially) compliant with Fortran 90/95. Packages making use of Fortran 90/95 features should use file extension ‘.f90’ or ‘.f95’ for the source files: the variable `PKG_FCFLAGS` specifies any special flags to be used. There is no guarantee that compiled Fortran 90/95 code can be mixed with any other type of compiled code, nor that a build of R will have support for such packages.

Some (but not) all compilers specified by the ‘FC’ macro will accept Fortran 2003 or 2008 code. For platforms using `gfortran`, you may need to include ‘-std=f2003’ or ‘-std=f2008’ in `PKG_FCFLAGS`: the default is ‘GNU Fortran’, Fortran 95 with non-standard extensions.

<sup>21</sup> Cygwin used `g77` up to 2011, and some pre-built versions of R for Unix OSes still do.



The Solaris `f95` compiler ‘accepts some Fortran 2003 features’. Such code should still use file extension ‘`.f90`’ or ‘`.f95`’.

## 1.3 Checking and building packages

Before using these tools, please check that your package can be installed and loaded. R CMD `check` will *inter alia* do this, but you may get more detailed error messages doing the checks directly.

**Note:** R CMD `check` and R CMD `build` run R with ‘`--vanilla`’, so none of the user’s startup files are read. If you need `R_LIBS` set (to find packages in a non-standard library) you can set it in the environment: also you can use files<sup>22</sup> ‘`~/R/check.Renviron`’ and ‘`~/R/build.Renviron`’ to set environment variables when using these utilities.

**Note to Windows users:** R CMD `build` may require you to have installed the Windows toolset (see the “R Installation and Administration” manual) and have it in your path, and R CMD `check` will make use of it if present. You may need to set `TMPDIR` to point to a suitable writable directory with a path not containing spaces – use forward slashes for the separators. Also, the directory needs to be on a case-honouring file system (some network-mounted file systems are not).

### 1.3.1 Checking packages

Using R CMD `check`, the R package checker, one can test whether *source* R packages work correctly. It can be run on one or more directories, or gzipped package `tar` archives<sup>23</sup> with extension ‘`.tar.gz`’ or ‘`.tgz`’. (Some platforms may allow other forms of compression and extensions ‘`.tar.bz2`’ and ‘`.tar.xz`’.)

This runs a series of checks, including

1. The package is installed. This will warn about missing cross-references and duplicate aliases in help files.
2. The file names are checked to be valid across file systems and supported operating system platforms.
3. The files and directories are checked for sufficient permissions (Unix-alikes only).
4. The files are checked for binary executables, using a suitable version of `file` if available<sup>24</sup>. (There may be rare false positives.)
5. The ‘`DESCRIPTION`’ file is checked for completeness, and some of its entries for correctness. Unless installation tests are skipped, checking is aborted if the package dependencies cannot be resolved at run time. (You may need to set `R_LIBS` if dependent packages are in a separate library tree.) One check is that the package name is not that of a standard package, nor one of the defunct standard packages (‘`ctest`’, ‘`eda`’, ‘`lqs`’, ‘`mle`’, ‘`modreg`’, ‘`mva`’, ‘`nls`’, ‘`stepfun`’ and ‘`ts`’). Another check is that all packages mentioned in `library` or `requires` or from which the ‘`NAMESPACE`’ file imports or are

<sup>22</sup> On systems which use sub-architectures, architecture-specific versions such as ‘`~/R/check.Renviron.i386`’ take precedence.

<sup>23</sup> This may require GNU `tar`: the command used can be set with environment variable `TAR`.

<sup>24</sup> A suitable `file.exe` is part of the Windows toolset.



called *via* `::` or `:::` are listed (in ‘Depends’, ‘Imports’, ‘Suggests’ or ‘Contains’): this is not an exhaustive check of the actual imports.

6. Available index information (in particular, for demos and vignettes) is checked for completeness.
7. The package subdirectories are checked for suitable file names and for not being empty. The checks on file names are controlled by the option ‘`--check-subdirs=value`’. This defaults to ‘`default`’, which runs the checks only if checking a tarball: the default can be overridden by specifying the value as ‘`yes`’ or ‘`no`’. Further, the check on the ‘`src`’ directory is only run if the package does not contain a ‘`configure`’ script (which corresponds to the value ‘`yes-maybe`’) and there is no ‘`src/Makefile`’ or ‘`src/Makefile.in`’.

To allow a ‘`configure`’ script to generate suitable files, files ending in ‘`.in`’ will be allowed in the ‘`R`’ directory.

A warning is given for directory names that look like R package check directories – many packages have been submitted to CRAN containing these.

8. The R files are checked for syntax errors. Bytes which are non-ASCII are reported as warnings, but these should be regarded as errors unless it is known that the package will always be used in the same locale.
9. It is checked that the package can be loaded, first with the usual default packages and then only with package **base** already loaded. It is checked that the namespace this can be loaded in an empty session with only the **base** namespace loaded. (Namespaces and packages can be loaded very early in the session, before the default packages are available, so packages should work then.)
10. The R files are checked for correct calls to `library.dynam`. Package startup functions are checked for correct argument lists and (incorrect) calls to functions which modify the search path or inappropriately generate messages. The R code is checked for possible problems using **codetools**. In addition, it is checked whether S3 methods have all arguments of the corresponding generic, and whether the final argument of replacement functions is called ‘`value`’. All foreign function calls (`.C`, `.Fortran`, `.Call` and `.External` calls) are tested to see if they have a **PACKAGE** argument, and if not, whether the appropriate DLL might be deduced from the namespace of the package. Any other calls are reported. (The check is generous, and users may want to supplement this by examining the output of `tools::checkFF("mypkg", verbose=TRUE)`, especially if the intention were to always use a **PACKAGE** argument)
11. The ‘`Rd`’ files are checked for correct syntax and metadata, including the presence of the mandatory fields (`\name`, `\alias`, `\title` and `\description`). The ‘`Rd`’ name and title are checked for being non-empty, and there is a check for missing cross-references (links).
12. A check is made for missing documentation entries, such as undocumented user-level objects in the package.
13. Documentation for functions, data sets, and S4 classes is checked for consistency with the corresponding code.
14. It is checked whether all function arguments given in `\usage` sections of ‘`Rd`’ files are documented in the corresponding `\arguments` section.

15. The ‘data’ directory is checked for non-ASCII characters and for the use of reasonable levels of compression.
16. C, C++ and FORTRAN source and header files are tested for portable (LF-only) line endings. If there is a ‘Makefile’ or ‘Makefile.in’ or ‘Makevars’ or ‘Makevars.in’ file under the ‘src’ directory, it is checked for portable line endings and the correct use of ‘\$(BLAS\_LIBS)’ and ‘\$(LAPACK\_LIBS)’

Compiled code is checked for symbols corresponding to functions which might terminate R or write to ‘stdout’/‘stderr’ instead of the console. Note that the latter might give false positives in that the symbols might be pulled in with external libraries and could never be called. Windows<sup>25</sup> users should note that the Fortran and C++ runtime libraries are examples of such external libraries.

17. Some checks are made of the contents of the ‘inst/doc’ directory. These always include checking for files that look like leftovers, and if suitable tools (such as `qpdf`) are available, checking that the PDF documentation is of minimal size.
18. The examples provided by the package’s documentation are run. (see [Chapter 2 \[Writing R documentation files\]](#), page 55, for information on using `\examples` to create executable example code.) If there is a file ‘tests/Examples/pkg-Ex.Rout.save’, the output of running the examples is compared to that file.

Of course, released packages should be able to run at least their own examples. Each example is run in a ‘clean’ environment (so earlier examples cannot be assumed to have been run), and with the variables T and F redefined to generate an error unless they are set in the example: See [Section “Logical vectors” in \*An Introduction to R\*](#).

19. If the package sources contain a ‘tests’ directory then the tests specified in that directory are run. (Typically they will consist of a set of ‘.R’ source files and target output files ‘.Rout.save’.) Please note that the comparison will be done in the end user’s locale, so the target output files should be ASCII if at all possible.
20. The code in package vignettes (see [Section 1.4 \[Writing package vignettes\]](#), page 32) is executed, and the vignette PDFs re-made from their sources as a check of completeness of the sources (unless there is a ‘BuildVignettes’ field in the package’s ‘DESCRIPTION’ file with a false value). If there is a target output file ‘.Rout.save’ in the vignette source directory, the output from running the code in that vignette is compared with the target output file and any differences are reported (but not recorded in the log file). (If the vignette sources are in the deprecated location ‘inst/doc’, do mark such target output files to not be installed in ‘.Rinstignore’.)

If there is an error<sup>26</sup> in executing the R code in vignette ‘foo.ext’, a log file ‘foo.ext.log’ is created in the check directory. The vignette PDFs are re-made in a copy of the package sources in the ‘vign\_test’ subdirectory of the check directory, so for further information on errors look in directory ‘pkgname/vign\_test/inst/doc’. (It is only retained if there are errors or if environment variable `_R_CHECK_CLEAN_VIGN_TEST_` is set to a false value.)

<sup>25</sup> on most other platforms such runtime libraries are dynamic, but static libraries are currently used on Windows because the toolchain is not a standard part of the OS.

<sup>26</sup> or if option ‘--use-valgrind’ is used or environment variable `_R_CHECK_ALWAYS_LOG_VIGNETTE_OUTPUT_` is set to a true value or if there are differences from a target output file

21. The PDF version of the package’s manual is created (to check that the ‘Rd’ files can be converted successfully). This needs L<sup>A</sup>T<sub>E</sub>X and suitable fonts and L<sup>A</sup>T<sub>E</sub>X packages to be installed. See [Section “Making the manuals” in \*R Installation and Administration\*](#).

All these tests are run with collation set to the C locale, and for the examples and tests with environment variable `LANGUAGE=en`: this is to minimize differences between platforms.

Use `R CMD check --help` to obtain more information about the usage of the R package checker. A subset of the checking steps can be selected by adding command-line options. It also allows customization by setting environment variables `_R_CHECK_*_:`, as described in [Section “Tools” in \*R Internals\*](#): a set of these customizations similar to those used by CRAN can be selected by the option ‘`--as-cran`’ (which works best if Internet access is available).

You do need to ensure that the package is checked in a suitable locale if it contains non-ASCII characters. Such packages are likely to fail some of the checks in a C locale, and `R CMD check` will warn if it spots the problem. You should be able to check any package in a UTF-8 locale (if one is available). Beware that although a C locale is rarely used at a console, it may be the default if logging in remotely or for batch jobs.

**Multiple sub-architectures:** On systems which support multiple sub-architectures (principally Windows and Mac OS X), `R CMD check` will install and check a package which contains compiled code under all available sub-architectures. (Use option ‘`--force-multiarch`’ to force this for packages without compiled code, which are otherwise only checked under the main sub-architecture.) This will run the loading tests, examples and ‘`tests`’ directory under each installed sub-architecture in turn, and give an error if any fail. Where environment variables (including perhaps `PATH`) need to be set differently for each sub-architecture, these can be set in architecture-specific files such as ‘`R_HOME/etc/i386/Renviron.site`’.

An alternative approach is to use `R CMD check --no-multiarch` to check the primary sub-architecture, and then to use something like `R --arch=x86_64 CMD check --extra-arch` or (Windows) `/path/to/R/bin/x64/Rcmd check --extra-arch` to run for each additional sub-architecture just the checks<sup>27</sup> which differ by sub-architecture.

### 1.3.2 Building package tarballs

Packages may be distributed in source form as “tarballs” (‘`.tar.gz`’ files) or in binary form. The source form can be installed on all platforms with suitable tools and is the usual form for Unix-like systems; the binary form is platform-specific, and is more common distribution form for the Windows and Mac platforms.

Using `R CMD build`, the R package builder, one can build R package tarballs from their sources (for example, for subsequent release).

Prior to actually building the package in the standard gzipped tar file format, a few diagnostic checks and cleanups are performed. In particular, it is tested whether object indices exist and can be assumed to be up-to-date, and C, C++ and FORTRAN source files and relevant make files are tested and converted to LF line-endings if necessary.

---

<sup>27</sup> loading, examples, tests, vignettes

Run-time checks whether the package works correctly should be performed using R CMD `check` prior to invoking the final build procedure.

To exclude files from being put into the package, one can specify a list of exclude patterns in file `‘.Rbuildignore’` in the top-level source directory. These patterns should be Perl-like regular expressions (see the help for `regexp` in R for the precise details), one per line, to be matched<sup>28</sup> against the file names<sup>29</sup> relative to the top-level package source directory. In addition, directories from source control systems<sup>30</sup>, directories with names ending `‘.Rcheck’` or `‘Old’` or `‘old’` and files `‘GNUMakefile’`, `‘Read-and-delete-me’` or with base names starting with `‘.#’`, or starting and ending with `‘#’`, or ending in `‘~’`, `‘.bak’` or `‘.swp’`, are excluded by default. In addition, those files in the `‘R’`, `‘demo’` and `‘man’` directories which are flagged by R CMD `check` as having invalid names will be excluded.

Use R CMD `build --help` to obtain more information about the usage of the R package builder.

Unless R CMD `build` is invoked with the `‘--no-vignettes’` option<sup>31</sup>, it will attempt to rebuild the vignettes (see [Section 1.4 \[Writing package vignettes\], page 32](#)) in the package. To do so it installs the current package into a temporary library tree, but any dependent packages need to be installed in an available library tree (see the Note: at the top of this section).

Similarly, if the `‘.Rd’` documentation files contain any `\Sexpr` macros (see [Section 2.12 \[Dynamic pages\], page 70](#)), the package will be temporarily installed to execute them. Post-execution binary copies of those pages containing build-time macros will be saved in `‘build/partial.rdb’`. If there are any install-time or render-time macros, a `‘.pdf’` version of the package manual will be built and installed in the `‘build/’` subdirectory. (This allows CRAN or other repositories to display the manual even if they are unable to install the package.) This can be suppressed by the option `‘--no-manual’` or if package’s description contains `‘BuildManual: no’` or similar.

One of the checks that R CMD `build` runs is for empty source directories. These are in most (but not all) cases unintentional, if they are intentional use the option `‘--keep-empty-dirs’` (or set the environment variable `_R_BUILD_KEEP_EMPTY_DIRS_` to `‘TRUE’`, or have a `‘BuildKeepEmpty’` field with a true value in the `‘DESCRIPTION’` file).

The `‘--resave-data’` option allows saved images (`‘.rda’` and `‘.RData’` files) in the `‘data’` directory to be optimized for size. It will also compress tabular files and convert `‘.R’` files to saved images. It can take values `no`, `gzip` (the default if this option is not supplied, which can be changed by setting the environment variable `_R_BUILD_RESERVE_DATA_`) and `best` (equivalent to giving it without a value), which chooses the most effective compression. Using `best` adds a dependence on R (`>= 2.10`) to the `‘DESCRIPTION’` file if `bzip2` or `xz` compression is selected for any of the files. If this is thought undesirable, `‘--resave-data=gzip’` (which is the default if that option is not supplied) will do what compression it can with `gzip`. A package can control how its data is resaved by supplying a `‘BuildResaveData’` field (with one of the values given earlier in this paragraph) in its `‘DESCRIPTION’` file.

<sup>28</sup> case-insensitively on Windows.

<sup>29</sup> including directory names as from R 2.13.0: earlier versions accepted the names of non-empty directories.

<sup>30</sup> called `‘CVS’` or `‘.svn’` or `‘.arch-ids’` or `‘.bzip’` or `‘.git’` or `‘.hg’`.

<sup>31</sup> or the package’s description contains `‘BuildVignettes: no’` or similar.

The ‘`--compact-vignettes`’ option will run `tools::compactPDF` over the PDF files in ‘`inst/doc`’ (and its subdirectories) to losslessly compress them. This is not enabled by default (it can be selected by environment variable `_R_BUILD_COMPACT_VIGNETTES_`) and needs `qpdf` (<http://qpdf.sourceforge.net/>) to be available.

It can be useful to run `R CMD check --check-subdirs=yes` on the built tarball as a final check on the contents.

Note that prior to R 2.13.0, `R CMD build` did some cleaning in the supplied source directory, but this was undocumented and is no longer done.

`R CMD build` requires a suitable `tar` program that can produce a compressed tarball: almost certainly one will have been found when R was configured on a Unix-alike (and the Windows toolset contains one), but if there are problems, set the environment variable `TAR` to the path to a suitable program or to “`internal`” if none is available.

### 1.3.3 Building binary packages

Binary packages are compressed copies of installed versions of packages. They contain compiled shared libraries rather than C, C++ or Fortran source code, and the R functions are included in their installed form. The format and filename are platform-specific; for example, a binary package for Windows is usually supplied as a ‘`.zip`’ file, and for the Mac platform the default binary package file extension is ‘`.tgz`’.

The recommended method of building binary packages is to use

`R CMD INSTALL --build pkg` where ‘`pkg`’ is either the name of a source tarball (in the usual ‘`.tar.gz`’ format) or the location of the directory of the package source to be built.

`R CMD INSTALL --build` operates by first installing the package and then packing the installed binaries into the appropriate binary package file for the particular platform.

By default, `R CMD INSTALL --build` will attempt to install the package into the default library tree for the local installation of R. This has two implications:

- If the installation is successful, it will overwrite any existing installation of the same package.
- The default library tree must have write permission; if not, the package will not install and the binary will not be created.

To prevent changes to the present working installation or to provide an install location with write access, create a suitably located directory with write access and use the `-l` option to build the package in the chosen location. The usage is then

`R CMD INSTALL -l location --build pkg`

where ‘`location`’ is the chosen directory with write access. The package will be installed as a subdirectory of ‘`location`’, and the package binary will be created in the current directory.

Other options for `R CMD INSTALL` can be found using `R CMD INSTALL --help`, and platform-specific details for special cases (e.g. handling Fortran sources on Mac OS X) are discussed in the platform-specific FAQs.

In earlier versions of R, `R CMD build --binary` could build a binary version of a package, but this approach is now deprecated in favour of `R CMD INSTALL --build`.

Finally, at least one web-based service is available for building binary packages from (checked) source code: WinBuilder (see <http://win-builder.r-project.org/>) is able to

build Windows binaries. Note that this is intended for developers on other platforms who do not have access to Windows but wish to provide binaries for the Windows platform.

## 1.4 Writing package vignettes

In addition to the help files in ‘`Rd`’ format, R packages allow the inclusion of documents in arbitrary other formats. The standard location for these is subdirectory ‘`inst/doc`’ of a source package, the contents will be copied to subdirectory ‘`doc`’ when the package is installed. Pointers from package help indices to the installed documents are automatically created. Documents in ‘`inst/doc`’ can be in arbitrary format, however we strongly recommend providing them in PDF format, so users on almost all platforms can easily read them. To ensure that they can be accessed from a browser (as an HTML index is provided), the file names should start with an ASCII letter and be comprised entirely of ASCII letters or digits or hyphen or underscore.

A special case are PDF documents with sources in Sweave format, which we call *package vignettes*. As from R 2.14.0 the preferred location for the Sweave sources is the subdirectory ‘`vignettes`’ of the source packages, but for compatibility with earlier versions of R, vignette sources will be looked for in ‘`inst/doc`’ if ‘`vignettes`’ does not exist.

Vignette sources are normally given the file extension ‘`.Rnw`’ or ‘`.Rtex`’, but for historical reasons extensions<sup>32</sup> ‘`.Snw`’ and ‘`.Stex`’ are also recognized as vignettes. Sweave allows the integration of  $\text{\LaTeX}$  documents: see the *Sweave* help page in R and the *Sweave* vignette in package *utils* for details on the document format. Package vignettes are tested by R CMD *check* by executing all R code chunks they contain (except those with option `eval=FALSE`). The R working directory for all vignette tests in R CMD *check* is a *copy* of the vignette source directory. Make sure all files needed to run the R code in the vignette (data sets, ...) are accessible by either placing them in the ‘`inst/doc`’ hierarchy of the source package or by using calls to `system.file()`. All other files needed to re-make the vignette PDFs (such as  $\text{\LaTeX}$  style files,  $\text{\BibTeX}$  input files and files for any figures not created by running the code in the vignette) must in the vignette source directory.

R CMD *build* will automatically<sup>33</sup> create PDF versions of the vignettes in ‘`inst/doc`’ for distribution with the package sources. By including the PDF version in the package sources it is not necessary that the vignette PDFs can be re-built at install time, i.e., the package author can use private R packages, screen snapshots and  $\text{\LaTeX}$  extensions which are only available on his machine.<sup>34</sup>

By default R CMD *build* will run *Sweave* on all files in Sweave format in ‘`vignettes`’, or if that does not exist, ‘`inst/doc`’ (but not in sub-directories). If no ‘`Makefile`’ is found in directory ‘`inst/doc`’, then `tools::texi2dvi(pdf = TRUE)` is run on all processed vignettes. Whenever a ‘`Makefile`’ is found, then R CMD *build* will try to run *make* after the *Sweave* runs. The first target in the ‘`Makefile`’ should take care of both creation of PDF files and cleaning up afterwards (including after *Sweave*), i.e., delete all files that shall not appear in the final package archive. Note that if the *make* step runs R it needs to be careful to respect

<sup>32</sup> and to avoid problems with case-insensitive file systems, lower-case versions of all these extensions.

<sup>33</sup> unless inhibited by using ‘`BuildVignettes: no`’ in the ‘`DESCRIPTION`’ file.

<sup>34</sup> provided the conditions of the package’s licence are met: many would see these as incompatible with an Open Source licence.



the environment values of `R_LIBS` and `R_HOME`<sup>35</sup>. Finally, if there is a ‘`Makefile`’ and it has a ‘`clean:`’ target, `make clean` is run.

All the usual *caveats* about including a ‘`Makefile`’ apply. It must be portable (no GNU extensions) and must work correctly with a parallel `make`: too many authors have written things like

```
## BAD EXAMPLE
all: pdf clean

pdf: ABC-intro.pdf ABC-details.pdf

%.pdf: %.tex
      texi2dvi --pdf $*

clean:
      rm *.tex ABC-details-*.pdf
```

which will start removing the source files whilst `pdflatex` is working.

Note that it is pointless (and potentially misleading since the files might be outdated) to include in ‘`inst/doc`’ R code files which would be generated from vignettes, as these will be re-generated when the package is installed (unless the vignette does not generate any R code, in which case it is also pointless/misleading).

Metadata lines can be placed in the source file, preferably in LaTeX comments in the preamble. One such is a `\VignetteIndexEntry` of the form

```
%\VignetteIndexEntry{Using Animal}
```

Others you may see are `\VignettePackage` (currently ignored), `\VignetteDepends` and `\VignetteKeyword` (which replaced `\VignetteKeywords`). These are processed at package installation time to create the saved data frame ‘`Meta/vignette.rds`’, but only the `\VignetteIndexEntry` and `\VignetteKeyword` statements are currently used.

At install time an HTML index for all vignettes in the package is automatically created from the `\VignetteIndexEntry` statements unless a file ‘`index.html`’ exists in directory ‘`inst/doc`’. This index is linked from the HTML help index for the package. If you do supply a ‘`inst/doc/index.html`’ file it should contain relative links only to files under the installed ‘`doc`’ directory, or perhaps (not really an index) to HTML help files or to the ‘`DESCRIPTION`’ file.

Sweave/Stangle allows the document to specify the `split=TRUE` option to create a single R file for each code chunk: this will not work for vignettes where it is assumed that each vignette source generates a single file with the vignette extension replaced by ‘`.R`’.

Do watch that PDFs are not too large – one in a CRAN package was 72MB! This is usually caused by the inclusion of overly detailed figures, which will not render well in PDF viewers. Sometimes it is much better to generate fairly high resolution bitmap (PNG, JPEG) figures and include those in the PDF document.

When R CMD `build` builds the vignette PDFs, it copies these and the vignette sources from directory ‘`vignettes`’ to ‘`inst/doc`’. To install any other files from the ‘`vignettes`’

<sup>35</sup> As from R 2.13.0, `R_HOME/bin` is prepended to the `PATH` so that references to `R` or `Rscript` in the ‘`Makefile`’ do make use of the currently running version of R.

directory, include a file ‘`vignettes/.install_extras`’ which specifies these as Perl-like regular expressions on one or more lines. (See the description of the ‘`.Rinstignore`’ file for full details.)

### 1.4.1 Encodings and vignettes

Vignette PDFs will in general include descriptive text, R input, R output and figures,  $\LaTeX$  include files and bibliographic references. As any of these may contain non-ASCII characters, the handling of encodings can become very complicated.

The vignette source file should be written in ASCII or contain a declaration of the encoding (see below). This applies even to comments within the source file, since `Sweave()` processes comments to look for options and metadata lines. When `Sweave()` or `Stangle()` is called on the vignette source, it will be converted<sup>36</sup> to the encoding of the current R session.

`Stangle()` will produce an R code file in the current locale’s encoding: for a non-ASCII vignette what that is recorded in a comment at the top of the file.

`Sweave()` will produce a ‘`.tex`’ file in the current locale’s encoding. That needs to be declared to  $\LaTeX$  via a line like

```
\usepackage[utf8]{inputenc}
```

R CMD `check` will warn about any non-ASCII vignettes it finds which do not have such a declaration. The problem is that this cannot be known in advance, so vignette PDFs may only be re-createable on the author’s own machine. R CMD `check` will report on any non-ASCII vignettes it finds which do not have such a declaration. (It is also possible to use the more recent ‘`inputenx`’  $\LaTeX$  package.)

`Sweave()` will also parse and evaluate the R code in each chunk. The R output will also be in the current locale, and should be covered by the ‘`inputenc`’ declaration. One thing people often forget is that the R output may not be ASCII even for ASCII R sources, for many possible reasons. One common one is the use of ‘fancy’ quotes: see the R help on `sQuote`: note carefully that it is not portable to declare UTF-8 or CP1252 to cover such quotes, as their encoding will depend on the locale used to run `Sweave()`: this can be circumvented by setting `options(useFancyQuotes="UTF-8")` in the vignette.

The final issue is the encoding of figures – this applies only to PDF figures and not PNG etc. The PDF figures will contain declarations for their encoding, but the Sweave option `pdf.encoding` may need to be set appropriately: see the help for the `pdf()` graphics device.

As a real example of the complexities, consider the **fortunes** package version ‘1.4-0’. That package did not have a declared encoding, and its vignette was in ASCII. However, the data it displays are read from a UTF-8 CSV file and will be assumed to be in the current encoding, so ‘`fortunes.tex`’ will be in UTF-8 in any locale. Had `read.table` been told the data were UTF-8, ‘`fortunes.tex`’ would have been in the locale’s encoding.

## 1.5 Submitting a package to CRAN

CRAN is a network of WWW sites holding the R distributions and contributed code, especially R packages. Users of R are encouraged to join in the collaborative project and to submit their own packages to CRAN.

<sup>36</sup> provided the encoding is known: currently if it is not, it is guessed to be Latin-1.



Before submitting a package *mypkg*, do run the following steps to test it is complete and will install properly. (Run from the directory containing ‘*mypkg*’ as a subdirectory.)

1. Run R CMD `build` to make the release ‘*.tar.gz*’ file.
2. Run R CMD `check --as-cran` on the ‘*.tar.gz*’ file to check that the package will install and will run its examples, and that the documentation is complete and can be processed. If the package contains code that needs to be compiled, try to enable a reasonable amount of diagnostic messaging (“warnings”) when compiling, such as e.g. ‘`-Wall -pedantic`’ for tools from GCC, the GNU Compiler Collection. If R was not configured accordingly, one can achieve this *via* personal ‘*Makevars*’ files. See [Section “Customizing package compilation” in \*R Installation and Administration\*](#),

Note that it is particularly important to use ‘`-Wall -pedantic`’ with C++ code: the GNU C++ compiler has many extensions which are not supported by other compilers, and this will report *some* of them (such as the misuse of variable-length arrays). If possible, check C++ code on a standards-conformant compiler.

Although there is now a 2011 version of the C++ standard, it is not yet implemented (nor is it likely to be widely available for some years) and portable C++ code needs to follow the 1998 standard (and not use features from C99).

Similarly, the 2011 C standard is unlikely to be widely implemented for several years.

3. Study the output from running your examples, in file ‘*mypkg.Rcheck/mypkg-Ex.Rout*’. Often warnings there indicate actual errors, and warnings about your mistakes (which the R developers are warning you that they are working around for you) will just annoy or confuse your users.

If your package has tests or vignettes, study their output too.

4. Look for any problems with help file conversions. For example, you should
  - Read through the PDF manual that was produced by R CMD `check` at ‘*mypkg.Rcheck/mypkg-manual.pdf*’, or produce another copy by R CMD `Rd2pdf mypkg`.
  - Look at the rendering of your help pages in text from within R.

Many aspects of help rendering changed in R 2.10.0, and in particular the interpretation of comment lines (which are rendered as blank lines, so do not put comment lines in the middle of a paragraph of text).

5. Ensure that the package sources are not unnecessarily large. In particular, R CMD `check` will report<sup>37</sup> on installed packages of more than 5Mb, detailing directories of more than 1Mb. It warns about inefficient compression of data: R CMD `build --resave-data` will compact data as best it can.

Watch out for unnecessary files in ‘*inst/doc*’: R CMD `check` will note files of types that probably should be installed, but it cannot distinguish PDF figures from PDF documents. If files need to be in ‘*inst/doc*’ but not installed, use a ‘*.Rinstignore*’ file.

The CRAN policy is that ‘*doc*’ directories should not exceed 5Mb, and where ‘*data*’ directories need to be more than 5–10Mb, consideration should be given to a separate

<sup>37</sup> provided a POSIX-compliant `du` program is found on the system: it is possible that some other `du` programs will incorrectly report double the actual size. This can be disabled by setting `_R_CHECK_PKG_SIZE_` to a false value.

package containing just the data. (Similarly for external data directories, large ‘`jar`’ files and other libraries that need to be installed.)

See below for ways to reduce the size of PDF files such as vignettes.

6. Ensure that checking the package takes no more time than is needed (the CRAN check farm is a resource shared between several thousand packages).

See below for ways to find out where your package checks are taking significant time.

Please ensure that you can run through the complete procedure with only warnings that you understand and have reasons not to eliminate. In principle, packages must pass `R CMD check` without warnings or significant notes to be admitted to the main CRAN package area. If there are warnings you cannot eliminate (for example because you believe them to be spurious) send an explanatory note as part of your covering email.

Also read the CRAN policies linked from <http://cran.r-project.org/web/packages/> and note that by submitting a package you are confirming that your package complies with them.

When all the testing is done, upload the ‘`.tar.gz`’ file, using ‘`anonymous`’ as log-in name and your e-mail address as password, to <ftp://CRAN.R-project.org/incoming/> (note: use ‘`ftp`’<sup>38</sup> and not ‘`sftp`’ to connect to this server, and passive ‘`ftp`’ is more often successful) and send a message to [CRAN@R-project.org](mailto:CRAN@R-project.org) about it (with your package name and version in the subject line, and please do not submit a package by email). For a new submission, please note in the message that you have read and agreed to the CRAN policies.

The CRAN maintainers will run these tests before putting a submission online. (They will use the latest development version of R, so if at all possible so should you.)

Please note that submissions without an accompanying email to [CRAN@R-project.org](mailto:CRAN@R-project.org) will not be processed, and that emails should not be sent personally to members of the CRAN team.

Note also that for running  $\text{\LaTeX}$ , the Debian GNU/Linux CRAN check systems use the Debian TeXLive<sup>39</sup> distribution (<http://packages.debian.org/en/sid/texlive>); the Fedora and Solaris check systems use current TeXLive; the Windows CRAN builder uses a reasonably recent version of MikTeX (including all packages available directly for MikTeX); the Mac OS X builders use a current full version of MacTeX, which includes all of the current TeXLive. Developers wanting to have their vignettes use  $\text{\TeX}$  packages or style files not (yet) included in these distributions should add<sup>40</sup> the style files to the ‘`vignettes`’ (or for the legacy layout, ‘`inst/doc`’) subdirectory of their package.

### 1.5.1 PDF size

There are a several tools available to reduce the size of PDF files, including Adobe Acrobat (not Reader), Apple’s Preview<sup>41</sup>, `qpdf` (<http://qpdf.sourceforge.net/>), and Ghostscript (which converts PDF to PDF by

<sup>38</sup> for Windows users the simplest way may be to open that URL in Internet Explorer and (depending on the version) follow the instructions to view it as a folder, then copy the submission to the folder.

<sup>39</sup> currently the long obsolete TeXLive 2009.

<sup>40</sup> if their license allows: this often requires also including the corresponding ‘`.dtx`’ file.

<sup>41</sup> Select ‘Save as’, and select ‘Reduce file size’ from the ‘Quartz filter’ menu’: this can be accessed in other ways, for example by Automator.

```
ps2pdf options -dAutoRotatePages=/None in.pdf out.pdf
```

and suitable options might be

```
-dPDFSETTINGS=/ebook
-dPDFSETTINGS=/screen
```

; see <http://www.ghostscript.com/doc/9.04/Ps2pdf.htm> for more such and consider all the options for image downsampling) as well as numerous commercial and shareware Windows programs. Note that these do not all try the same size-reduction strategies, and Acrobat and `ps2pdf` can sometimes do much better at reducing the size of embedded bitmap images, and `ps2pdf` does not use PDF object compression (see below).

Since `qpdf` is fairly readily available (e.g. it has binaries for Windows and packages in Debian/Ubuntu, and is installed as part of the CRAN Mac OS X distribution of R), there is an option ‘`--compact-vignettes`’ to R CMD build to run `qpdf` over PDF files under ‘`inst/doc`’ and replace them if at least 10Kb and 10% is saved. The full path to the `qpdf` command can be supplied as environment variable `R_QPDF` (and is on the CRAN binary of R for Mac OS X). This option can take values ‘`qpdf`’ (the default) as well as ‘`gs`’ or ‘`both`’ to try harder to reduce the size. These should definitely be tried before submission to CRAN for packages with more than 250Kb of PDF files: as ‘`gs`’ may make lossy changes such as downsampling bitmap images, do examine the results and if necessary use `ps2pdf` or `tools::compactPDF` directly.

Most of the large PDFs we have encountered have been large because of the inclusion of figures, for example complex figures from R (where ‘`.png`’ versions may be more appropriate, and PDF compression was not used by `pdf()` prior to R 2.14.0, so it may help to re-generate them) and screendumps. However, some have been unnecessarily large due to `pdftex` settings. The modern default is to use both PDF compression and PDF object compression (which needs PDF version 1.5 from 2003): this is the default in most  $\text{\TeX}$  distributions but not  $\text{\LaTeX}$ . It can be overridden by code in the preamble of an Sweave or  $\text{\LaTeX}$  file: see how this is done for the R reference manual at <https://svn.r-project.org/R/trunk/doc/manual/refman.top>.

## 1.5.2 Package timing

There are several ways to find out where time is being spent in the check process. Start by setting the environment variable `_R_CHECK_TIMINGS_` to ‘0’. This will report the total CPU times (not Windows) and elapsed times for installation and running examples, tests and vignettes, under each sub-architecture if appropriate. For tests and vignettes, it reports the time for each as well as the total.

Setting `_R_CHECK_TIMINGS_` to a non-zero value sets a threshold (in seconds elapsed time) for reporting timings.

If you need to look in more detail at the timings for examples, use option ‘`--timings`’ to R CMD check. This generates a file called ‘`mypkg.Rcheck/mypkg-Ex.timings`’ containing timings for each help files (as given by `system.time()`). It is a tab-delimited file which can be read into R for further analysis.

Timings for the tests and vignette runs are given at the bottom of the corresponding log file: note that log files for successful vignette runs are only retained if `_R_CHECK_ALWAYS_LOG_VIGNETTE_OUTPUT_` is set to a true value.

### 1.5.3 Windows external software

Note that CRAN does not accept submissions of precompiled binaries due to security concerns, and does not allow binary executables in source packages. Maintainers who need additional software for the Windows binaries of their packages on CRAN have three options

1. To arrange for installation of the package to download the additional software from a URL, as e.g. package **Cairo** does.
2. To negotiate with Uwe Ligges to host the additional components on WinBuilder, and write a ‘`configure.win`’ file to install them. There are used to be many examples, e.g. package **rgdal** (however nowadays CRAN prefers to use a uniform cross-compilation approach for software such as GDAL).
3. To negotiate with Brian Ripley to host the package on CRAN extras, as was done for package **BRugs** versions 0.5-x.

Be aware that in all cases license requirements will need to be met so you may need to supply the sources for the additional components (and will if your package has a GPL-like license).

Also be aware that there are both 32- and 64-bit builds of R for Windows with a combined distribution of binary packages, so the CRAN team will be unwilling to support a package that works under just one of the architectures.

## 1.6 Package namespaces

R has a namespace management system for code in packages. This system allows the package writer to specify which variables in the package should be *exported* to make them available to package users, and which variables should be *imported* from other packages.

The mechanism for specifying a namespace for a package is to place a ‘`NAMESPACE`’ file in the top level package directory. This file contains *namespace directives* describing the imports and exports of the namespace. Additional directives register any shared objects to be loaded and any S3-style methods that are provided. Note that although the file looks like R code (and often has R-style comments) it is not processed as R code. Only very simple conditional processing of `if` statements is implemented.

Packages are loaded and attached to the search path by calling `library` or `require`. Only the exported variables are placed in the attached frame. Loading a package that imports variables from other packages will cause these other packages to be loaded as well (unless they have already been loaded), but they will *not* be placed on the search path by these implicit loads.

Namespaces are *sealed* once they are loaded. Sealing means that imports and exports cannot be changed and that internal variable bindings cannot be changed. Sealing allows a simpler implementation strategy for the namespace mechanism. Sealing also allows code analysis and compilation tools to accurately identify the definition corresponding to a global variable reference in a function body.

The namespace controls the search strategy for variables used by functions in the package. If not found locally, R searches the package namespace first, then the imports, then the base namespace and then the normal search path.

If a ‘`NAMESPACE`’ file is not present, then one is generated automatically when the package is built or installed, all objects are exported, and all packages listed in the Imports or

Depends fields in the ‘DESCRIPTION’ file are imported. This is only intended as a temporary measure whilst packages are converted to have a ‘NAMESPACE’ file and will be removed in due course. A hand-crafted ‘NAMESPACE’ should be added to any existing package which does not have one.

Prior to version 2.14.0, namespaces were optional in packages. In such packages searches for non-local variables started with the search path, so a package’s own functions could be masked by those of a package appearing earlier.

As from R 2.14.0 all packages have a namespace, and a default ‘NAMESPACE’ file is generated on installation if there is not one in the sources. However, not all versions of R will read the ‘NAMESPACE’ file if the package contains not R code.

### 1.6.1 Specifying imports and exports

Exports are specified using the `export` directive in the ‘NAMESPACE’ file. A directive of the form

```
export(f, g)
```

specifies that the variables `f` and `g` are to be exported. (Note that variable names may be quoted, and reserved words and non-standard names such as `[<-.fractions` must be.)

For packages with many variables to export it may be more convenient to specify the names to export with a regular expression using `exportPattern`. The directive

```
exportPattern("[^\\\\.]")
```

exports all variables that do not start with a period. However, such broad patterns are not recommended for production code: it is better to list all exports or use narrowly-defined groups. (As from R 2.13.0 this pattern applies to S4 classes, but did not in earlier versions of R.) Beware of patterns which include names starting with a period: some of these are internal-only variables and should never be exported, e.g. ‘`__S3MethodsTable__`’. (Such objects are excluded from pattern matches in recent versions of R, so such patterns are safer for packages only to be used with R 2.14.0 or later.)

Packages implicitly import the base namespace. Variables exported from other packages with namespaces need to be imported explicitly using the directives `import` and `importFrom`. The `import` directive imports all exported variables from the specified package(s). Thus the directives

```
import(foo, bar)
```

specifies that all exported variables in the packages `foo` and `bar` are to be imported. If only some of the exported variables from a package are needed, then they can be imported using `importFrom`. The directive

```
importFrom(foo, f, g)
```

specifies that the exported variables `f` and `g` of the package `foo` are to be imported.

It is possible to export variables from a namespace that it has imported from other namespaces.

If a package only needs a few objects from another package it can use a fully qualified variable reference in the code instead of a formal import. A fully qualified reference to the function `f` in package `foo` is of the form `foo::f`. This is slightly less efficient than a formal import and also loses the advantage of recording all dependencies in the ‘NAMESPACE’ file, so this approach is usually not recommended. Evaluating `foo::f` will cause package `foo`

to be loaded, but not attached, if it was not loaded already—this can be an advantage in delaying the loading of a rarely used package.

Using `foo:::f` instead of `foo::f` allows access to unexported objects. This is generally not recommended, as the semantics of unexported objects may be changed by the package author in routine maintenance.

## 1.6.2 Registering S3 methods

The standard method for S3-style `UseMethod` dispatching might fail to locate methods defined in a package that is imported but not attached to the search path. To ensure that these methods are available the packages defining the methods should ensure that the generics are imported and register the methods using `S3method` directives. If a package defines a function `print.foo` intended to be used as a `print` method for class `foo`, then the directive

```
S3method(print, foo)
```

ensures that the method is registered and available for `UseMethod` dispatch, and the function `print.foo` does not need to be exported. Since the generic `print` is defined in `base` it does not need to be imported explicitly.

(Note that function and class names may be quoted, and reserved words and non-standard names such as `[<-` and `function` must be.)

## 1.6.3 Load hooks

There are a number of hooks called as packages are loaded, attached, detached, and unloaded. See `help(".onLoad")` for more details.

Since loading and attaching are distinct operations, separate hooks are provided for each. These hook functions are called `.onLoad` and `.onAttach`. They both take arguments<sup>42</sup> `libname` and `pkgname`; they should be defined in the namespace but not exported.

Packages use the `.Last.lib` function (provided it is exported from the namespace) when `detach` is called on the package. It is called with a single argument, the full path to the installed package. There is also a hook `.onUnload` which is called when the namespace is unloaded (*via* a call to `unloadNamespace`, perhaps called by `detach(unload=TRUE)`) with argument the full path to the installed package's directory. `.onUnload` should be defined in the name space and not exported, but `.Last.lib` does need to be exported.

Packages are not likely to need `.onAttach` (except perhaps for a start-up banner); code to set options and load shared objects should be placed in a `.onLoad` function, or use made of the `useDynLib` directive described next.

There can be one or more `useDynLib` directives which allows shared objects that need to be loaded to be specified in the 'NAMESPACE' file.<sup>43</sup> The directive

```
useDynLib(foo)
```

registers the shared object `foo`<sup>44</sup> for loading with `library.dynam`. Loading of registered object(s) occurs after the package code has been loaded and before running the load hook

<sup>42</sup> they will be called with two unnamed arguments, in that order.

<sup>43</sup> NB: this will only be read in all versions of R if the package contains R code in a 'R' directory.

<sup>44</sup> Note that this is the basename of the shared object, and the appropriate extension ('.so' or '.dll') will be added.



function. Packages that would only need a load hook function to load a shared object can use the `useDynLib` directive instead.

User-level hooks are also available: see the help on function `setHook`.

The `useDynLib` directive also accepts the names of the native routines that are to be used in R *via* the `.C`, `.Call`, `.Fortran` and `.External` interface functions. These are given as additional arguments to the directive, for example,

```
useDynLib(foo, myRoutine, myOtherRoutine)
```

By specifying these names in the `useDynLib` directive, the native symbols are resolved when the package is loaded and R variables identifying these symbols are added to the package's namespace with these names. These can be used in the `.C`, `.Call`, `.Fortran` and `.External` calls in place of the name of the routine and the `PACKAGE` argument. For instance, we can call the routine `myRoutine` from R with the code

```
.Call(myRoutine, x, y)
```

rather than

```
.Call("myRoutine", x, y, PACKAGE = "foo")
```

There are at least two benefits to this approach. Firstly, the symbol lookup is done just once for each symbol rather than each time the routine is invoked. Secondly, this removes any ambiguity in resolving symbols that might be present in several compiled DLLs.

In some circumstances, there will already be an R variable in the package with the same name as a native symbol. For example, we may have an R function in the package named `myRoutine`. In this case, it is necessary to map the native symbol to a different R variable name. This can be done in the `useDynLib` directive by using named arguments. For instance, to map the native symbol name `myRoutine` to the R variable `myRoutine_sym`, we would use

```
useDynLib(foo, myRoutine_sym = myRoutine, myOtherRoutine)
```

We could then call that routine from R using the command

```
.Call(myRoutine_sym, x, y)
```

Symbols without explicit names are assigned to the R variable with that name.

In some cases, it may be preferable not to create R variables in the package's namespace that identify the native routines. It may be too costly to compute these for many routines when the package is loaded if many of these routines are not likely to be used. In this case, one can still perform the symbol resolution correctly using the DLL, but do this each time the routine is called. Given a reference to the DLL as an R variable, say `dll`, we can call the routine `myRoutine` using the expression

```
.Call(dll$myRoutine, x, y)
```

The `$` operator resolves the routine with the given name in the DLL using a call to `getNativeSymbol`. This is the same computation as above where we resolve the symbol when the package is loaded. The only difference is that this is done each time in the case of `dll$myRoutine`.

In order to use this dynamic approach (e.g., `dll$myRoutine`), one needs the reference to the DLL as an R variable in the package. The DLL can be assigned to a variable by using the `variable = dllName` format used above for mapping symbols to R variables. For example, if we wanted to assign the DLL reference for the DLL `foo` in the example above to the variable `myDLL`, we would use the following directive in the 'NAMESPACE' file:



```
myDLL = useDynLib(foo, myRoutine_sym = myRoutine, myOtherRoutine)
```

Then, the R variable `myDLL` is in the package's namespace and available for calls such as `myDLL$dynRoutine` to access routines that are not explicitly resolved at load time.

If the package has registration information (see [Section 5.4 \[Registering native routines\]](#), [page 97](#)), then we can use that directly rather than specifying the list of symbols again in the `useDynLib` directive in the 'NAMESPACE' file. Each routine in the registration information is specified by giving a name by which the routine is to be specified along with the address of the routine and any information about the number and type of the parameters. Using the `.registration` argument of `useDynLib`, we can instruct the namespace mechanism to create R variables for these symbols. For example, suppose we have the following registration information for a DLL named `myDLL`:

```
R_CMethodDef cMethods[] = {
  {"foo", (DL_FUNC) &foo, 4, {REALSXP, INTSXP, STRSXP, LGLSXP}},
  {"bar_sym", (DL_FUNC) &bar, 0},
  {NULL, NULL, 0}
};

R_CallMethodDef callMethods[] = {
  {"R_call_sym", (DL_FUNC) &R_call, 4},
  {"R_version_sym", (DL_FUNC) &R_version, 0},
  {NULL, NULL, 0}
};
```

Then, the directive in the 'NAMESPACE' file

```
useDynLib(myDLL, .registration = TRUE)
```

causes the DLL to be loaded and also for the R variables `foo`, `bar_sym`, `R_call_sym` and `R_version_sym` to be defined in the package's namespace.

Note that the names for the R variables are taken from the entry in the registration information and do not need to be the same as the name of the native routine. This allows the creator of the registration information to map the native symbols to non-conflicting variable names in R, e.g. `R_version` to `R_version_sym` for use in an R function such as

```
R_version <- function()
{
  .Call(R_version_sym)
}
```

Using argument `.fixes` allows an automatic prefix to be added to the registered symbols, which can be useful when working with an existing package. For example, package **KernSmooth** has

```
useDynLib(KernSmooth, .registration = TRUE, .fixes = "F_")
```

which makes the R variables corresponding to the FORTRAN symbols `F_bkde` and so on, and so avoid clashes with R code in the name space.

More information about this symbol lookup, along with some approaches for customizing it, is available from <http://www.omegahat.org/examples/RDotCall>.

### 1.6.4 An example

As an example consider two packages named **foo** and **bar**. The R code for package **foo** in file ‘foo.R’ is

```
x <- 1
f <- function(y) c(x,y)
foo <- function(x) .Call("foo", x, PACKAGE="foo")
print.foo <- function(x, ...) cat("<a foo>\n")
```

Some C code defines a C function compiled into DLL **foo** (with an appropriate extension). The ‘NAMESPACE’ file for this package is

```
useDynLib(foo)
export(f, foo)
S3method(print, foo)
```

The second package **bar** has code file ‘bar.R’

```
c <- function(...) sum(...)
g <- function(y) f(c(y, 7))
h <- function(y) y+9
```

and ‘NAMESPACE’ file

```
import(foo)
export(g, h)
```

Calling `library(bar)` loads **bar** and attaches its exports to the search path. Package **foo** is also loaded but not attached to the search path. A call to `g` produces

```
> g(6)
[1] 1 13
```

This is consistent with the definitions of `c` in the two settings: in **bar** the function `c` is defined to be equivalent to `sum`, but in **foo** the variable `c` refers to the standard function `c` in **base**.

### 1.6.5 Summary – converting an existing package

To summarize, converting a pre-2.14.0 package to use a namespace involves several simple steps:

- Identify the public definitions and place them in `export` directives.
- Identify S3-style method definitions and write corresponding `S3method` declarations.
- Identify dependencies and replace any `require` calls by `import` directives (and make appropriate changes in the `Depends` and `Imports` fields of the ‘DESCRIPTION’ file).
- Replace `.First.lib` functions with `.onLoad`/`.onAttach` functions or use a `useDynLib` directive in the ‘NAMESPACE’ file.

The first two of these are done automatically, but a package author can usually improve on R's guesswork.

R CMD build will add a basic 'NAMESPACE' file to a package. If this is edited, do remove the first line (as the comment in the file says).

### 1.6.6 Namespaces with S4 classes and methods

Some additional steps are needed for packages which make use of formal (S4-style) classes and methods (unless these are purely used internally). The package should have `Depends: methods` in its 'DESCRIPTION' file and any classes and methods which are to be exported need to be declared in the 'NAMESPACE' file. For example, the **stats4** package has

```
export(mle)
importFrom("graphics", plot)
importFrom("stats", optim, qchisq)
## For these, we define methods or (AIC, BIC, nobs) an implicit generic:
importFrom("stats", AIC, BIC, coef, confint, logLik, nobs, profile,
           update, vcov)
exportClasses(mle, profile.mle, summary.mle)
## All methods for imported generics:
exportMethods(coef, confint, logLik, plot, profile, summary, show, update, vcov)
## implicit generics which do not have any methods here
export(AIC, BIC, nobs)
```

All S4 classes to be used outside the package need to be listed in an `exportClasses` directive. Alternatively, they can be specified using `exportClassPattern`.<sup>45</sup> in the same style as for `exportPattern`. To export methods for generics from other packages an `exportMethods` directive can be used.

Note that exporting methods on a generic in the namespace will also export the generic, and exporting a generic in the namespace will also export its methods. If the generic function is not local to this package, either because it was imported as a generic function or because the non-generic version has been made generic solely to add S4 methods to it (as for functions such as `plot` in the example above), it can be declared *via* either or both of `export` or `exportMethods`, but the latter is clearer (and is used in the **stats4** example above). In particular, for primitive functions there is no generic function, so `export` would export the primitive, which makes no sense. On the other hand, if the generic is local to this package, it is more natural to export the function itself using `export()`, and this *must* be done if an implicit generic is created without setting any methods for it (as is the case for `AIC` in **stats4**).

A non-local generic function is only exported to ensure that calls to the function will dispatch the methods from this package (and that is not done or required when the methods are for primitive functions). For this reason, you do not need to document such implicitly created generic functions, and `undoc` in package **tools** will not report them.

If a package uses S4 classes and methods exported from another package, but does not import the entire namespace of the other package, it needs to import the classes and methods explicitly, with directives

<sup>45</sup> As from R 2.13.0 this defaults to the same pattern as `exportPattern`: use something like `exportClassPattern("^$")` to override this.

```
importClassesFrom(package, ...)
importMethodsFrom(package, ...)
```

listing the classes and functions with methods respectively. Suppose we had two small packages **A** and **B** with **B** using **A**. Then they could have `NAMESPACE` files

```
export(f1, ng1)
exportMethods("[" )
exportClasses(c1)
```

and

```
importFrom(A, ng1)
importClassesFrom(A, c1)
importMethodsFrom(A, f1)
export(f4, f5)
exportMethods(f6, "[")
exportClasses(c1, c2)
```

respectively.

Note that `importMethodsFrom` will also import any generics defined in the namespace on those methods.

It is important if you export S4 methods that the corresponding generics are available: the requirements on this are stricter as from R 2.15.0. You may for example need to import `plot` from **graphics** to make visible a function to be converted into its implicit generic. But it is better practice to make use of the generics exported by **stats4** as this enables multiple packages to unambiguously set methods on those generics.

## 1.7 Writing portable packages

Portable packages should have simple file names: use only alphanumeric ASCII characters and `.`, and avoid those names not allowed under Windows which are mentioned above.

R CMD `check` provides a basic set of checks, but often further problems emerge when people try to install and use packages submitted to CRAN – many of these involve compiled code. Here are some further checks that you can do to make your package more portable.

- If your package has a `'configure'` script, provide a `'configure.win'` script to be used on Windows. The CRAN binary packages for Windows are built automatically, and if your package does not build without intervention it is unlikely to be easily available to a high proportion of R users.
- If your package has a `'Makevars'` or `'Makefile'` file, make sure that you use only portable make features. Such files should be LF-terminated (including the final line of the file) and not make use of GNU extensions. Commonly misused GNU extensions are conditional inclusions (`ifeq` and the like), `${shell ...}` and `${wildcard ...}`, and the use of `+=` and `:=`. Also, the use of `$<` other than in implicit rules is a GNU extension. Unfortunately makefiles which use GNU extensions often run on other platforms but do not have the intended results.

The use of `${shell ...}` can be avoided by using backticks, e.g.

```
PKG_CPPFLAGS = 'gsl-config --cflags'
```

which works in all versions of `make` known<sup>46</sup> to be used with R.

If you really must assume GNU `make`, declare it in the 'DESCRIPTION' file by

```
SystemRequirements: GNU make
```

Since the only viable `make` for Windows is GNU `make`, it is permissible to use GNU extensions in files 'Makevars.win' or 'Makefile.win'.

- Make use of the abilities of your compilers to check the standards-conformance of your code. For example, `gcc` can be used with options '`-Wall -pedantic`' to alert you to potential problems. This is particularly important for C++, where `g++ -Wall -pedantic` will alert you to the use of GNU extensions which fail to compile on most other C++ compilers. R assumes a C99 compiler as from version 2.12.0, but if you want your package to be portable to earlier versions you should write in C90. (In practice C99 has been available on most platforms since *ca* 2007 but old versions of `gcc` were still in use for R 2.11.x.)

If you use FORTRAN 77, `ftnchek` (<http://www.dsm.fordham.edu/~ftnchek/>) provides thorough testing of conformance to the standard.

- Do be very careful with passing arguments between R, C and FORTRAN code. In particular, `long` in C will be 32-bit on most R platforms (including those mostly used by the CRAN maintainers), but 64-bit on many modern Unix and Linux platforms. It is rather unlikely that the use of `long` in C code has been thought through: if you need a longer type than `int` you should use a configure test for a C99 type such as `int_fast64_t` (and failing that, `long long`<sup>47</sup>) and typedef your own type to be `long` or `long long`, or use another suitable type (such as `size_t`).

It is not safe to assume that `long` and pointer types are the same size, and they are not on 64-bit Windows. If you need to convert pointers to and from integers use the C99 integer types `intptr_t` and `uintptr_t` (which are defined in the header `<stdint.h>` and are not required to be implemented by the C99 standard).

Note that `integer` in FORTRAN corresponds to `int` in C on all R platforms.

- Under no circumstances should your compiled code ever call `abort` or `exit`: these terminate the user's R process, quite possibly including all his unsaved work. One usage that could call `abort` is the `assert` macro in C or C++ functions, which should never be active in production code. The normal way to ensure that is to define the macro `NDEBUG`, and as from R 2.15.0 R CMD INSTALL does so as part of the compilation flags. If you wish to use `assert` during development. you can include `-UNDEBUG` in `PKG_CPPFLAGS`. Note that your own 'src/Makefile' or makefiles in sub-directories may also need to define `NDEBUG`.

This applies not only to your own code but to any external software you compile in or link to. Such code may contain references to `abort` or `exit` that can never be called, but if any are found in the package's shared object/DLL, they are reported by R CMD `check`.

<sup>46</sup> GNU `make`, BSD `make` as in FreeBSD and `bsdmake` on Darwin, AT&T `make` as implemented on Solaris.

<sup>47</sup> but note that `long long` is not a standard C++ type, and C++ compilers set up for strict checking will reject it.

- Compiled code should not write to ‘`stdout`’ or ‘`stderr`’ and C++ and Fortran I/O should not be used. Where `R CMD check` can detect this in the package’s shared object/DLL it will report it: as with the previous item such calls may come from external software and may never be called. (This seems particularly common on Mac OS X and Windows with static libraries: however on Windows the use of Fortran I/O usually results in the detection of `_assert` and `exit`.)
- Errors in memory allocation and reading/writing outside arrays are very common causes of crashes (e.g., segfaults) on some machines. See [Section 4.3.2 \[Using valgrind\]](#), [page 87](#) for a tool which can be used to look for this.
- Many platforms will allow unsatisfied entry points in compiled code, but will crash the application (here R) if they are ever used. Some (notably Windows) will not. Looking at the output of

```
nm -pg mypkg.so # or other extension such as '.sl'
```

and checking if any of the symbols marked `U` is unexpected is a good way to avoid this.

- Conflicts between symbols in DLLs are handled in very platform-specific ways. Good ways to avoid trouble are to make as many symbols as possible static (check with `nm -pg`), and to use unusual names, as well as ensuring you have used the `PACKAGE` argument that `R CMD check` checks for.
- Do not use (hard or symbolic) file links in your package sources. `R CMD build` packages the tarball with the ‘`-h`’ `tar` flag which is documented to dereference links so this is not usually a problem, but versions 1.24 and later of GNU `tar` dereference some links to hard links which may not be handled correctly by `R CMD INSTALL`.
- If you do not yourself have a Windows system, submit your source package to Win-Builder (<http://win-builder.r-project.org/>) before distribution (including submission to CRAN).

### 1.7.1 Encoding issues

Care is needed if your package contains non-ASCII text, and in particular if it is intended to be used in more than one locale. It is possible to mark the encoding used in the ‘DESCRIPTION’ file and in ‘.Rd’ files, as discussed elsewhere in this manual.

First, consider carefully if you really need non-ASCII text. Many users of R will only be able to view correctly text in their native language group (e.g. Western European, Eastern European, Simplified Chinese) and ASCII. Other characters may not be rendered at all, rendered incorrectly, or cause your R code to give an error. For documentation, marking the encoding and including ASCII transliterations is likely to do a reasonable job. The set of characters which is commonly supported is wider than it used to be around 2000, but non-Latin alphabets (Greek, Russian, Georgian, . . .) are still often problematic and those with double-width characters (Chinese, Japanese, Korean) often need specialist fonts to render correctly.

Several CRAN packages have messages in their R code in French (and a few in German). A better way to tackle this is to use the internationalization facilities discussed elsewhere in this manual.

Function `showNonASCIIfile` in package `tools` can help in finding non-ASCII bytes in files.

From R 2.10.0 there is a portable way to have arbitrary text in character strings (only) in your R code, which is to supply them in Unicode as `\uxxxx` escapes. If there are any characters not in the current encoding the parser will encode the character string as UTF-8 and mark it as such. This applies also to character strings in datasets: they can be prepared using `\uxxxx` escapes or encoded in UTF-8 in a UTF-8 locale, or even converted to UTF-8 via `'iconv()'`. If you do this, make sure you have `'R (>= 2.10)'` (or later) in the `'Depends:'` field of the `'DESCRIPTION'` file.

R sessions running in non-UTF-8 locales will if possible re-encode such strings for display (and this is done by `RGui` on Windows, for example). Suitable fonts will need to be selected or made available<sup>48</sup> both for the console/terminal and graphics devices such as `'X11()'` and `'windows()'`. Using `'postscript'` or `'pdf'` will choose a default 8-bit encoding depending on the language of the UTF-8 locale, and your users would need to be told how to select the `'encoding'` argument.

If you want to run R CMD `check` on a Unix-alike over a package that sets a package encoding in its `'DESCRIPTION'` file you may need to specify a suitable locale *via* environment variable `R_ENCODING_LOCALES`. The default is equivalent to the value

```
"latin1=en_US:latin2=pl_PL:UTF-8=en_US.UTF-8:latin9=fr_FR.iso885915@euro"
```

(which is appropriate for a system based on `glibc`) except that if the current locale is UTF-8 then the package code is translated to UTF-8 for syntax checking.

### 1.7.2 Binary distribution

If you want to distribute a binary version of a package on Windows or Mac OS X, there are further checks you need to do to check it is portable: it is all too easy to depend on external software on your own machine that other users will not have.

For Windows, check what other DLLs your package's DLL depends on ('imports' from in the DLL tools' parlance). A convenient GUI-based tool to do so is 'Dependency Walker' (<http://www.dependencywalker.com/>) for both 32-bit and 64-bit DLLs – note that this will report as missing links to R's own DLLs such as `'R.dll'` and `'Rblas.dll'`. For 32-bit DLLs only, the command-line tool `pedump.exe -i` (in `'Rtools*.exe'`) can be used, and for the brave, the `objdump` tool in the appropriate toolchain will also reveal what DLLs are imported from. If you use a toolchain other than one provided by the R developers or use your own makefiles, watch out in particular for dependencies on the toolchain's runtime DLLs such as `'libgfortran'`, `'libstdc++'` and `'libgcc_s'`.

For Mac OS X, using R CMD `otool -L` on the package's shared objects under `'libs'` will show what they depend on: watch for any dependencies in `'/usr/local/lib'`, notably `'libgfortran.2.dylib'`.

## 1.8 Diagnostic messages

Now that diagnostic messages can be made available for translation, it is important to write them in a consistent style. Using the tools described in the next section to extract all the messages can give a useful overview of your consistency (or lack of it).

Some guidelines follow.

<sup>48</sup> Typically on a Unix-alike this is done by telling `fontconfig` where to find suitable fonts to select glyphs from.



- Messages are sentence fragments, and not viewed in isolation. So it is conventional not to capitalize the first word and not to end with a period (or other punctuation).
- Try not to split up messages into small pieces. In C error messages use a single format string containing all English words in the messages.

In R error messages do not construct a message with `paste` (such messages will not be translated) but *via* multiple arguments to `stop` or `warning`, or *via* `gettextf`.

- Do not use colloquialisms such as “can’t” and “don’t”.
- If possible, make quotation marks part of your message as different languages have different conventions. In R messages this means not using `sQuote` or `dQuote` except where the argument is a variable.

Conventionally single quotation marks are used for quotations such as

`'ord'` must be a positive integer, at most the number of knots

and double quotation marks when referring to an R character string such as

`'format'` must be "normal" or "short" - using "normal"

Since ASCII does not contain directional quotation marks, it is best to use ‘’ and let the translator (including automatic translation) use directional quotations where available. The range of quotation styles is immense: unfortunately we cannot reproduce them in a portable `texinfo` document. But as a taster, some languages use ‘up’ and ‘down’ (comma) quotes rather than left or right quotes, and some use guillemets (and some use what Adobe calls ‘guillemotleft’ to start and others use it to end).

- Occasionally messages need to be singular or plural (and in other languages there may be no such concept or several plural forms – Slovenian has four). So avoid constructions such as was once used in `library`

```
if((length(nopkgs) > 0) && !missing(lib.loc)) {
  if(length(nopkgs) > 1)
    warning("libraries ",
            paste(sQuote(nopkgs), collapse = ", "),
            " contain no packages")
  else
    warning("library ", paste(sQuote(nopkgs)),
            " contains no package")
}
```

and was replaced by

```
if((length(nopkgs) > 0) && !missing(lib.loc)) {
  pkglist <- paste(sQuote(nopkgs), collapse = ", ")
  msg <- sprintf(ngettext(length(nopkgs),
                           "library %s contains no packages",
                           "libraries %s contain no packages"),
                 pkglist)
  warning(msg, domain=NA)
}
```

Note that it is much better to have complete clauses as here, since in another language one might need to say ‘There is no package in library %s’ or ‘There are no packages in libraries %s’.

## 1.9 Internationalization

There are mechanisms to translate the R- and C-level error and warning messages. There are only available if R is compiled with NLS support (which is requested by `configure` option ‘`--enable-nls`’, the default).

The procedures make use of `msgfmt` and `xgettext` which are part of GNU `gettext` and this will need to be installed: Windows users can find pre-compiled binaries at the GNU archive mirrors and packaged with the `poEdit` package (<http://poedit.sourceforge.net/download.php#win32>).

### 1.9.1 C-level messages

The process of enabling translations is

- In a header file that will be included in all the C files containing messages that should be translated, declare

```
#include <R.h> /* to include Rconfig.h */

#ifdef ENABLE_NLS
#include <libintl.h>
#define _(String) dgettext ("pkg", String)
/* replace pkg as appropriate */
#else
#define _(String) (String)
#endif
```

- For each message that should be translated, wrap it in `_(...)`, for example

```
error(_("'ord' must be a positive integer"));
```

If you want to use different messages for singular and plural forms, you need to add

```
#ifndef ENABLE_NLS
#define dngettext(pkg, String, StringP, N) (N > 1 ? StringP: String)
#endif
```

and mark strings by

```
dngettext(("pkg", <singular string>, <plural string>, n)
```

(This is only supported from R 2.10.0, so packages which use it need to depend on R ( $\geq 2.10$ ).)

- In the package’s ‘src’ directory run

```
xgettext --keyword=_ -o pkg.pot *.c
```

The file ‘src/pkg.pot’ is the template file, and conventionally this is shipped as ‘po/pkg.pot’. A translator to another language makes a copy of this file and edits it (see the `gettext` manual) to produce say ‘ll.po’, where *ll* is the code for the language in which the translation is to be used. (This file would be shipped in the ‘po’ directory.) Next run `msgfmt` on ‘ll.po’ to produce ‘ll.mo’, and copy that to ‘inst/po/ll/LC\_MESSAGES/pkg.mo’. Now when the package is loaded after installation it will look for translations of its messages in the ‘po/lang/LC\_MESSAGES/pkg.mo’ file for any language *lang* that matches the user’s preferences (*via* the setting of the `LANGUAGE` environment variable or from the locale settings).

### 1.9.2 R messages

Mechanisms are also available to support the automatic translation of R `stop`, `warning` and `message` messages. They make use of message catalogs in the same way as C-level messages, but using domain `R-pkg` rather than `pkg`. Translation of character strings inside `stop`, `warning` and `message` calls is automatically enabled, as well as other messages enclosed in calls to `gettext` or `gettextf`. (To suppress this, use argument `domain=NA`.)

Tools to prepare the ‘`R-pkg.pot`’ file are provided in package `tools`: `xgettext2pot` will prepare a file from all strings occurring inside `gettext/gettextf`, `stop`, `warning` and `message` calls. Some of these are likely to be spurious and so the file is likely to need manual editing. `xgettext` extracts the actual calls and so is more useful when tidying up error messages.

Translation of messages which might be singular or plural can be very intricate: languages can have up to four different forms. The R function `ngettext` provides an interface to the C function of the same name, and will choose an appropriate singular or plural form for the selected language depending on the value of its first argument `n`. It is safest to use `domain="R-pkg"` explicitly in calls to `ngettext`, and necessary unless they are calls directly from a function in the package.

### 1.9.3 Installing translations

Once the template files have been created, translations can be made. Conventional translations have file extension ‘`.po`’ and are placed in the ‘`po`’ subdirectory of the package with a name that is either ‘`ll.po`’ or ‘`R-ll.po`’ for translations of the C and R messages respectively to language with code ‘`ll`’.

See [Section “Localization of messages” in \*R Installation and Administration\*](#), for details of language codes.

Translations need to be prepared and installed in ‘`inst/po/`’ to be usable once the package is installed. To do this use the appropriate lines of

```
mkdir -p inst/po/ll/LC_MESSAGES
msgfmt -c --statistics -o inst/po/ll/LC_MESSAGES/R-pkg.mo po/R-ll.po
msgfmt -c --statistics -o inst/po/ll/LC_MESSAGES/pkg.mo po/ll.po
```

from the package’s top-level directory. Using ‘`-c`’ does some useful validity checks, and ‘`--statistics`’ notes the coverage.

### 1.9.4 Makefile support

There is some makefile support in the ‘`po`’ directory of the R sources. To use this to create the template files, use

```
mkdir -p pkgdir/po
```

where ‘`pkgdir`’ is the top-level directory of the package sources. If the package has C source files in its ‘`src`’ directory that are marked for translation, use

```
touch pkgdir/po/pkg.pot
```

to create a dummy template file. Then

```
cd R_BUILD_DIR/po
make pkg-update PKG=pkg PKGDIR=pkgdir
```

will create a template file of R messages and update any template of C messages. It will also prepare and install a translation for the ‘en@quot’ pseudo-language, which if selected interprets (single and double) quotes in their directional forms in suitable (e.g. UTF-8) locales.

If translations to new languages are added in the ‘*pkgdir/po*’ directory, running the same **make** command will check and then install the translations.

If the package sources are updated, the same **make** command will update the template files, merge the changes into the translation ‘.po’ files and then installed the updated translations. You will often see that merging marks translations as ‘fuzzy’ and this is reported in the coverage statistics. As fuzzy translations are *not* used, this is an indication that the translation files need human attention.

This support is only for Unix-alikes, and the tools did not work correctly on at least one Mac OS X system.

## 1.10 CITATION files

An installed file named ‘CITATION’ will be used by the `citation()` function. (To be installed, it needed to be in the ‘*inst*’ subdirectory of the package sources.)

The ‘CITATION’ file is parsed as R code (in the package’s declared encoding, or in ASCII if none is declared). If no such file is present, **citation** auto-generates citation information from the package ‘DESCRIPTION’ metadata, and an example of what that would look like as a ‘CITATION’ file can be seen in recommended package **nlme** (see below): recommended packages **boot**, **cluster** and **mgcv** have further examples.

A ‘CITATION’ file will contain calls to function **bibentry** (new style, only works with R 2.12.0 or later), or to the functions **citHeader**, **citEntry** and (optionally) **citFooter** (old style).

Here is that for **nlme**, re-formatted:

```
citHeader("To cite package 'nlme' in publications use:")

year <- sub(".*(2[[:digit:]]{3})-.*", "\\1", meta$Date, perl = TRUE)
vers <- paste("R package version", meta$Version)

citEntry(entry="Manual",
  title = "nlme: Linear and Nonlinear Mixed Effects Models",
  author = personList(as.person("Jose Pinheiro"),
                      as.person("Douglas Bates"),
                      as.person("Saikat DebRoy"),
                      as.person("Deepayan Sarkar"),
                      person("R Development Core Team")),
  year = year,
  note = vers,

  textVersion =
    paste("Jose Pinheiro, Douglas Bates, Saikat DebRoy,",
          "Deepayan Sarkar and the R Development Core Team (",
          year,
```

```
)". nlme: Linear and Nonlinear Mixed Effects Models. ",
vers, ".", sep=""))
```

Note the way that information that may need to be updated is picked up from the ‘DESCRIPTION’ file – it is tempting to hardcode such information, but it normally then gets outdated. See `?bibentry` for further details of the information which can be provided.

The ‘CITATION’ file should itself produce no output when `source-d`.

## 1.11 Package types

The ‘DESCRIPTION’ file has an optional field `Type` which if missing is assumed to be `Package`, the sort of extension discussed so far in this chapter. Currently two other types are recognized, both of which need write permission in the R installation tree.

### 1.11.1 Frontend

This is a rather general mechanism, designed for adding new front-ends such as the former **gnomeGUI** package (see the ‘Archive’ area on CRAN). If a ‘configure’ file is found in the top-level directory of the package it is executed, and then if a `Makefile` is found (often generated by ‘configure’), `make` is called. If R CMD `INSTALL --clean` is used `make clean` is called. No other action is taken.

R CMD `build` can package up this type of extension, but R CMD `check` will check the type and skip it.

### 1.11.2 Translation

Conventionally, a translation package for language *ll* is called **Translation-ll** and has `Type: Translation`. It needs to contain the directories ‘share/locale/ll’ and ‘library/pkgname/po/ll’, or at least those for which translations are available. The files ‘.mo’ are installed in the parallel places in the R installation tree.

For example, a package **Translation-it** might be prepared from an installed (and tested) version of R by

```
mkdir Translation-it
cd Translation-it
(cd "$R_HOME"; tar cf - share/locale/it library/*/po/it) | tar xf -
# the next step is not needed on Windows
msgfmt -c -o share/locale/it/LC_MESSAGES/RGui.mo $R_SRC_HOME/po/RGui-it.gmo
# create a DESCRIPTION file
cd ..
R CMD build Translation-it
```

It is probably appropriate to give the package a version number based on the version of R which has been translated. So the ‘DESCRIPTION’ file might look like

```
Package: Translation-it
Type: Translation
Version: 2.2.1-1
Title: Italian Translations for R 2.2.1
Description: Italian Translations for R 2.2.1
Author: The translators
```

```
Maintainer: Some Body <somebody@some.where.net>  
License: GPL (>= 2)
```

## 1.12 Services

Several members of the R project have set up services to assist those writing R packages, particularly those intended for public distribution.

[win-builder.r-project.org](http://win-builder.r-project.org) offers the automated preparation of (32/64-bit) Windows binaries from well-tested source packages.

R-Forge ([R-Forge.r-project.org](http://R-Forge.r-project.org)) and RForge ([www.rforge.net](http://www.rforge.net)) are similar services with similar names. Both provide source-code management through SVN, daily building and checking, mailing lists and a repository that can be accessed *via* `install.packages` (they can be selected by `setRepositories` and the GUI menus that use it). Package developers have the opportunity to present their work on the basis of project websites or news announcements. Mailing lists, forums or wikis provide useRs with convenient instruments for discussions and for exchanging information between developers and/or interested useRs.

## 2 Writing R documentation files

### 2.1 Rd format

R objects are documented in files written in “R documentation” (Rd) format, a simple markup language much of which closely resembles (La)TeX, which can be processed into a variety of formats, including L<sup>A</sup>T<sub>E</sub>X, HTML and plain text. The translation is carried out by functions in the **tools** package called by the script `Rdconv` in `'R_HOME/bin'` and by the installation scripts for packages.

The R distribution contains more than 1300 such files which can be found in the `'src/library/pkg/man'` directories of the R source tree, where *pkg* stands for one of the standard packages which are included in the R distribution.

As an example, let us look at a simplified version of `'src/library/base/man/load.Rd'` which documents the R function `load`.

```
% File src/library/base/man/load.Rd
\name{load}
\alias{load}
\title{Reload Saved Datasets}
\description{
  Reload the datasets written to a file with the function
  \code{save}.
}
\usage{
load(file, envir = parent.frame())
}
\arguments{
  \item{file}{a connection or a character string giving the
    name of the file to load.}
  \item{envir}{the environment where the data should be
    loaded.}
}
\seealso{
  \code{\link{save}}.
}
\examples{
## save all data
save(list = ls(), file= "all.RData")

## restore the saved values to the current environment
load("all.RData")

## restore the saved values to the workspace
load("all.RData", .GlobalEnv)
}
\keyword{file}
```

An ‘Rd’ file consists of three parts. The header gives basic information about the name of the file, the topics documented, a title, a short textual description and R usage information for the objects documented. The body gives further information (for example, on the function’s arguments and return value, as in the above example). Finally, there is an optional footer with keyword information. The header is mandatory.



Information is given within a series of *sections* with standard names (and user-defined sections are also allowed). Unless otherwise specified<sup>1</sup> these should occur only once in an ‘Rd’ file (in any order), and the processing software will retain only the first occurrence of a standard section in the file, with a warning.

See “**Guidelines for Rd files**” for guidelines for writing documentation in ‘Rd’ format which should be useful for package writers. The R generic function `prompt` is used to construct a bare-bones ‘Rd’ file ready for manual editing. Methods are defined for documenting functions (which fill in the proper function and argument names) and data frames. There are also functions `promptData`, `promptPackage`, `promptClass`, and `promptMethods` for other types of ‘Rd’ file.

The general syntax of ‘Rd’ files is summarized below. For a detailed technical discussion of current ‘Rd’ syntax, see “**Parsing Rd files**”. Note that there have been a number of changes to the ‘Rd’ format over the years, which can be important if a package is intended to be used with earlier versions of R: see earlier versions of this manual if a package is intended to be used with R before 2.10.0.

‘Rd’ files consists of three types of text input. The most common is L<sup>A</sup>T<sub>E</sub>X-like, with the backslash used as a prefix on markup (e.g. `\alias`), and braces used to indicate arguments (e.g. `{load}`). The least common type of text is verbatim text, where no markup is processed. The third type is R-like, intended for R code, but allowing some embedded macros. Quoted strings within R-like text are handled specially: regular character escapes such as `\n` may be entered as-is. Only markup starting with `\l` (e.g. `\link`) or `\v` (e.g. `\var`) will be recognized within quoted strings. The rarely used vertical tab `\v` must be entered as `\\v`.

Each macro defines the input type for its argument. For example, the file initially uses L<sup>A</sup>T<sub>E</sub>X-like syntax, and this is also used in the `\description` section, but the `\usage` section uses R-like syntax, and the `\alias` macro uses verbatim syntax. Comments run from a percent symbol `%` to the end of the line in all types of text (as on the first line of the `load` example).

Because backslashes, braces and percent symbols have special meaning, to enter them into text sometimes requires escapes using a backslash. In general balanced braces do not need to be escaped, but percent symbols always do. For the complete list of macros and rules for escapes, see “**Parsing Rd files**”.

### 2.1.1 Documenting functions

The basic markup commands used for documenting R objects (in particular, functions) are given in this subsection.

`\name{name}`

*name* typically<sup>2</sup> is the basename of the ‘Rd’ file containing the documentation. It is the “name” of the ‘Rd’ object represented by the file and has to be unique in a package. To avoid problems with indexing the package manual, it may not contain ‘!’ ‘|’ nor ‘@’, and to avoid possible problems with the HTML help system it should not contain ‘/’ nor a space. (L<sup>A</sup>T<sub>E</sub>X special characters are

<sup>1</sup> e.g. `\alias`, `\keyword` and `\note` sections.

<sup>2</sup> There can be exceptions: for example ‘Rd’ files are not allowed to start with a dot, and have to be uniquely named on a case-insensitive file system.

allowed, but may not be collated correctly in the index.) There can only be one `\name` entry in a file, and it must not contain any markup. Entries in the package manual will be in alphabetic<sup>3</sup> order of the `\name` entries.

#### `\alias{topic}`

The `\alias` sections specify all “topics” the file documents. This information is collected into index data bases for lookup by the on-line (plain text and HTML) help systems. The *topic* can contain spaces, but (for historical reasons) leading and trailing spaces will be stripped. Percent and left brace need to be escaped by a backslash.

There may be several `\alias` entries. Quite often it is convenient to document several R objects in one file. For example, file ‘`Normal.Rd`’ documents the density, distribution function, quantile function and generation of random variates for the normal distribution, and hence starts with

```
\name{Normal}
\alias{Normal}
\alias{dnorm}
\alias{pnorm}
\alias{qnorm}
\alias{rnorm}
```

Also, it is often convenient to have several different ways to refer to an R object, and an `\alias` does not need to be the name of an object.

Note that the `\name` is not necessarily a topic documented, and if so desired it needs to have an explicit `\alias` entry (as in this example).

#### `\title{Title}`

Title information for the ‘`Rd`’ file. This should be capitalized and not end in a period; try to limit its length to at most 65 characters for widest compatibility.

Since R version 2.12.0 markup has been supported in the text, but use of characters other than English text and punctuation (e.g., ‘`<`’) may limit portability.

There must be one (and only one) `\title` section in a help file.

#### `\description{...}`

A short description of what the function(s) do(es) (one paragraph, a few lines only). (If a description is too long and cannot easily be shortened, the file probably tries to document too much at once.) This is mandatory except for package-overview files.

#### `\usage{fun(arg1, arg2, ...)}`

One or more lines showing the synopsis of the function(s) and variables documented in the file. These are set in typewriter font. This is an R-like command.

The usage information specified should match the function definition *exactly* (such that automatic checking for consistency between code and documentation is possible).

---

<sup>3</sup> in the current locale, and with special treatment for L<sup>A</sup>T<sub>E</sub>X special characters and with any ‘*pkgname-package*’ topic moved to the top of the list.

It is no longer advisable to use `\synopsis` for the actual synopsis and show modified synopses in the `\usage`. Support for `\synopsis` will be removed eventually. To indicate that a function can be used in several different ways, depending on the named arguments specified, use section `\details`. E.g., ‘`abline.Rd`’ contains

```
\details{
  Typical usages are
  \preformatted{
    abline(a, b, untf = FALSE, \dots)
    .....
  }
```

Use `\method{generic}{class}` to indicate the name of an S3 method for the generic function *generic* for objects inheriting from class "*class*". In the printed versions, this will come out as *generic* (reflecting the understanding that methods should not be invoked directly but *via* method dispatch), but `codoc()` and other QC tools always have access to the full name.

For example, ‘`print.ts.Rd`’ contains

```
\usage{
  \method{print}{ts}(x, calendar, \dots)
}
```

which will print as

Usage:

```
## S3 method for class 'ts':
print(x, calendar, ...)
```

Usage for replacement functions should be given in the style of `dim(x) <- value` rather than explicitly indicating the name of the replacement function ("`dim<-`" in the above). Similarly, one can use `\method{generic}{class}(arglist) <- value` to indicate the usage of an S3 replacement method for the generic replacement function "*generic*<-" for objects inheriting from class "*class*".

Usage for S3 methods for extracting or replacing parts of an object, S3 methods for members of the Ops group, and S3 methods for user-defined (binary) infix operators ('`%xxx%`') follows the above rules, using the appropriate function names. E.g., ‘`Extract.factor.Rd`’ contains

```
\usage{
  \method{[]}{factor}(x, \dots, drop = FALSE)
  \method{[[]}{factor}(x, \dots)
  \method{[]}{factor}(x, \dots) <- value
}
```

which will print as

Usage:

```
## S3 method for class 'factor':
x[... , drop = FALSE]
## S3 method for class 'factor':
x[[...]]
## S3 replacement method for class 'factor':
x[...] <- value
```

`\S3method` is accepted as an alternative to `\method`.

`\arguments{...}`

Description of the function's arguments, using an entry of the form

```
\item{arg_i}{Description of arg_i.}
```

for each element of the argument list. (Note that there is no whitespace between the three parts of the entry.) There may be optional text outside the `\item` entries, for example to give general information about groups of parameters.

`\details{...}`

A detailed if possible precise description of the functionality provided, extending the basic information in the `\description` slot.

`\value{...}`

Description of the function's return value.

If a list with multiple values is returned, you can use entries of the form

```
\item{comp_i}{Description of comp_i.}
```

for each component of the list returned. Optional text may precede<sup>4</sup> this list (see for example the help for `rle`). Note that `\value` is implicitly a `\describe` environment, so that environment should not be used for listing components, just individual `\item{...}` entries.

`\references{...}`

A section with references to the literature. Use `\url{...}` or `\href{...}{...}` for web pointers.

`\note{...}`

Use this for a special note you want to have pointed out. Multiple `\note` sections are allowed, but might be confusing to the end users.

For example, 'pie.Rd' contains

```
\note{
  Pie charts are a very bad way of displaying information.
  The eye is good at judging linear measures and bad at
  judging relative areas.
  .....
}
```

---

<sup>4</sup> Text between or after list items was discarded prior to R 2.10.0, and is discouraged.

`\author{...}`

Information about the author(s) of the ‘Rd’ file. Use `\email{}` without extra delimiters (such as ‘( )’ or ‘< >’) to specify email addresses, or `\url{}` or `\href{}{}` for web pointers.

`\seealso{...}`

Pointers to related R objects, using `\code{\link{...}}` to refer to them (`\code` is the correct markup for R object names, and `\link` produces hyperlinks in output formats which support this. See [Section 2.3 \[Marking text\]](#), page 63, and [Section 2.5 \[Cross-references\]](#), page 66).

`\examples{...}`

Examples of how to use the function. Code in this section is set in typewriter font without reformatting and is run by `example()` unless marked otherwise (see below).

Examples are not only useful for documentation purposes, but also provide test code used for diagnostic checking of R code. By default, text inside `\examples{}` will be displayed in the output of the help page and run by `example()` and by R CMD `check`. You can use `\dontrun{}` for text that should only be shown, but not run, and `\dontshow{}` for extra commands for testing that should not be shown to users, but will be run by `example()`. (Previously this was called `\testonly`, and that is still accepted.)

Text inside `\dontrun{}` is verbatim, but the other parts of the `\examples` section are R-like text.

For example,

```
x <- runif(10)      # Shown and run.
\dontrun{plot(x)}   # Only shown.
\dontshow{log(x)}   # Only run.
```

Thus, example code not included in `\dontrun` must be executable! In addition, it should not use any system-specific features or require special facilities (such as Internet access or write permission to specific directories). Text included in `\dontrun` is indicated by comments in the processed help files: it need not be valid R code but the escapes must still be used for `%`, `\` and unpaired braces as in other verbatim text.

Example code must be capable of being run by `example`, which uses `source`. This means that it should not access ‘`stdin`’, e.g. to `scan()` data from the example file.

Data needed for making the examples executable can be obtained by random number generation (for example, `x <- rnorm(100)`), or by using standard data sets listed by `data()` (see `?data` for more info).

Finally, there is `\donttest`, used (at the beginning of a separate line) to mark code that should be run by `examples()` but not by R CMD `check`. This should be needed only occasionally but can be used for code which might fail in circumstances that are hard to test for, for example in some locales. (Use e.g. `capabilities()` to test for features needed in the examples wherever possible, and you can also use `try()` or `trycatch()`.)

**\keyword{key}**

There can be zero or more `\keyword` sections per file. Each `\keyword` section should specify a single keyword, preferably one of the standard keywords as listed in file ‘KEYWORDS’ in the R documentation directory (default ‘*R\_HOME/doc*’). Use e.g. `RShowDoc("KEYWORDS")` to inspect the standard keywords from within R. There can be more than one `\keyword` entry if the R object being documented falls into more than one category, or none.

The special keyword ‘`internal`’ marks a page of internal objects that are not part of the package’s API. If the help page for object `foo` has keyword ‘`internal`’, then `help(foo)` gives this help page, but `foo` is excluded from several object indices, including the alphabetical list of objects in the HTML help system.

`help.search()` can search by keyword, including user-defined values: however the ‘Search Engine & Keywords’ HTML page accessed *via* `help.start()` provides single-click access only to a pre-defined list of keywords.

### 2.1.2 Documenting data sets

The structure of ‘`Rd`’ files which document R data sets is slightly different. Sections such as `\arguments` and `\value` are not needed but the format and source of the data should be explained.

As an example, let us look at ‘`src/library/datasets/man/rivers.Rd`’ which documents the standard R data set `rivers`.

```
\name{rivers}
\docType{data}
\alias{rivers}
\title{Lengths of Major North American Rivers}
\description{
  This data set gives the lengths (in miles) of 141 \dQuote{major}
  rivers in North America, as compiled by the US Geological
  Survey.
}
\usage{rivers}
\format{A vector containing 141 observations.}
\source{World Almanac and Book of Facts, 1975, page 406.}
\references{
  McNeil, D. R. (1977) \emph{Interactive Data Analysis}.
  New York: Wiley.
}
\keyword{datasets}
```

This uses the following additional markup commands.

**\docType{...}**

Indicates the “type” of the documentation object. Always ‘`data`’ for data sets, and ‘`package`’ for ‘`pkg-package.Rd`’ overview files. Documentation for S4 methods and classes uses ‘`methods`’ (from `promptMethods()`) and ‘`class`’ (from `promptClass()`).

`\format{...}`

A description of the format of the data set (as a vector, matrix, data frame, time series, ...). For matrices and data frames this should give a description of each column, preferably as a list or table. See [Section 2.4 \[Lists and tables\]](#), [page 65](#), for more information.

`\source{...}`

Details of the original source (a reference or URL). In addition, section `\references` could give secondary sources and usages.

Note also that when documenting data set *bar*,

- The `\usage` entry is always *bar* or (for packages which do not use lazy-loading of data) `data(bar)`. (In particular, only document a *single* data object per ‘Rd’ file.)
- The `\keyword` entry should always be ‘*datasets*’.

If *bar* is a data frame, documenting it as a data set can be initiated *via* `prompt(bar)`. Otherwise, the `promptData` function may be used.

### 2.1.3 Documenting S4 classes and methods

There are special ways to use the ‘?’ operator, namely ‘`class?topic`’ and ‘`methods?topic`’, to access documentation for S4 classes and methods, respectively. This mechanism depends on conventions for the topic names used in `\alias` entries. The topic names for S4 classes and methods respectively are of the form

```
class-class
generic,signature_list-method
```

where *signature\_list* contains the names of the classes in the signature of the method (without quotes) separated by ‘,’ (without whitespace), with ‘ANY’ used for arguments without an explicit specification. E.g., ‘`genericFunction-class`’ is the topic name for documentation for the S4 class “`genericFunction`”, and ‘`coerce,ANY,NULL-method`’ is the topic name for documentation for the S4 method for `coerce` for signature `c("ANY", "NULL")`.

Skeletons of documentation for S4 classes and methods can be generated by using the functions `promptClass()` and `promptMethods()` from package **methods**. If it is necessary or desired to provide an explicit function declaration (in a `\usage` section) for an S4 method (e.g., if it has “surprising arguments” to be mentioned explicitly), one can use the special markup

```
\S4method{generic}{signature_list}(argument_list)
(e.g., ‘\S4method{coerce}{ANY,NULL}(from, to)’).
```

To make full use of the potential of the on-line documentation system, all user-visible S4 classes and methods in a package should at least have a suitable `\alias` entry in one of the package’s ‘Rd’ files. If a package has methods for a function defined originally somewhere else, and does not change the underlying default method for the function, the package is responsible for documenting the methods it creates, but not for the function itself or the default method.

An S4 replacement method is documented in the same way as an S3 one: see the description of `\method` in [Section 2.1.1 \[Documenting functions\]](#), [page 56](#).

See `help("Documentation", package = "methods")` for more information on using and creating on-line documentation for S4 classes and methods.



### 2.1.4 Documenting packages

Packages may have an overview help page with an `\alias pkgname-package`, e.g. ‘`utils-package`’ for the `utils` package, when `package?pkgname` will open that help page. If a topic named `pkgname` does not exist in another ‘Rd’ file, it is helpful to use this as an additional `\alias`.

Skeletons of documentation for a package can be generated using the function `promptPackage()`. If the `final = TRUE` argument is used, then the ‘Rd’ file will be generated in final form, containing the information that would be produced up to `library(help = pkgname)`. Otherwise (the default) comments will be inserted giving suggestions for content.

Apart from the mandatory `\name` and `\title` and the `pkgname-package` alias, the only requirement for the package overview page is that it include a `\docType{package}` statement. All other content is optional. We suggest that it should be a short overview, to give a reader unfamiliar with the package enough information to get started. More extensive documentation is better placed into a package vignette (see [Section 1.4 \[Writing package vignettes\]](#), page 32) and referenced from this page, or into individual man pages for the functions, datasets, or classes.

## 2.2 Sectioning

To begin a new paragraph or leave a blank line in an example, just insert an empty line (as in (La)TeX). To break a line, use `\cr`.

In addition to the predefined sections (such as `\description{}`, `\value{}`, etc.), you can “define” arbitrary ones by `\section{section_title}{...}`. For example

```
\section{Warning}{
  You must not call this function unless ...
}
```

For consistency with the pre-assigned sections, the section name (the first argument to `\section`) should be capitalized (but not all upper case). Whitespace between the first and second braced expressions is not allowed. Markup (e.g. `\code`) within the section title may cause problems with the latex conversion (depending on the version of macro packages such as ‘`hyperref`’) and so should be avoided.

The `\subsection` macro takes arguments in the same format as `\section`, but is used within a section, so it may be used to nest subsections within sections or other subsections. There is no predefined limit on the nesting level, but formatting is not designed for more than 3 levels (i.e. subsections within subsections within sections).

Note that additional named sections are always inserted at a fixed position in the output (before `\note`, `\seealso` and the examples), no matter where they appear in the input (but in the same order amongst themselves as in the input).

## 2.3 Marking text

The following logical markup commands are available for emphasizing or quoting text.

`\emph{text}`

`\strong{text}`

Emphasize *text* using *italic* and **bold** font if possible; `\strong` is regarded as stronger (more emphatic).

`\bold{text}`

Set *text* in **bold** font if possible.

`\sQuote{text}`

`\dQuote{text}`

Portably single or double quote *text* (without hard-wiring the characters used for quotation marks).

Each of the above commands takes L<sup>A</sup>T<sub>E</sub>X-like input, so other macros may be used within *text*.

The following logical markup commands are available for indicating specific kinds of text. Except as noted, these take verbatim text input, and so other macros may not be used within them. Some characters will need to be escaped (see [Section 2.8 \[Insertions\]](#), [page 68](#)).

`\code{text}`

Indicate text that is a literal example of a piece of an R program, e.g., a fragment of R code or the name of an R object. Text is entered in R-like syntax, and displayed using **typewriter** font if possible. Macros `\var` and `\link` are interpreted within *text*.

`\preformatted{text}`

Indicate text that is a literal example of a piece of a program. Text is displayed using **typewriter** font if possible. Formatting, e.g. line breaks, is preserved.

Due to limitations in L<sup>A</sup>T<sub>E</sub>X as of this writing, this macro may not be nested within other markup macros other than `\dQuote` and `\sQuote`, as errors or bad formatting may result.

`\kbd{keyboard-characters}`

Indicate keyboard input, using *slanted typewriter* font if possible, so users can distinguish the characters they are supposed to type from computer output. Text is entered verbatim.

`\samp{text}`

Indicate text that is a literal example of a sequence of characters, entered verbatim. No wrapping or reformatting will occur. Displayed using **typewriter** font if possible.

`\verb{text}`

Indicate text that is a literal example of a sequence of characters, with no interpretation of e.g. `\var`, but which will be included within word-wrapped text. Displayed using **typewriter** font if possible.

`\pkg{package_name}`

Indicate the name of an R package. L<sup>A</sup>T<sub>E</sub>X-like.

|  |   |
|--|---|
| <code>\file{file_name}</code>                      | Indicate the name of a file. Text is L <sup>A</sup> T <sub>E</sub> X-like, so backslash needs to be escaped. Displayed using a distinct font if possible.   |
| <code>\email{email_address}</code>                 | Indicate an electronic mail address. L <sup>A</sup> T <sub>E</sub> X-like, will be rendered as a hyperlink in HTML and PDF conversion. Displayed using <b>typewriter</b> font if possible.                                |
| <code>\url{uniform_resource_locator}</code>        | Indicate a uniform resource locator (URL) for the World Wide Web. The argument is handled verbatim, and rendered as a hyperlink in HTML and PDF conversion. Displayed using <b>typewriter</b> font if possible.           |
| <code>\href{uniform_resource_locator}{text}</code> | Indicate a hyperlink to the World Wide Web. The first argument is handled verbatim, and is used as the URL in the hyperlink, with the second argument of L <sup>A</sup> T <sub>E</sub> X-like text displayed to the user. |
| <code>\var{metasyntactic_variable}</code>          | Indicate a metasyntactic variable. In some cases this will be rendered distinctly, e.g. in italic, but not in all <sup>5</sup> . L <sup>A</sup> T <sub>E</sub> X-like.  |
| <code>\env{environment_variable}</code>            | Indicate an environment variable. Verbatim. Displayed using <b>typewriter</b> font if possible  |
| <code>\option{option}</code>                       | Indicate a command-line option. Verbatim. Displayed using <b>typewriter</b> font if possible.   |
| <code>\command{command_name}</code>                | Indicate the name of a command. L <sup>A</sup> T <sub>E</sub> X-like, so <code>\var</code> is interpreted. Displayed using <b>typewriter</b> font if possible.  |
| <code>\dfn{term}</code>                            | Indicate the introductory or defining use of a term. L <sup>A</sup> T <sub>E</sub> X-like.  |
| <code>\cite{reference}</code>                      | Indicate a reference without a direct cross-reference <i>via</i> <code>\link</code> (see <a href="#">Section 2.5 [Cross-references]</a> , page 66), such as the name of a book. L <sup>A</sup> T <sub>E</sub> X-like.     |
| <code>\acronym{acronym}</code>                     | Indicate an acronym (an abbreviation written in all capital letters), such as GNU. L <sup>A</sup> T <sub>E</sub> X-like.  |

## 2.4 Lists and tables

The `\itemize` and `\enumerate` commands take a single argument, within which there may be one or more `\item` commands. The text following each `\item` is formatted as one or more paragraphs, suitably indented and with the first paragraph marked with a bullet point (`\itemize`) or a number (`\enumerate`).

---

<sup>5</sup> Currently it is rendered differently only in HTML conversions, and L<sup>A</sup>T<sub>E</sub>X conversion outside ‘`\usage`’ and ‘`\examples`’ environments.

Note that unlike argument lists, `\item` in these formats is followed by a space and the text (not enclosed in braces). For example

```
\enumerate{
  \item A database consists of one or more records, each with one or
    more named fields.
  \item Regular lines start with a non-whitespace character.
  \item Records are separated by one or more empty lines.
}
```

`\itemize` and `\enumerate` commands may be nested.

The `\describe` command is similar to `\itemize` but allows initial labels to be specified. Each `\item` takes two arguments, the label and the body of the item, in exactly the same way as an argument or value `\item`. `\describe` commands are mapped to `<DL>` lists in HTML and `\description` lists in L<sup>A</sup>T<sub>E</sub>X.

The `\tabular` command takes two arguments. The first gives for each of the columns the required alignment ('l' for left-justification, 'r' for right-justification or 'c' for centring.) The second argument consists of an arbitrary number of lines separated by `\cr`, and with fields separated by `\tab`. For example:

```
\tabular{rllll}{
  [,1] \tab Ozone      \tab numeric \tab Ozone (ppb)\cr
  [,2] \tab Solar.R    \tab numeric \tab Solar R (lang)\cr
  [,3] \tab Wind       \tab numeric \tab Wind (mph)\cr
  [,4] \tab Temp       \tab numeric \tab Temperature (degrees F)\cr
  [,5] \tab Month      \tab numeric \tab Month (1--12)\cr
  [,6] \tab Day        \tab numeric \tab Day of month (1--31)
}
```

There must be the same number of fields on each line as there are alignments in the first argument, and they must be non-empty (but can contain only spaces). (There is no white-space between `\tabular` and the first argument, nor between the two arguments.)

## 2.5 Cross-references

The markup `\link{foo}` (usually in the combination `\code{\link{foo}}`) produces a hyperlink to the help for *foo*. Here *foo* is a *topic*, that is the argument of `\alias` markup in another 'Rd' file (possibly in another package). Hyperlinks are supported in some of the formats to which 'Rd' files are converted, for example HTML and PDF, but ignored in others, e.g. the text format.

One main usage of `\link` is in the `\seealso` section of the help page, see [Section 2.1 \[Rd format\]](#), page 55.

Note that whereas leading and trailing spaces are stripped when extracting a topic from a `\alias`, they are not stripped when looking up the topic of a `\link`.

You can specify a link to a different topic than its name by `\link[=dest]{name}` which links to topic *dest* with name *name*. This can be used to refer to the documentation for S3/4 classes, for example `\code{"\link[=abc-class]{abc}"}` would be a way to refer to the documentation of an S4 class "abc" defined in your package, and `\code{"\link[=terms.object]{terms}"}` to the S3 "terms" class (in package `stats`). To

make these easy to read in the source file, `\code{"\linkS4class{abc}"}` expands to the form given above.

There are two other forms of optional argument specified as `\link[pkg]{foo}` and `\link[pkg:bar]{foo}` to link to the package *pkg*, to files ‘*foo.html*’ and ‘*bar.html*’ respectively. These are rarely needed, perhaps to refer to not-yet-installed packages (but there the HTML help system will resolve the link at run time) or in the normally undesirable event that more than one package offers help on a topic<sup>6</sup> (in which case the present package has precedence so this is only needed to refer to other packages). They are currently only used in HTML help (and ignored for hyperlinks in L<sup>A</sup>T<sub>E</sub>X conversions of help pages), and link to the file rather than the topic (since there is no way to know which topics are in which files in an uninstalled package). The **only** reason to use these forms for base and recommended packages is to force a reference to a package that might be further down the search path. Because they have been frequently misused, the HTML help system looks for topic *foo* in package *pkg* if it does not find file ‘*foo.html*’.

## 2.6 Mathematics

Mathematical formulae should be set beautifully for printed documentation yet we still want something useful for text and HTML online help. To this end, the two commands `\eqn{latex}{ascii}` and `\deqn{latex}{ascii}` are used. Whereas `\eqn` is used for “inline” formulae (corresponding to T<sub>E</sub>X’s  $\dots$ ), `\deqn` gives “displayed equations” (as in L<sup>A</sup>T<sub>E</sub>X’s `displaymath` environment, or T<sub>E</sub>X’s  $\dots$ ). Both arguments are treated as verbatim text.

Both commands can also be used as `\eqn{latexascii}` (only *one* argument) which then is used for both *latex* and *ascii*. No whitespace is allowed between command and the first argument, nor between the first and second arguments.

The following example is from ‘Poisson.Rd’:

```
\deqn{p(x) = \frac{\lambda^x e^{-\lambda}}{x!}}{%
      p(x) = \lambda^x exp(-\lambda)/x!}
for \eqn{x = 0, 1, 2, \ldots}.
```

For the L<sup>A</sup>T<sub>E</sub>X manual, this becomes

$$p(x) = \lambda^x \frac{e^{-\lambda}}{x!}$$

for  $x = 0, 1, 2, \dots$

For text on-line help we get

```
p(x) = lambda^x exp(-lambda)/x!
for x = 0, 1, 2, ....
```

Greek letters (both cases) will be rendered in HTML if preceded by a backslash, `\dots` and `\ldots` will be rendered as ellipses and `\sqrt`, `\ge` and `\le` as mathematical symbols.

<sup>6</sup> a common example in CRAN packages is `\link[mgcv]{gam}`.

Note that only basic L<sup>A</sup>T<sub>E</sub>X can be used, there being no provision to specify L<sup>A</sup>T<sub>E</sub>X style files such as the AMS extensions.

## 2.7 Figures

To include figures in help pages, use the `\figure` markup. There are three forms.

The two commonly used simple forms are `\figure{filename}` and `\figure{filename}{alternate text}`. This will include a copy of the figure in either HTML or L<sup>A</sup>T<sub>E</sub>X output. In text output, the alternate text will be displayed instead. (When the second argument is omitted, the filename will be used.) Both the filename and the alternate text will be parsed verbatim, and should not include special characters that are significant in HTML or L<sup>A</sup>T<sub>E</sub>X.

The expert form is `\figure{filename}{options: string}`. (The word ‘options:’ must be typed exactly as shown and followed by at least one space.) In this form, the *string* is copied into the HTML `img` tag as attributes following the `src` attribute, or into the second argument of the `\Figure` macro in LaTeX, which by default is used as options to an `\includegraphics` call. As it is unlikely that any single string would suffice for both display modes, the expert form would normally be wrapped in conditionals. It is up to the author to make sure that legal HTML/LaTeX is used. For example, to include a logo in both HTML (using the simple form) and LaTeX (using the expert form), the following could be used:

```
\if{html}{\figure{logo.jpg}{Our logo}}
\if{latex}{\figure{logo.jpg}{options: width=0.5in}}
```

The files containing the figures should be stored in the directory ‘`man/figures`’. Files with extensions ‘`.jpg`’, ‘`.pdf`’, ‘`.png`’ and ‘`.svg`’ from that directory will be copied to the ‘`help/figures`’ directory at install time. (Figures in PDF format will not display in most HTML browsers, but might be the best choice in reference manuals.) Specify the filename relative to ‘`man/figures`’ in the `\figure` directive.

## 2.8 Insertions

Use `\R` for the R system itself. Use `\dots` for the dots in function argument lists ‘`...`’, and `\ldots` for ellipsis dots in ordinary text.<sup>7</sup> These can be followed by `{}`, and should be unless followed by whitespace.

After an unescaped ‘`%`’, you can put your own comments regarding the help text. The rest of the line (but not the newline at the end) will be completely disregarded. Therefore, you can also use it to make part of the “help” invisible.

You can produce a backslash (‘`\`’) by escaping it by another backslash. (Note that `\cr` is used for generating line breaks.)

The “comment” character ‘`%`’ and unpaired braces<sup>8</sup> *almost always* need to be escaped by ‘`\`’, and ‘`\\`’ can be used for backslash and needs to be when there two or more adjacent

<sup>7</sup> There is only a fine distinction between `\dots` and `\ldots`. It is technically incorrect to use `\ldots` in code blocks and `tools::checkRd` will warn about this—on the other hand the current converters treat them the same way in code blocks, and elsewhere apart from the small distinction between the two in LaTeX.

<sup>8</sup> See the examples section in the file ‘`Paren.Rd`’ for an example.

backslashes). In R-like code quoted strings are handled slightly differently; see “[Parsing Rd files](#)” for details – in particular braces should not be escaped in quoted strings.

All of ‘% { } \’ should be escaped in L<sup>A</sup>T<sub>E</sub>X-like text.

Text which might need to be represented differently in different encodings should be marked by `\enc`, e.g. `\enc{Jöreskog}{Joreskog}` (with no whitespace between the braces) where the first argument will be used where encodings are allowed and the second should be ASCII (and is used for e.g. the text conversion in locales that cannot represent the encoded form). (This is intended to be used for individual words, not whole sentences or paragraphs.)

## 2.9 Indices

The `\alias` command (see [Section 2.1.1 \[Documenting functions\]](#), page 56) is used to specify the “topics” documented, which should include *all* R objects in a package such as functions and variables, data sets, and S4 classes and methods (see [Section 2.1.3 \[Documenting S4 classes and methods\]](#), page 62). The on-line help system searches the index data base consisting of all alias topics.

In addition, it is possible to provide “concept index entries” using `\concept`, which can be used for `help.search()` lookups. E.g., file ‘`cor.test.Rd`’ in the standard package `stats` contains

```
\concept{Kendall correlation coefficient}
\concept{Pearson correlation coefficient}
\concept{Spearman correlation coefficient}
```

so that e.g. `??Spearman` will succeed in finding the help page for the test for association between paired samples using Spearman’s  $\rho$ .

(Note that `help.search()` only uses “sections” of documentation objects with no additional markup.)

If you want to cross reference such items from other help files *via* `\link`, you need to use `\alias` and not `\concept`.

## 2.10 Platform-specific documentation

Sometimes the documentation needs to differ by platform. Currently two OS-specific options are available, ‘`unix`’ and ‘`windows`’, and lines in the help source file can be enclosed in

```
#ifdef OS
...
#endif
```

or

```
#ifndef OS
...
#endif
```

for OS-specific inclusion or exclusion. Such blocks should not be nested, and should be entirely within a block (that, is between the opening and closing brace of a section or item), or at top-level contain one or more complete sections.

If the differences between platforms are extensive or the R objects documented are only relevant to one platform, platform-specific ‘Rd’ files can be put in a ‘`unix`’ or ‘`windows`’ subdirectory.



## 2.11 Conditional text

Occasionally the best content for one output format is different from the best content for another. For this situation, the `\if{format}{text}` or `\ifelse{format}{text}{alternate}` markup is used. Here *format* is a comma separated list of formats in which the *text* should be rendered. The *alternate* will be rendered if the format does not match. Both *text* and *alternate* may be any sequence of text and markup.

Currently the following formats are recognized: `example`, `html`, `latex` and `text`. These select output for the corresponding targets. (Note that `example` refers to extracted example code rather than the displayed example in some other format.) Also accepted are `TRUE` (matching all formats) and `FALSE` (matching no formats). These could be the output of the `\Sexpr` macro (see [Section 2.12 \[Dynamic pages\]](#), page 70).

The `\out{literal}` macro would usually be used within the *text* part of `\if{format}{text}`. It causes the renderer to output the literal text exactly, with no attempt to escape special characters. For example, use the following to output the markup necessary to display the Greek letter in L<sup>A</sup>T<sub>E</sub>X or HTML, and the text string `alpha` in other formats:

```
\if{latex}{\out{\alpha}}\ifelse{html}{\out{&alpha;}}{\alpha}
```

## 2.12 Dynamic pages

Two new macros supporting dynamically generated man pages were introduced in R 2.10.0, `\Sexpr` and `\RdOpts`. These are modelled after Sweave, and are intended to contain executable R expressions in the ‘Rd’ file.

The main argument to `\Sexpr` must be valid R code that can be executed. It may also take options in square brackets before the main argument. Depending on the options, the code may be executed at package build time, package install time, or man page rendering time.

The options follow the same format as in Sweave, but different options are supported. Currently the allowed options and their defaults are:

- `eval=TRUE` Whether the R code should be evaluated.
- `echo=FALSE` Whether the R code should be echoed. If `TRUE`, a display will be given in a preformatted block. For example, `\Sexpr[echo=TRUE]{ x <- 1 }` will be displayed as
 

```
> x <- 1
```
- `keep.source=TRUE` Whether to keep the author’s formatting when displaying the code, or throw it away and use a deparsed version.
- `results=text` How should the results be displayed? The possibilities are:
  - `results=text` Apply `as.character()` to the result of the code, and insert it as a text element.
  - `results=verbatim` Print the results of the code just as if it was executed at the console, and include the printed results verbatim. (Invisible results will not print.)
  - `results=rd` The result is assumed to be a character vector containing markup to be passed to `parse_Rd()`, with the result inserted in place. This could be used to

insert computed aliases, for instance. As of R 2.13.1-patched, `parse_Rd()` is called first with `fragment=FALSE` to allow a single Rd section macro to be inserted. If that fails, it is called again with `fragment=TRUE`, the older behavior.

- `results=hide` Insert no output.
- `strip.white=TRUE` Remove leading and trailing white space from each line of output if `strip.white=TRUE`. With `strip.white=all`, also remove blank lines.
- `stage=install` Control when this macro is run. Possible values are
  - `stage=build` The macro is run when building a source tarball.
  - `stage=install` The macro is run when installing from source.
  - `stage=render` The macro is run when displaying the help page.

Conditionals such as `#ifdef` (see [Section 2.10 \[Platform-specific sections\]](#), page 69) are applied after the `build` macros but before the `install` macros. In some situations (e.g. installing directly from a source directory without a tarball, or building a binary package) the above description is not literally accurate, but authors can rely on the sequence being `build`, `#ifdef`, `install`, `render`, with all stages executed.

Code is only run once in each stage, so a `\Sexpr[results=rd]` macro can output an `\Sexpr` macro designed for a later stage, but not for the current one or any earlier stage.

- `width`, `height`, `fig` These options are currently allowed but ignored.

The `\RdOpts` macro is used to set new defaults for options to apply to following uses of `\Sexpr`.

For more details, see the online document [“Parsing Rd files”](#).

## 2.13 User-defined macros

Two new macros supporting user-defined macros were introduced in R 2.12.0. The `\newcommand` and `\renewcommand` macros allow new macros to be defined within an Rd file. These are similar but not identical to the same-named LaTeX macros.

They each take two arguments which are parsed verbatim. The first is the name of the new macro including the initial backslash, and the second is the macro definition. As in  $\text{\LaTeX}$ , `\newcommand` requires that the new macro not have been previously defined, whereas `\renewcommand` allows existing macros (including all built-in ones) to be replaced.

Also as in  $\text{\LaTeX}$ , the new macro may be defined to take arguments, and numeric placeholders such as `#1` are used in the macro definition. However, unlike  $\text{\LaTeX}$ , the number of arguments is determined automatically from the highest placeholder number seen in the macro definition. For example, a macro definition containing `#1` and `#3` (but no other placeholders) will define a three argument macro (whose second argument will be ignored). As in  $\text{\LaTeX}$ , at most 9 arguments may be defined. If the `#` character is followed by a non-digit it will have no special significance. All arguments to user-defined macros will be parsed as verbatim text, and simple text-substitution will be used to replace the place-holders, after which the replacement text will be parsed.

For example, the ‘NEWS.Rd’ file currently uses the definition

```
\newcommand{\PR}{\Sexpr[results=rd]{tools:::Rd_expr_PR(#1)}}
```

which defines `\PR` to be a single argument macro; then code like

```
\PR{1234}
```

will expand to

```
\Sexpr[results=rd]{tools:::Rd_expr_PR(1234)}
```

when parsed.

## 2.14 Encoding

Rd files are text files and so it is impossible to deduce the encoding they are written in unless ASCII: files with 8-bit characters could be UTF-8, Latin-1, Latin-9, KOI8-R, EUC-JP, *etc.* So an `\encoding{}` section must be used to specify the encoding if it is not ASCII. (The `\encoding{}` section must be on a line by itself, and in particular one containing no non-ASCII characters. The encoding declared in the ‘DESCRIPTION’ file will be used if none is declared in the file.) The ‘Rd’ files are converted to UTF-8 before parsing and so the preferred encoding for the files themselves is now UTF-8.

Wherever possible, avoid non-ASCII chars in ‘Rd’ files, and even symbols such as ‘<’, ‘>’, ‘\$’, ‘^’, ‘&’, ‘|’, ‘@’, ‘~’, and ‘\*’ outside verbatim environments (since they may disappear in fonts designed to render text). (Function `showNonASCIIfile` in package **tools** can help in finding non-ASCII bytes in the files.)

For convenience, encoding names ‘latin1’ and ‘latin2’ are always recognized: these and ‘UTF-8’ are likely to work fairly widely. However, this does not mean that all characters in UTF-8 will be recognized, and the coverage of non-Latin characters<sup>9</sup> is fairly low. Using  $\LaTeX$  `inputenc` (see `?Rd2pdf` in R) will give greater coverage of UTF-8.

The `\enc` command (see [Section 2.8 \[Insertions\]](#), page 68) can be used to provide transliterations which will be used in conversions that do not support the declared encoding.

The  $\LaTeX$  conversion converts the file to UTF-8 from the declared encoding, and includes a

```
\inputencoding{utf8}
```

command, and this needs to be matched by a suitable invocation of the `\usepackage{inputenc}` command. The R utility `R CMD Rd2pdf` looks at the converted code and includes the encodings used: it might for example use

```
\usepackage[utf8]{inputenc}
```

(Use of `utf8` as an encoding requires  $\LaTeX$  dated 2003/12/01 or later. Also, the use of Cyrillic characters in ‘UTF-8’ appears to also need ‘`\usepackage[T2A]{fontenc}`’, and R `CMD Rd2pdf` includes this conditionally on the file ‘`t2aenc.def`’ being present and environment variable `_R_CYRILLIC_TEX_` being set.)

Note that this mechanism works best with Latin letters: the coverage of UTF-8 in  $\LaTeX$  is quite low.

---

<sup>9</sup> R 2.9.0 added support for UTF-8 Cyrillic characters in  $\LaTeX$ , but on some OSes this will need Cyrillic support added to  $\LaTeX$ , so environment variable `_R_CYRILLIC_TEX_` may need to be set to a non-empty value to enable this.

## 2.15 Processing Rd format

There are several commands to process Rd files from the system command line.

Using R CMD `Rdconv` one can convert R documentation format to other formats, or extract the executable examples for run-time testing. The currently supported conversions are to plain text, HTML and L<sup>A</sup>T<sub>E</sub>X as well as extraction of the examples.

R CMD `Rd2pdf` generates PDF output from documentation in ‘Rd’ files, which can be specified either explicitly or by the path to a directory with the sources of a package. In the latter case, a reference manual for all documented objects in the package is created, including the information in the ‘DESCRIPTION’ files.

R CMD `Sweave` and R CMD `Stangle` process ‘Sweave’ documentation files (usually with extension ‘.Snw’ or ‘.Rnw’): R CMD `Stangle` is used to extract the R code fragments.

The exact usage and a detailed list of available options for all of these commands can be obtained by running R CMD `command --help`, e.g., `R CMD Rdconv --help`. All available commands can be listed using `R --help` (or `Rcmd --help` under Windows).

All of these work under Windows. You may need to have installed the tools to build packages from source as described in the “R Installation and Administration” manual, although typically all that is needed is a L<sup>A</sup>T<sub>E</sub>X installation.

## 2.16 Editing Rd files

It can be very helpful to prepare ‘.Rd’ files using an editor which knows about their syntax and will highlight commands, indent to show the structure and detect mis-matched braces, and so on.

The system most commonly used for this is some version of Emacs (including XEmacs) with the ESS package (<http://ess.r-project.org/>: it is often installed with Emacs but may need to be loaded, or even installed, separately).

Another is the Eclipse IDE with the Stat-ET plugin (<http://www.walware.de/goto/statet>), and (on Windows only) Tinn-R (<http://sourceforge.net/projects/tinn-r/>).

People have also used L<sup>A</sup>T<sub>E</sub>X mode in an editor, as ‘.Rd’ files are rather similar to L<sup>A</sup>T<sub>E</sub>X files.

Some R front-ends provide editing support for ‘.Rd’ files, for example RStudio (<http://rstudio.org/>).

## 3 Tidying and profiling R code

R code which is worth preserving in a package and perhaps making available for others to use is worth documenting, tidying up and perhaps optimizing. The last two of these activities are the subject of this chapter.

### 3.1 Tidying R code

R treats function code loaded from packages and code entered by users differently. By default code entered by users has the source code stored internally, and when the function is listed, the original source is reproduced. Loading code from a package (by default) discards the source code, and the function listing is re-created from the parse tree of the function.

Normally keeping the source code is a good idea, and in particular it avoids comments being removed from the source. However, we can make use of the ability to re-create a function listing from its parse tree to produce a tidy version of the function, for example with consistent indentation and spaces around operators. If the original source does not follow the standard format this tidied version can be much easier to read.

We can subvert the keeping of source in two ways.

1. The option `keep.source` can be set to `FALSE` before the code is loaded into R.
2. The stored source code can be removed by calling the `removeSource()` function, for example by

```
myfun <- removeSource(myfun)
```

In each case if we then list the function we will get the standard layout.

Suppose we have a file of functions ‘myfuncs.R’ that we want to tidy up. Create a file ‘tidy.R’ containing

```
source("myfuncs.R", keep.source = FALSE)
dump(ls(all = TRUE), file = "new.myfuncs.R")
```

and run R with this as the source file, for example by `R --vanilla < tidy.R` or by pasting into an R session. Then the file ‘new.myfuncs.R’ will contain the functions in alphabetical order in the standard layout. Warning: comments in your functions will be lost.

The standard format provides a good starting point for further tidying. Although the deparsing cannot do so, we recommend the consistent use of the preferred assignment operator ‘<-’ (rather than ‘=’) for assignment. Many package authors use a version of Emacs (on a Unix-alike or Windows) to edit R code, using the ESS[S] mode of the ESS Emacs package. See [Section “R coding standards” in \*R Internals\*](#) for style options within the ESS[S] mode recommended for the source code of R itself.

### 3.2 Profiling R code for speed

It is possible to profile R code on Windows and most<sup>1</sup> Unix-alike versions of R.

The command `Rprof` is used to control profiling, and its help page can be consulted for full details. Profiling works by recording at fixed intervals<sup>2</sup> (by default every 20 msecs)

<sup>1</sup> R has to be built to enable this, but the option ‘--enable-R-profiling’ is the default.

<sup>2</sup> For Unix-alikes these are intervals of CPU time, and for Windows of elapsed time.

which R function is being used, and recording the results in a file (default ‘Rprof.out’ in the working directory). Then the function `summaryRprof` or the command-line utility R CMD `Rprof Rprof.out` can be used to summarize the activity.

As an example, consider the following code (from Venables & Ripley, 2002, pp. 225–6).

```
library(MASS); library(boot)
storm.fm <- nls(Time ~ b*Viscosity/(Wt - c), stormer,
               start = c(b=30.401, c=2.2183))
st <- cbind(stormer, fit=fitted(storm.fm))
storm.bf <- function(rs, i) {
  st$Time <- st$fit + rs[i]
  tmp <- nls(Time ~ (b * Viscosity)/(Wt - c), st,
             start = coef(storm.fm))
  tmp$m$getAllPars()
}
rs <- scale(resid(storm.fm), scale = FALSE) # remove the mean
Rprof("boot.out")
storm.boot <- boot(rs, storm.bf, R = 4999) # slow enough to profile
Rprof(NULL)
```

Having run this we can summarize the results by

```
R CMD Rprof boot.out
```

```
Each sample represents 0.02 seconds.
Total run time: 22.52 seconds.
```

```
Total seconds: time spent in function and callees.
Self seconds: time spent in function alone.
```

| %     | total   | %    | self    |                       |
|-------|---------|------|---------|-----------------------|
| total | seconds | self | seconds | name                  |
| 100.0 | 25.22   | 0.2  | 0.04    | "boot"                |
| 99.8  | 25.18   | 0.6  | 0.16    | "statistic"           |
| 96.3  | 24.30   | 4.0  | 1.02    | "nls"                 |
| 33.9  | 8.56    | 2.2  | 0.56    | "<Anonymous>"         |
| 32.4  | 8.18    | 1.4  | 0.36    | "eval"                |
| 31.8  | 8.02    | 1.4  | 0.34    | ".Call"               |
| 28.6  | 7.22    | 0.0  | 0.00    | "eval.parent"         |
| 28.5  | 7.18    | 0.3  | 0.08    | "model.frame"         |
| 28.1  | 7.10    | 3.5  | 0.88    | "model.frame.default" |
| 17.4  | 4.38    | 0.7  | 0.18    | "sapply"              |
| 15.0  | 3.78    | 3.2  | 0.80    | "nlsModel"            |
| 12.5  | 3.16    | 1.8  | 0.46    | "lapply"              |
| 12.3  | 3.10    | 2.7  | 0.68    | "assign"              |
| ...   |         |      |         |                       |

| %<br>self | self<br>seconds | %<br>total | total<br>seconds | name                  |
|-----------|-----------------|------------|------------------|-----------------------|
| 5.7       | 1.44            | 7.5        | 1.88             | "inherits"            |
| 4.0       | 1.02            | 96.3       | 24.30            | "nls"                 |
| 3.6       | 0.92            | 3.6        | 0.92             | "\$"                  |
| 3.5       | 0.88            | 28.1       | 7.10             | "model.frame.default" |
| 3.2       | 0.80            | 15.0       | 3.78             | "nlsModel"            |
| 2.8       | 0.70            | 9.8        | 2.46             | "qr.coef"             |
| 2.7       | 0.68            | 12.3       | 3.10             | "assign"              |
| 2.5       | 0.64            | 2.5        | 0.64             | ".Fortran"            |
| 2.5       | 0.62            | 7.1        | 1.80             | "qr.default"          |
| 2.2       | 0.56            | 33.9       | 8.56             | "<Anonymous>"         |
| 2.1       | 0.54            | 5.9        | 1.48             | "unlist"              |
| 2.1       | 0.52            | 7.9        | 2.00             | "FUN"                 |
| ...       |                 |            |                  |                       |

(Function names are not quoted on Windows.) This often produces surprising results and can be used to identify bottlenecks or pieces of R code that could benefit from being replaced by compiled code.

Two warnings: profiling does impose a small performance penalty, and the output files can be very large if long runs are profiled at the default sampling interval.

Profiling short runs can sometimes give misleading results. R from time to time performs *garbage collection* to reclaim unused memory, and this takes an appreciable amount of time which profiling will charge to whichever function happens to provoke it. It may be useful to compare profiling code immediately after a call to `gc()` with a profiling run without a preceding call to `gc`.

More detailed analysis of the output can be achieved by the tools in the CRAN packages **proftools** and **profr**: in particular these allow call graphs to be studied.

### 3.3 Profiling R code for memory use

Measuring memory use in R code is useful either when the code takes more memory than is conveniently available or when memory allocation and copying of objects is responsible for slow code. There are three ways to profile memory use over time in R code. All three require R to have been compiled with ‘`--enable-memory-profiling`’, which is not the default, but is currently used for the Mac OS X and Windows binary distributions. All can be misleading, for different reasons.

In understanding the memory profiles it is useful to know a little more about R’s memory allocation. Looking at the results of `gc()` shows a division of memory into `Vcells` used to store the contents of vectors and `Ncells` used to store everything else, including all the administrative overhead for vectors such as type and length information. In fact the vector contents are divided into two pools. Memory for small vectors (by default 128 bytes or less) is obtained in large chunks and then parcelled out by R; memory for larger vectors is obtained directly from the operating system.

Some memory allocation is obvious in interpreted code, for example,

```
y <- x + 1
```

allocates memory for a new vector `y`. Other memory allocation is less obvious and occurs because R is forced to make good on its promise of ‘call-by-value’ argument passing. When an argument is passed to a function it is not immediately copied. Copying occurs (if



necessary) only when the argument is modified. This can lead to surprising memory use. For example, in the ‘survey’ package we have

```
print.svycoxph <- function (x, ...)
{
  print(x$survey.design, varnames = FALSE, design.summaries = FALSE,
    ...)
  x$call <- x$printcall
  NextMethod()
}
```

It may not be obvious that the assignment to `x$call` will cause the entire object `x` to be copied. This copying to preserve the call-by-value illusion is usually done by the internal C function `duplicate`.

The main reason that memory-use profiling is difficult is garbage collection. Memory is allocated at well-defined times in an R program, but is freed whenever the garbage collector happens to run.

### 3.3.1 Memory statistics from Rprof

The sampling profiler `Rprof` described in the previous section can be given the option `memory.profiling=TRUE`. It then writes out the total R memory allocation in small vectors, large vectors, and cons cells or nodes at each sampling interval. It also writes out the number of calls to the internal function `duplicate`, which is called to copy R objects. `summaryRprof` provides summaries of this information. The main reason that this can be misleading is that the memory use is attributed to the function running at the end of the sampling interval. A second reason is that garbage collection can make the amount of memory in use decrease, so a function appears to use little memory. Running under `gctorture` helps with both problems: it slows down the code to effectively increase the sampling frequency and it makes each garbage collection release a smaller amount of memory. Changing the memory limits with `mem.limits()` may also be useful, to see how the code would run under different memory conditions.

### 3.3.2 Tracking memory allocations

The second method of memory profiling uses a memory-allocation profiler, `Rprofmem()`, which writes out a stack trace to an output file every time a large vector is allocated (with a user-specified threshold for ‘large’) or a new page of memory is allocated for the R heap. Summary functions for this output are still being designed.

Running the example from the previous section with

```
> Rprofmem("boot.memprof", threshold=1000)
> storm.boot <- boot(rs, storm.bf, R = 4999)
> Rprofmem(NULL)
```

shows that apart from some initial and final work in `boot` there are no vector allocations over 1000 bytes.

### 3.3.3 Tracing copies of an object

The third method of memory profiling involves tracing copies made of a specific (presumably large) R object. Calling `tracemem` on an object marks it so that a message is printed to standard output when the object is copied *via* `duplicate` or coercion to another type, or when a new object of the same size is created in arithmetic operations. The main reason

that this can be misleading is that copying of subsets or components of an object is not tracked. It may be helpful to use `tracemem` on these components.

In the example above we can run `tracemem` on the data frame `st`

```
> tracemem(st)
[1] "<0x9abd5e0>"
> storm.boot <- boot(rs, storm.bf, R = 4)
memtrace[0x9abd5e0->0x92a6d08]: statistic boot
memtrace[0x92a6d08->0x92a6d80]: $<-.data.frame $<- statistic boot
memtrace[0x92a6d80->0x92a6df8]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x9271318]: statistic boot
memtrace[0x9271318->0x9271390]: $<-.data.frame $<- statistic boot
memtrace[0x9271390->0x9271408]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x914f558]: statistic boot
memtrace[0x914f558->0x914f5f8]: $<-.data.frame $<- statistic boot
memtrace[0x914f5f8->0x914f670]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x972cbf0]: statistic boot
memtrace[0x972cbf0->0x972cc68]: $<-.data.frame $<- statistic boot
memtrace[0x972cc68->0x972cd08]: $<-.data.frame $<- statistic boot
memtrace[0x9abd5e0->0x98ead98]: statistic boot
memtrace[0x98ead98->0x98eae10]: $<-.data.frame $<- statistic boot
memtrace[0x98eae10->0x98eae88]: $<-.data.frame $<- statistic boot
```

The object is duplicated fifteen times, three times for each of the `R+1` calls to `storm.bf`. This is surprising, since none of the duplications happen inside `nls`. Stepping through `storm.bf` in the debugger shows that all three happen in the line

```
st$Time <- st$fit + rs[i]
```

Data frames are slower than matrices and this is an example of why. Using `tracemem(st$Viscosity)` does not reveal any additional copying.

## 3.4 Profiling compiled code

Profiling compiled code is highly system-specific, but this section contains some hints gleaned from various R users. Some methods need to be different for a compiled executable and for dynamic/shared libraries/objects as used by R packages. We know of no good way to profile DLLs on Windows.

### 3.4.1 Linux

Options include using `sprof` for a shared object, and `oprofile` (see <http://oprofile.sourceforge.net/>) for any executable or shared object.

#### 3.4.1.1 sprof

You can select shared objects to be profiled with `sprof` by setting the environment variable `LD_PROFILE`. For example

```
% setenv LD_PROFILE /path/to/R_HOME/library/stats/libs/stats.so
R
... run the boot example
% sprof /path/to/R_HOME/library/stats/libs/stats.so \
    /var/tmp/path/to/R_HOME/library/stats/libs/stats.so.profile
```

Flat profile:

Each sample counts as 0.01 seconds.

| %<br>time | cumulative<br>seconds | self<br>seconds | calls | self<br>us/call | total<br>us/call | name           |
|-----------|-----------------------|-----------------|-------|-----------------|------------------|----------------|
| 76.19     | 0.32                  | 0.32            | 0     | 0.00            |                  | numeric_deriv  |
| 16.67     | 0.39                  | 0.07            | 0     | 0.00            |                  | nls_iter       |
| 7.14      | 0.42                  | 0.03            | 0     | 0.00            |                  | getListElement |

```
rm /path/to/R_HOME/library/stats/libs/stats.so.profile
... to clean up ...
```

It is possible that root access is needed to create the directories used for the profile data.

### 3.4.1.2 oprofile

oprofile works by running a daemon which collects information. The daemon must be started as root, e.g.

```
% su
% opcontrol --no-vmlinux
% (optional, some platforms) opcontrol --callgraph=5
% opcontrol --start
% exit
```

Then as a user

```
% R
... run the boot example
% opcontrol --dump
% oprofile -l /path/to/R_HOME/library/stats/libs/stats.so
...
samples %      symbol name
1623    75.5939 anonymous symbol from section .plt
349     16.2552 numeric_deriv
113     5.2632  nls_iter
62      2.8878  getListElement
% oprofile -l /path/to/R_HOME/bin/exec/R
...
samples %      symbol name
76052   11.9912 Rf_eval
54670   8.6198  Rf_findVarInFrame3
37814   5.9622  Rf_allocVector
31489   4.9649  Rf_duplicate
28221   4.4496  Rf_protect
26485   4.1759  Rf_cons
23650   3.7289  Rf_matchArgs
21088   3.3250  Rf_findFun
19995   3.1526  findVarLocInFrame
14871   2.3447  Rf_evalList
13794   2.1749  R_Newhashpjw
13522   2.1320  R_gc_internal
...
```

Shutting down the profiler and clearing the records needs to be done as root. You can use `opannotate` to annotate the source code with the times spent in each section, if the appropriate source code was compiled with debugging support, and `opreport -c` to generate a callgraph (if collection was enabled and the platform supports this).

### 3.4.2 Solaris

On 64-bit (only) Solaris, the standard profiling tool `gprof` collects information from shared objects compiled with ‘`-pg`’.

### 3.4.3 Mac OS X

Developers have recommended `sample` (or `Sampler.app`, which is a GUI version) and `Shark` (see <http://developer.apple.com/tools/sharkoptimize.html> and [http://developer.apple.com/tools/shark\\_optimize.html](http://developer.apple.com/tools/shark_optimize.html)).

## 4 Debugging

This chapter covers the debugging of R extensions, starting with the ways to get useful error information and moving on to how to deal with errors that crash R. For those who prefer other styles there are contributed packages such as **debug** on CRAN (described in an article in *R-News* 3/3). (There are notes from 2002 provided by Roger Peng at <http://www.biostat.jhsph.edu/~rpeng/docs/R-debug-tools.pdf> which provide complementary examples to those given here.)

### 4.1 Browsing

Most of the R-level debugging facilities are based around the built-in browser. This can be used directly by inserting a call to `browser()` into the code of a function (for example, using `fix(my_function)`). When code execution reaches that point in the function, control returns to the R console with a special prompt. For example

```
> fix(summary.data.frame) ## insert browser() call after for() loop
> summary(women)
Called from: summary.data.frame(women)
Browse[1]> ls()
[1] "digits" "i"      "lbs"      "lw"      "maxsum" "nm"      "nr"      "nv"
[9] "object" "sms"     "z"
Browse[1]> maxsum
[1] 7
Browse[1]>
      height      weight
Min.   :58.0   Min.   :115.0
1st Qu.:61.5   1st Qu.:124.5
Median :65.0   Median :135.0
Mean   :65.0   Mean   :136.7
3rd Qu.:68.5   3rd Qu.:148.0
Max.   :72.0   Max.   :164.0
> rm(summary.data.frame)
```

At the browser prompt one can enter any R expression, so for example `ls()` lists the objects in the current frame, and entering the name of an object will<sup>1</sup> print it. The following commands are also accepted

- **n**  
Enter ‘step-through’ mode. In this mode, hitting return executes the next line of code (more precisely one line and any continuation lines). Typing **c** will continue to the end of the current context, e.g. to the end of the current loop or function.
- **c**  
In normal mode, this quits the browser and continues execution, and just return works in the same way. **cont** is a synonym.

---

<sup>1</sup> With the exceptions of the commands listed below: an object of such a name can be printed *via* an explicit call to `print`.

- **where**

This prints the call stack. For example

```
> summary(women)
Called from: summary.data.frame(women)
Browse[1]> where
where 1: summary.data.frame(women)
where 2: summary(women)

Browse[1]>
```

- **Q**

Quit both the browser and the current expression, and return to the top-level prompt.

Errors in code executed at the browser prompt will normally return control to the browser prompt. Objects can be altered by assignment, and will keep their changed values when the browser is exited. If really necessary, objects can be assigned to the workspace from the browser prompt (by using `<<-` if the name is not already in scope).

## 4.2 Debugging R code

Suppose your R program gives an error message. The first thing to find out is what R was doing at the time of the error, and the most useful tool is `traceback()`. We suggest that this is run whenever the cause of the error is not immediately obvious. Daily, errors are reported to the R mailing lists as being in some package when `traceback()` would show that the error was being reported by some other package or base R. Here is an example from the regression suite.

```
> success <- c(13,12,11,14,14,11,13,11,12)
> failure <- c(0,0,0,0,0,0,0,2,2)
> resp <- cbind(success, failure)
> predictor <- c(0, 5^(0:7))
> glm(resp ~ 0+predictor, family = binomial(link="log"))
Error: no valid set of coefficients has been found: please supply starting values
> traceback()
3: stop("no valid set of coefficients has been found: please supply
    starting values", call. = FALSE)
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart,
    mustart = mustart, offset = offset, family = family, control = control,
    intercept = attr(mt, "intercept") > 0)
1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
```

The calls to the active frames are given in reverse order (starting with the innermost). So we see the error message comes from an explicit check in `glm.fit`. (`traceback()` shows you all the lines of the function calls, which can be limited by setting option `"deparse.max.lines"`.)

Sometimes the traceback will indicate that the error was detected inside compiled code, for example (from `?nls`)

```
Error in nls(y ~ a + b * x, start = list(a = 0.12345, b = 0.54321), trace = TRUE) :
  step factor 0.000488281 reduced below 'minFactor' of 0.000976563
> traceback()
2: .Call(R_nls_iter, m, ctrl, trace)
1: nls(y ~ a + b * x, start = list(a = 0.12345, b = 0.54321), trace = TRUE)
```

This will be the case if the innermost call is to `.C`, `.Fortran`, `.Call`, `.External` or `.Internal`, but as it is also possible for such code to evaluate R expressions, this need not be the innermost call, as in

```
> traceback()
9: gm(a, b, x)
8: .Call(R_numeric_deriv, expr, theta, rho, dir)
7: numericDeriv(form[[3]], names(ind), env)
6: getRHS()
5: assign("rhs", getRHS(), envir = thisEnv)
4: assign("resid", .swts * (lhs - assign("rhs", getRHS(), envir = thisEnv)),
  envir = thisEnv)
3: function (newPars)
  {
    setPars(newPars)
    assign("resid", .swts * (lhs - assign("rhs", getRHS(), envir = thisEnv)),
          envir = thisEnv)
    assign("dev", sum(resid^2), envir = thisEnv)
    assign("QR", qr(.swts * attr(rhs, "gradient")), envir = thisEnv)
    return(QR$rank < min(dim(QR$qr)))
  }(c(-0.00760232418963883, 1.00119632515036))
2: .Call(R_nls_iter, m, ctrl, trace)
1: nls(yeps ~ gm(a, b, x), start = list(a = 0.12345, b = 0.54321))
```

Occasionally `traceback()` does not help, and this can be the case if S4 method dispatch is involved. Consider the following example

```
> xyd <- new("xyloc", x=runif(20), y=runif(20))
Error in as.environment(pkg) : no item called "package:S4nswv"
on the search list
Error in initialize(value, ...) : S language method selection got
an error when called from internal dispatch for function 'initialize'
> traceback()
2: initialize(value, ...)
1: new("xyloc", x = runif(20), y = runif(20))
```

which does not help much, as there is no call to `as.environment` in `initialize` (and the note “called from internal dispatch” tells us so). In this case we searched the R sources for the quoted call, which occurred in only one place, `methods:::asEnvironmentPackage`. So now we knew where the error was occurring. (This was an unusually opaque example.)

The error message

```
evaluation nested too deeply: infinite recursion / options(expressions=)?
```

can be hard to handle with the default value (5000). Unless you know that there actually is deep recursion going on, it can help to set something like

```
options(expressions=500)
```

and re-run the example showing the error.

Sometimes there is warning that clearly is the precursor to some later error, but it is not obvious where it is coming from. Setting `options(warn = 2)` (which turns warnings into errors) can help here.

Once we have located the error, we have some choices. One way to proceed is to find out more about what was happening at the time of the crash by looking a *post-mortem* dump. To do so, set `options(error=dump.frames)` and run the code again. Then invoke `debugger()` and explore the dump. Continuing our example:



```
> options(error = dump.frames)
> glm(resp ~ 0 + predictor, family = binomial(link = "log"))
Error: no valid set of coefficients has been found: please supply starting values
```

which is the same as before, but an object called `last.dump` has appeared in the workspace. (Such objects can be large, so remove it when it is no longer needed.) We can examine this at a later time by calling the function `debugger`.

```
> debugger()
Message: Error: no valid set of coefficients has been found: please supply starting values
Available environments had calls:
1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart, mus
3: stop("no valid set of coefficients has been found: please supply starting values
Enter an environment number, or 0 to exit Selection:
```

which gives the same sequence of calls as `traceback`, but in outer-first order and with only the first line of the call, truncated to the current width. However, we can now examine in more detail what was happening at the time of the error. Selecting an environment opens the browser in that frame. So we select the function call which spawned the error message, and explore some of the variables (and execute two function calls).

```
Enter an environment number, or 0 to exit Selection: 2
Browsing in the environment with call:
  glm.fit(x = X, y = Y, weights = weights, start = start, etas
Called from: debugger.look(ind)
Browse[1]> ls()
 [1] "aic"      "boundary" "coefold"   "control"   "conv"
 [6] "dev"      "dev.resids" "devold"    "EMPTY"     "eta"
[11] "etastart" "family"    "fit"       "good"      "intercept"
[16] "iter"     "linkinv"   "mu"        "mu.eta"    "mu.eta.val"
[21] "mustart"  "n"         "ngoodobs"  "nobs"      "nvars"
[26] "offset"   "start"     "valideta"  "validmu"   "variance"
[31] "varmu"    "w"         "weights"   "x"         "xnames"
[36] "y"        "ynames"    "z"

Browse[1]> eta
      1      2      3      4      5
0.000000e+00 -2.235357e-06 -1.117679e-05 -5.588393e-05 -2.794197e-04
      6      7      8      9
-1.397098e-03 -6.985492e-03 -3.492746e-02 -1.746373e-01

Browse[1]> valideta(eta)
[1] TRUE

Browse[1]> mu
      1      2      3      4      5      6      7      8
1.0000000 0.9999978 0.9999888 0.9999441 0.9997206 0.9986039 0.9930389 0.9656755
      9
0.8397616

Browse[1]> validmu(mu)
[1] FALSE

Browse[1]> c
Available environments had calls:
1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart
3: stop("no valid set of coefficients has been found: please supply starting v

Enter an environment number, or 0 to exit Selection: 0
> rm(last.dump)
```

Because `last.dump` can be looked at later or even in another R session, post-mortem debugging is possible even for batch usage of R. We do need to arrange for the dump to be

saved: this can be done either using the command-line flag ‘`--save`’ to save the workspace at the end of the run, or *via* a setting such as

```
> options(error = quote({dump.frames(to.file=TRUE); q()}))
```

See the help on `dump.frames` for further options and a worked example.

An alternative error action is to use the function `recover()`:

```
> options(error = recover)
> glm(resp ~ 0 + predictor, family = binomial(link = "log"))
Error: no valid set of coefficients has been found: please supply starting values

Enter a frame number, or 0 to exit

1: glm(resp ~ 0 + predictor, family = binomial(link = "log"))
2: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart

Selection:
```

which is very similar to `dump.frames`. However, we can examine the state of the program directly, without dumping and re-loading the dump. As its help page says, `recover` can be routinely used as the error action in place of `dump.calls` and `dump.frames`, since it behaves like `dump.frames` in non-interactive use.

Post-mortem debugging is good for finding out exactly what went wrong, but not necessarily why. An alternative approach is to take a closer look at what was happening just before the error, and a good way to do that is to use `debug`. This inserts a call to the browser at the beginning of the function, starting in step-through mode. So in our example we could use

```
> debug(glm.fit)
> glm(resp ~ 0 + predictor, family = binomial(link = "log"))
debugging in: glm.fit(x = X, y = Y, weights = weights, start = start, etastart = etastart,
  mustart = mustart, offset = offset, family = family, control = control,
  intercept = attr(mt, "intercept") > 0)
debug: {
## lists the whole function
Browse[1]>
debug: x <- as.matrix(x)
...
Browse[1]> start
[1] -2.235357e-06
debug: eta <- drop(x %*% start)
Browse[1]> eta
      1      2      3      4      5
0.000000e+00 -2.235357e-06 -1.117679e-05 -5.588393e-05 -2.794197e-04
      6      7      8      9
-1.397098e-03 -6.985492e-03 -3.492746e-02 -1.746373e-01
Browse[1]>
debug: mu <- linkinv(eta <- eta + offset)
Browse[1]> mu
      1      2      3      4      5      6      7      8
1.0000000 0.9999978 0.9999888 0.9999441 0.9997206 0.9986039 0.9930389 0.9656755
      9
0.8397616
```

(The prompt `Browse[1]>` indicates that this is the first level of browsing: it is possible to step into another function that is itself being debugged or contains a call to `browser()`.)

`debug` can be used for hidden functions and S3 methods by e.g. `debug(stats:::predict.Arima)`. (It cannot be used for S4 methods, but an alternative is given on the help page for `debug`.) Sometimes you want to debug a function defined inside another function, e.g. the function `arimafn` defined inside `arima`. To do so, set `debug` on the outer function (here `arima`) and step through it until the inner function has been defined. Then call `debug` on the inner function (and use `c` to get out of step-through mode in the outer function).

To remove debugging of a function, call `undebug` with the argument previously given to `debug`; debugging otherwise lasts for the rest of the R session (or until the function is edited or otherwise replaced).

`trace` can be used to temporarily insert debugging code into a function, for example to insert a call to `browser()` just before the point of the error. To return to our running example

```
## first get a numbered listing of the expressions of the function
> page(as.list(body(glm.fit)), method="print")
> trace(glm.fit, browser, at=22)
Tracing function "glm.fit" in package "stats"
[1] "glm.fit"
> glm(resp ~ 0 + predictor, family = binomial(link = "log"))
Tracing glm.fit(x = X, y = Y, weights = weights, start = start,
  etastart = etastart, .... step 22
Called from: eval(expr, envir, enclos)
Browse[1]> n
## and single-step from here.
> untrace(glm.fit)
```

For your own functions, it may be as easy to use `fix` to insert temporary code, but `trace` can help with functions in a namespace (as can `fixInNamespace`). Alternatively, use `trace(,edit=TRUE)` to insert code visually.

## 4.3 Using `gctorture` and `valgrind`

Errors in memory allocation and reading/writing outside arrays are very common causes of crashes (e.g., segfaults) on some machines. Often the crash appears long after the invalid memory access: in particular damage to the structures which R itself has allocated may only become apparent at the next garbage collection (or even at later garbage collections after objects have been deleted).

### 4.3.1 Using `gctorture`

We can help to detect memory problems earlier by running garbage collection as often as possible. This is achieved by `gctorture(TRUE)`, which as described on its help page

Provokes garbage collection on (nearly) every memory allocation. Intended to ferret out memory protection bugs. Also makes R run *very* slowly, unfortunately.

The reference to ‘memory protection’ is to missing C-level calls to `PROTECT`/`UNPROTECT` (see [Section 5.9.1 \[Garbage Collection\]](#), [page 106](#)) which if missing allow R objects to be garbage-collected when they are still in use. But it can also help with other memory-related errors.

Normally running under `gctorture(TRUE)` will just produce a crash earlier in the R program, hopefully close to the actual cause. See the next section for how to decipher such crashes.

It is possible to run all the examples, tests and vignettes covered by R CMD `check` under `gctorture(TRUE)` by using the option `'--use-gct'`.

The function `gctorture2` provides more refined control over the GC torture process. Its arguments `step`, `wait` and `inhibit_release` are documented on its help page. Environment variables can also be used to turn on GC torture: `R_GCTORTURE` corresponds to the `step` argument to `gctorture`, `R_GCTORTURE_WAIT` to `wait`, and `R_GCTORTURE_INHIBIT_RELEASE` to `inhibit_release`.

If R is configured with `'--enable-strict-barrier'` then a variety of tests for the integrity of the write barrier are enabled. In addition tests to help detect protect issues are enabled as well:

- All GCs are full GCs.
- New nodes in small node pages are marked as `NEWSXP` on creation.
- After a GC all free nodes that are not of type `NEWSXP` are marked as type `FREESXP` and their previous type is recorded.
- Most calls to accessor functions check their `SEXP` inputs and `SEXP` outputs and signal an error if a `FREESXP` is found. The address of the node and the old type are included in the error message.

Used with a debugger and with `gctorture` or `gctorture2` this mechanism can be helpful in isolating memory protect problems.

### 4.3.2 Using valgrind

If you have access to Linux on an `'ix86'`, `'x86_64'`, `'ppc32'`, `'ppc64'` or `'s390x'` platform, or Mac OS 10.5/6/7 on `'i386'` or `'x86_64'` you can use `valgrind` (<http://www.valgrind.org/>, pronounced to rhyme with 'tinned') to check for possible problems. To run some examples under `valgrind` use something like

```
R -d valgrind --vanilla < mypkg-Ex.R
R -d "valgrind --tool=memcheck --leak-check=full" --vanilla < mypkg-Ex.R
```

where `'mypkg-Ex.R'` is a set of examples, e.g. the file created in `'mypkg.Rcheck'` by R CMD `check`. Occasionally this reports memory reads of 'uninitialised values' that are the result of compiler optimization, so can be worth checking under an unoptimized compile: for maximal information use a build with debugging symbols. We know there will be some small memory leaks from `readline` and R itself — these are memory areas that are in use right up to the end of the R session. Expect this to run around 20x slower than without `valgrind`, and in some cases even slower than that. Earlier versions (at least) of `valgrind` are not happy with many optimized BLASes that use CPU-specific instructions (3D now, SSE, SSE2, SSE3 and similar) so you may need to build a version of R specifically to use with `valgrind`.

On platforms supported by `valgrind` you can build a version of R with extra instrumentation to help `valgrind` detect errors in the use of memory allocated from the R heap. The configure option is `'--with-valgrind-instrumentation=level'`, where `level` is 0, 1, or 2. Level 0 is the default and does not add anything. Level 1 will detect use of uninitialised

memory and has little impact on speed. Level 2 will detect many other memory use bugs but makes R much slower when running under `valgrind`. Using this in conjunction with `gctorture` can be even more effective (and even slower).

An example of `valgrind` output is

```
==12539== Invalid read of size 4
==12539==    at 0x1CDF6CBE: csc_compTr (Mutils.c:273)
==12539==    by 0x1CE07E1E: tsc_transpose (dtCMatrix.c:25)
==12539==    by 0x80A67A7: do_dotcall (dotcode.c:858)
==12539==    by 0x80CACE2: Rf_eval (eval.c:400)
==12539==    by 0x80CB5AF: R_execClosure (eval.c:658)
==12539==    by 0x80CB98E: R_execMethod (eval.c:760)
==12539==    by 0x1B93DEFA: R_standardGeneric (methods_list_dispatch.c:624)
==12539==    by 0x810262E: do_standardGeneric (objects.c:1012)
==12539==    by 0x80CAD23: Rf_eval (eval.c:403)
==12539==    by 0x80CB2F0: Rf_applyClosure (eval.c:573)
==12539==    by 0x80CADCC: Rf_eval (eval.c:414)
==12539==    by 0x80CAA03: Rf_eval (eval.c:362)
==12539== Address 0x1COD2EA8 is 280 bytes inside a block of size 1996 alloc'd
==12539==    at 0x1B9008D1: malloc (vg_replace_malloc.c:149)
==12539==    by 0x80F1B34: GetNewPage (memory.c:610)
==12539==    by 0x80F7515: Rf_allocVector (memory.c:1915)
...
```

This example is from an instrumented version of R, while tracking down a bug in the **Matrix** package in January, 2006. The first line indicates that R has tried to read 4 bytes from a memory address that it does not have access to. This is followed by a C stack trace showing where the error occurred. Next is a description of the memory that was accessed. It is inside a block allocated by `malloc`, called from `GetNewPage`, that is, in the internal R heap. Since this memory all belongs to R, `valgrind` would not (and did not) detect the problem in an uninstrumented build of R. In this example the stack trace was enough to isolate and fix the bug, which was in `tsc_transpose`, and in this example running under `gctorture()` did not provide any additional information. When the stack trace is not sufficiently informative the option ‘`--db-attach=yes`’ to `valgrind` may be helpful. This starts a post-mortem debugger (by default `gdb`) so that variables in the C code can be inspected (see [Section 4.4.2 \[Inspecting R objects\]](#), page 90).

It is possible to run all the examples, tests and vignettes covered by R CMD `check` under `valgrind` by using the option ‘`--use-valgrind`’. If you do this you will need to select the `valgrind` options some other way, for example by having a ‘`~/valgrindrc`’ file containing

```
--tool=memcheck
--memcheck:leak-check=full
```

or setting the environment variable `VALGRIND_OPTS`.

On Mac OS X you may need to ensure that debugging symbols are made available (so `valgrind` reports line numbers in files). This can usually be done with the `valgrind` option ‘`--dysmutil=yes`’ to ask for the symbols to be dumped when the ‘`.so`’ file is loaded. This will not work where packages are installed into a system area (such as the ‘`R.framework`’) and can be slow. Installing packages with R CMD `INSTALL --dsym` installs the dumped symbols. (This can also be done by setting environment variable `PKG_MAKE_DSYM` to a non-empty value.)

## 4.4 Debugging compiled code

Sooner or later programmers will be faced with the need to debug compiled code loaded into R. This section is geared to platforms using `gdb` with code compiled by `gcc`, but similar things are possible with front-ends to `gdb` such as `ddd` and `insight`, and other debuggers such as Sun's `dbx`.

Consider first ‘crashes’, that is when R terminated unexpectedly with an illegal memory access (a ‘segfault’ or ‘bus error’), illegal instruction or similar. Unix-alike versions of R use a signal handler which aims to give some basic information. For example

```
*** caught segfault ***
address 0x20000028, cause 'memory not mapped'

Traceback:
 1: .identC(class1[[1]], class2)
 2: possibleExtends(class(sloti), classi, ClassDef2 = getClassDef(classi,
where = where))
 3: validObject(t(cu))
 4: stopifnot(validObject(cu <- as(tu, "dtCMatrix")), validObject(t(cu)),
validObject(t(tu)))

Possible actions:
1: abort (with core dump)
2: normal R exit
3: exit R without saving workspace
4: exit R saving workspace
Selection: 3
```

Since the R process may be damaged, the only really safe option is the first.

Another cause of a ‘crash’ is to overrun the C stack. R tries to track that in its own code, but it may happen in third-party compiled code. For modern POSIX-compliant OSes we can safely catch that and return to the top-level prompt.

```
> .C("aaa")
Error: segfault from C stack overflow
>
```

However, C stack overflows are fatal under Windows and normally defeat attempts at debugging on that platform.

If you have a crash which gives a core dump you can use something like

```
gdb /path/to/R/bin/exec/R core.12345
```

to examine the core dump. If core dumps are disabled or to catch errors that do not generate a dump one can run R directly under a debugger by for example

```
$ R -d gdb --vanilla
...
gdb> run
```

at which point R will run normally, and hopefully the debugger will catch the error and return to its prompt. This can also be used to catch infinite loops or interrupt very long-running code. For a simple example

```

> for(i in 1:1e7) x <- rnorm(100)
[hit Ctrl-C]
Program received signal SIGINT, Interrupt.
0x00397682 in _int_free () from /lib/tls/libc.so.6
(gdb) where
#0 0x00397682 in _int_free () from /lib/tls/libc.so.6
#1 0x00397eba in free () from /lib/tls/libc.so.6
#2 0xb7cf2551 in R_gc_internal (size_needed=313)
    at /users/ripley/R/svn/R-devel/src/main/memory.c:743
#3 0xb7cf3617 in Rf_allocVector (type=13, length=626)
    at /users/ripley/R/svn/R-devel/src/main/memory.c:1906
#4 0xb7c3f6d3 in PutRNGstate ()
    at /users/ripley/R/svn/R-devel/src/main/RNG.c:351
#5 0xb7d6c0a5 in do_random2 (call=0x94bf7d4, op=0x92580e8, args=0x9698f98,
    rho=0x9698f28) at /users/ripley/R/svn/R-devel/src/main/random.c:183
...

```

Some “tricks” worth knowing follow:

#### 4.4.1 Finding entry points in dynamically loaded code

Under most compilation environments, compiled code dynamically loaded into R cannot have breakpoints set within it until it is loaded. To use a symbolic debugger on such dynamically loaded code under Unix-alikes use

- Call the debugger on the R executable, for example by `R -d gdb`.
- Start R.
- At the R prompt, use `dyn.load` or `library` to load your shared object.
- Send an interrupt signal. This will put you back to the debugger prompt.
- Set the breakpoints in your code.
- Continue execution of R by typing `signal ORET`.

Under Windows signals may not be able to be used, and if so the procedure is more complicated. See the `rw-FAQ` and [www.stats.uwo.ca/faculty/murdoch/software/debuggingR/gdb.shtml](http://www.stats.uwo.ca/faculty/murdoch/software/debuggingR/gdb.shtml).

#### 4.4.2 Inspecting R objects when debugging

The key to inspecting R objects from compiled code is the function `PrintValue(SEXP s)` which uses the normal R printing mechanisms to print the R object pointed to by `s`, or the safer version `R_PV(SEXP s)` which will only print ‘objects’.

One way to make use of `PrintValue` is to insert suitable calls into the code to be debugged.

Another way is to call `R_PV` from the symbolic debugger. (`PrintValue` is hidden as `Rf_PrintValue`.) For example, from `gdb` we can use

```
(gdb) p R_PV(ab)
```

using the object `ab` from the convolution example, if we have placed a suitable breakpoint in the convolution C code.

To examine an arbitrary R object we need to work a little harder. For example, let



```
R> DF <- data.frame(a = 1:3, b = 4:6)
```

By setting a breakpoint at `do_get` and typing `get("DF")` at the R prompt, one can find out the address in memory of DF, for example

```
Value returned is $1 = (SEXP *) 0x40583e1c
(gdb) p *$1
$2 = {
  sxpinfo = {type = 19, obj = 1, named = 1, gp = 0,
    mark = 0, debug = 0, trace = 0, = 0},
  attrib = 0x40583e80,
  u = {
    vecsxp = {
      length = 2,
      type = {c = 0x40634700 "0>X@D>X@0>X@", i = 0x40634700,
        f = 0x40634700, z = 0x40634700, s = 0x40634700},
      truelength = 1075851272,
    },
    primsxp = {offset = 2},
    symsxp = {pname = 0x2, value = 0x40634700, internal = 0x40203008},
    listsxp = {carval = 0x2, cdrval = 0x40634700, tagval = 0x40203008},
    envsxp = {frame = 0x2, enclos = 0x40634700},
    closxp = {formals = 0x2, body = 0x40634700, env = 0x40203008},
    promsxp = {value = 0x2, expr = 0x40634700, env = 0x40203008}
  }
}
```

(Debugger output reformatted for better legibility).

Using `R_PV()` one can “inspect” the values of the various elements of the SEXP, for example,

```
(gdb) p R_PV($1->attrib)
$names
[1] "a" "b"

$row.names
[1] "1" "2" "3"

$class
[1] "data.frame"

$3 = void
```

To find out where exactly the corresponding information is stored, one needs to go “deeper”:

```
(gdb) set $a = $1->attrib
(gdb) p $a->u.listsxp.tagval->u.symsxp.pname->u.vecsxp.type.c
$4 = 0x405d40e8 "names"
(gdb) p $a->u.listsxp.carval->u.vecsxp.type.s[1]->u.vecsxp.type.c
$5 = 0x40634378 "b"
(gdb) p $1->u.vecsxp.type.s[0]->u.vecsxp.type.i[0]
$6 = 1
(gdb) p $1->u.vecsxp.type.s[1]->u.vecsxp.type.i[1]
$7 = 5
```

Another alternative available from R 2.13.0 on is the `R_inspect` function which shows the low-level structure of the objects recursively (addresses differ from the above as this example is created on another machine):

```
(gdb) p R_inspect($1)
@100954d18 19 VECSXP g0c2 [OBJ,NAM(2),ATT] (len=2, tl=0)
  @100954d50 13 INTSXP g0c2 [NAM(2)] (len=3, tl=0) 1,2,3
  @100954d88 13 INTSXP g0c2 [NAM(2)] (len=3, tl=0) 4,5,6
ATTRIB:
  @102a70140 02 LISTSXP g0c0 []
    TAG: @10083c478 01 SYMSXP g0c0 [MARK,NAM(2),gp=0x4000] "names"
    @100954dc0 16 STRSXP g0c2 [NAM(2)] (len=2, tl=0)
      @10099df28 09 CHARSEX g0c1 [MARK,gp=0x21] "a"
      @10095e518 09 CHARSEX g0c1 [MARK,gp=0x21] "b"
    TAG: @100859e60 01 SYMSXP g0c0 [MARK,NAM(2),gp=0x4000] "row.names"
    @102a6f868 13 INTSXP g0c1 [NAM(2)] (len=2, tl=1) -2147483648,-3
    TAG: @10083c948 01 SYMSXP g0c0 [MARK,gp=0x4000] "class"
    @102a6f838 16 STRSXP g0c1 [NAM(2)] (len=1, tl=1)
      @1008c6d48 09 CHARSEX g0c2 [MARK,gp=0x21,ATT] "data.frame"
```

In general the representation of each object follows the format:

```
@<address> <type-nr> <type-name> <gc-info> [<flags>] ...
```

For a more fine-grained control over the the depth of the recursion and the output of vectors `R_inspect3` takes additional two integer parameters: maximum depth and the maximal number of elements that will be printed for scalar vectors. The defaults in `R_inspect` are currently -1 (no limit) and 5 respectively.

## 5 System and foreign language interfaces

### 5.1 Operating system access

Access to operating system functions is *via* the R functions `system` and `system2`. The details will differ by platform (see the on-line help), and about all that can safely be assumed is that the first argument will be a string `command` that will be passed for execution (not necessarily by a shell) and the second argument to `system` will be `internal` which if true will collect the output of the command into an R character vector.

The function `system.time` is available for timing. Timing on child processes is only available on Unix-alikes, and may not be reliable there.

### 5.2 Interface functions `.C` and `.Fortran`

These two functions provide an interface to compiled code that has been linked into R, either at build time or *via* `dyn.load` (see [Section 5.3 \[dyn.load and dyn.unload\]](#), page 95). They are primarily intended for compiled C and FORTRAN 77 code respectively, but the `.C` function can be used with other languages which can generate C interfaces, for example C++ (see [Section 5.6 \[Interfacing C++ code\]](#), page 101).

The first argument to each function is a character string specifying the symbol name as known<sup>1</sup> to C or FORTRAN, that is the function or subroutine name. (That the symbol is loaded can be tested by, for example, `is.loaded("cg")`. Use the name you pass to `.C` or `.Fortran` rather than the translated symbol name.)

There can be up to 65 further arguments giving R objects to be passed to compiled code. Normally these are copied before being passed in, and copied again to an R list object when the compiled code returns. If the arguments are given names, these are used as names for the components in the returned list object (but not passed to the compiled code).

The following table gives the mapping between the modes of R atomic vectors and the types of arguments to a C function or FORTRAN subroutine.

| R storage mode | C type          | FORTRAN type     |
|----------------|-----------------|------------------|
| logical        | int *           | INTEGER          |
| integer        | int *           | INTEGER          |
| double         | double *        | DOUBLE PRECISION |
| complex        | Rcomplex *      | DOUBLE COMPLEX   |
| character      | char **         | CHARACTER*255    |
| raw            | unsigned char * | none             |

Do please note the first two. On the 64-bit Unix/Linux/OS X platforms, `long` is 64-bit whereas `int` and `INTEGER` are 32-bit. Code ported from S-PLUS (which uses `long *` for `logical` and `integer`) will not work on all 64-bit platforms (although it may appear to work on some, including Windows). Note also that if your compiled code is a mixture of C functions and FORTRAN subprograms the argument types must match as given in the table above.

---

<sup>1</sup> possibly after some platform-specific translation, e.g. adding leading or trailing underscores.

C type `Rcomplex` is a structure with `double` members `r` and `i` defined in the header file `'R_ext/Complex.h'` included by `'R.h'`. (On most platforms this is stored in a way compatible with the C99 `double complex` type: however, it may not be possible to pass `Rcomplex` to a C99 function expecting a `double complex` argument. Nor need it be compatible with a C++ `complex` type. Moreover, the compatibility can depend on the optimization level set for the compiler.)

Only a single character string can be passed to or from FORTRAN, and the success of this is compiler-dependent. Other R objects can be passed to `.C`, but it is much better to use one of the other interfaces.

It is possible to pass numeric vectors of storage mode `double` to C as `float *` or to FORTRAN as `REAL` by setting the attribute `Csingle`, most conveniently by using the R functions `as.single`, `single` or `mode`. This is intended only to be used to aid interfacing existing C or FORTRAN code.

Logical values are sent as 0 (`FALSE`), 1 (`TRUE`) or `INT_MIN = -2147483648` (`NA`, but only if `NAOK` is true), and the compiled code should return one of these three values. (Non-zero values other than `INT_MIN` are mapped to `TRUE`.)

Unless formal argument `NAOK` is true, all the other arguments are checked for missing values `NA` and for the IEEE special values `NaN`, `Inf` and `-Inf`, and the presence of any of these generates an error. If it is true, these values are passed unchecked.

Argument `DUP` can be used to suppress copying. It is dangerous: see the on-line help for arguments against its use. It is not possible to pass numeric vectors as `float *` or `REAL` if `DUP = FALSE`, and character vectors cannot be used.

Argument `PACKAGE` confines the search for the symbol name to a specific shared object (or use `"base"` for code compiled into R). Its use is highly desirable, as there is no way to avoid two package writers using the same symbol name, and such name clashes are normally sufficient to cause R to crash. (If it is not present and the call is from the body of a function defined in a package namespace, the shared object loaded by the first (if any) `useDynLib` directive will be used.)

Note that the compiled code should not return anything except through its arguments: C functions should be of type `void` and FORTRAN subprograms should be subroutines.

To fix ideas, let us consider a very simple example which convolves two finite sequences. (This is hard to do fast in interpreted R code, but easy in C code.) We could do this using `.C` by

```
void convolve(double *a, int *na, double *b, int *nb, double *ab)
{
    R_len_t i, j, nab = *na + *nb - 1;

    for(i = 0; i < nab; i++)
        ab[i] = 0.0;
    for(i = 0; i < *na; i++)
        for(j = 0; j < *nb; j++)
            ab[i + j] += a[i] * b[j];
}
```

called from R by

```
conv <- function(a, b)
  .C("convolve",
    as.double(a),
    as.integer(length(a)),
    as.double(b),
    as.integer(length(b)),
    ab = double(length(a) + length(b) - 1))$ab
```

Note that we take care to coerce all the arguments to the correct R storage mode before calling `.C`; mistakes in matching the types can lead to wrong results or hard-to-catch errors.

Special care is needed in handling **character** vector arguments in C (or C++). Since only `DUP = TRUE` is allowed, on entry the contents of the elements are duplicated and assigned to the elements of a `char **` array, and on exit the elements of the C array are copied to create new elements of a character vector. This means that the contents of the character strings of the `char **` array can be changed, including to `\0` to shorten the string, but the strings cannot be lengthened. It is possible to allocate a new string *via* `R_alloc` and replace an entry in the `char **` array by the new string. However, when character vectors are used other than in a read-only way, the `.Call` interface is much to be preferred.

Passing character strings to FORTRAN code needs even more care, and should be avoided where possible. Only the first element of the character vector is passed in, as a fixed-length (255) character array. Up to 255 characters are passed back to a length-one character vector. How well this works (or even if it works at all) depends on the C and FORTRAN compilers on each platform.

### 5.3 `dyn.load` and `dyn.unload`

Compiled code to be used with R is loaded as a shared object (Unix-alikes including Mac OS X, see [Section 5.5 \[Creating shared objects\]](#), page 99 for more information) or DLL (Windows).

The shared object/DLL is loaded by `dyn.load` and unloaded by `dyn.unload`. Unloading is not normally necessary, but it is needed to allow the DLL to be re-built on some platforms, including Windows.

The first argument to both functions is a character string giving the path to the object. Programmers should not assume a specific file extension for the object/DLL (such as `‘.so’`) but use a construction like

```
file.path(path1, path2, paste0("mylib", .Platform$dynlib.ext))
```

for platform independence. On Unix-alike systems the path supplied to `dyn.load` can be an absolute path, one relative to the current directory or, if it starts with `‘~’`, relative to the user’s home directory.

Loading is most often done automatically based on the `useDynLib()` declaration in the `‘NAMESPACE’` file, but may be done explicitly *via* a call to `library.dynam`. This has the form

```
library.dynam("libname", package, lib.loc)
```

where `libname` is the object/DLL name *with the extension omitted*. Note that the first argument, `chname`, should **not** be `package` since this will not work if the package is installed under another name.

Under some Unix-alike systems there is a choice of how the symbols are resolved when the object is loaded, governed by the arguments `local` and `now`. Only use these if really necessary: in particular using `now=FALSE` and then calling an unresolved symbol will terminate R unceremoniously.

R provides a way of executing some code automatically when a object/DLL is either loaded or unloaded. This can be used, for example, to register native routines with R's dynamic symbol mechanism, initialize some data in the native code, or initialize a third party library. On loading a DLL, R will look for a routine within that DLL named `R_init_lib` where `lib` is the name of the DLL file with the extension removed. For example, in the command

```
library.dynam("mylib", package, lib.loc)
```

R looks for the symbol named `R_init_mylib`. Similarly, when unloading the object, R looks for a routine named `R_unload_lib`, e.g., `R_unload_mylib`. In either case, if the routine is present, R will invoke it and pass it a single argument describing the DLL. This is a value of type `DllInfo` which is defined in the `'Rdynload.h'` file in the `'R_ext'` directory.

Note that there are some implicit restrictions on this mechanism as the basename of the DLL needs to be both a valid file name and valid as part of a C entry point (e.g. it cannot contain `'.'`): for portable code it is best to confine DLL names to be ASCII alphanumeric plus underscore. As from R 2.15.0, if entry point `R_init_lib` is not found it is also looked for with `'.'` replaced by `'_'`.

The following example shows templates for the initialization and unload routines for the `mylib` DLL.

```
#include <R.h>
#include <Rinternals.h>
#include <R_ext/Rdynload.h>

void
R_init_mylib(DllInfo *info)
{
    /* Register routines,
       allocate resources. */
}

void
R_unload_mylib(DllInfo *info)
{
    /* Release resources. */
}
```

If a shared object/DLL is loaded more than once the most recent version is used. More generally, if the same symbol name appears in several shared objects, the most recently loaded occurrence is used. The `PACKAGE` argument and registration (see the next section) provide good ways to avoid any ambiguity in which occurrence is meant.

On Unix-alikes the paths used to resolve dynamically linked dependent libraries are fixed (for security reasons) when the process is launched, so `dyn.load` will only look for such libraries in the locations set by the ‘R’ shell script (*via* ‘`etc/ldpaths`’) and in the OS-specific defaults.

Windows allows more control (and less security) over where dependent DLLs are looked for. On all versions this includes the `PATH` environment variable, but with lowest priority: note that it does not include the directory from which the DLL was loaded. It is possible to add a single path with quite high priority *via* the `DLLpath` argument to `dyn.load`. This is (by default) used by `library.dynam` to include the package’s ‘`libs/i386`’ or ‘`libs/x64`’ directory in the DLL search path.

## 5.4 Registering native routines

By ‘native’ routine, we mean an entry point in compiled code.

In calls to `.C`, `.Call`, `.Fortran` and `.External`, R must locate the specified native routine by looking in the appropriate shared object/DLL. By default, R uses the operating system-specific dynamic loader to lookup the symbol in all loaded DLLs and elsewhere. Alternatively, the author of the DLL can explicitly register routines with R and use a single, platform-independent mechanism for finding the routines in the DLL. One can use this registration mechanism to provide additional information about a routine, including the number and type of the arguments, and also make it available to R programmers under a different name. In the future, registration may be used to implement a form of “secure” or limited native access.

To register routines with R, one calls the C routine `R_registerRoutines`. This is typically done when the DLL is first loaded within the initialization routine `R_init_dll` *name* described in [Section 5.3 \[dyn.load and dyn.unload\], page 95](#). `R_registerRoutines` takes 5 arguments. The first is the `DllInfo` object passed by R to the initialization routine. This is where R stores the information about the methods. The remaining 4 arguments are arrays describing the routines for each of the 4 different interfaces: `.C`, `.Call`, `.Fortran` and `.External`. Each argument is a NULL-terminated array of the element types given in the following table:

|                        |                                  |
|------------------------|----------------------------------|
| <code>.C</code>        | <code>R_CMethodDef</code>        |
| <code>.Call</code>     | <code>R_CallMethodDef</code>     |
| <code>.Fortran</code>  | <code>R_FortranMethodDef</code>  |
| <code>.External</code> | <code>R_ExternalMethodDef</code> |

Currently, the `R_ExternalMethodDef` is the same as `R_CallMethodDef` type and contains fields for the name of the routine by which it can be accessed in R, a pointer to the actual native symbol (i.e., the routine itself), and the number of arguments the routine expects. For routines with a variable number of arguments invoked *via* the `.External` interface, one specifies `-1` for the number of arguments which tells R not to check the actual number passed. For example, if we had a routine named `myCall` defined as

```
SEXP myCall(SEXP a, SEXP b, SEXP c);
```

we would describe this as

```
R_CallMethodDef callMethods[] = {
    {"myCall", (DL_FUNC) &myCall, 3},
```



```

    {NULL, NULL, 0}
};

```

along with any other routines for the `.Call` interface.

Routines for use with the `.C` and `.Fortran` interfaces are described with similar data structures, but which have two additional fields for describing the type and “style” of each argument. Each of these can be omitted. However, if specified, each should be an array with the same number of elements as the number of parameters for the routine. The types array should contain the `SEXP` types describing the expected type of the argument. (Technically, the elements of the types array are of type `R_NativePrimitiveArgType` which is just an unsigned integer.) The R types and corresponding type identifiers are provided in the following table:

|                        |                        |
|------------------------|------------------------|
| <code>numeric</code>   | <code>REALSXP</code>   |
| <code>integer</code>   | <code>INTSXP</code>    |
| <code>logical</code>   | <code>LGLSXP</code>    |
| <code>single</code>    | <code>SINGLESXP</code> |
| <code>character</code> | <code>STRSXP</code>    |
| <code>list</code>      | <code>VECSXP</code>    |

Consider a C routine, `myC`, declared as

```
void myC(double *x, int *n, char **names, int *status);
```

We would register it as

```

R_CMethodDef cMethods[] = {
    {"myC", (DL_FUNC) &myC, 4, {REALSXP, INTSXP, STRSXP, LGLSXP}},
    {NULL, NULL, 0}
};

```

One can also specify whether each argument is used simply as input, or as output, or as both input and output. The style field in the description of a method is used for this. The purpose is to allow<sup>2</sup> R to transfer values more efficiently across the R-C/FORTRAN interface by avoiding copying values when it is not necessary. Typically, one omits this information in the registration data.

Having created the arrays describing each routine, the last step is to actually register them with R. We do this by calling `R_registerRoutines`. For example, if we have the descriptions above for the routines accessed by the `.C` and `.Call` we would use the following code:

```

void
R_init_myLib(DllInfo *info)
{
    R_registerRoutines(info, cMethods, callMethods, NULL, NULL);
}

```

This routine will be invoked when R loads the shared object/DLL named `myLib`. The last two arguments in the call to `R_registerRoutines` are for the routines accessed by `.Fortran` and `.External` interfaces. In our example, these are given as `NULL` since we have no routines of these types.

---

<sup>2</sup> but this is not currently done.

When R unloads a shared object/DLL, its registrations are automatically removed. There is no other facility for unregistering a symbol.

Examples of registering routines can be found in the different packages in the R source tree (e.g., **stats**). Also, there is a brief, high-level introduction in *R News* (volume 1/3, September 2001, pages 20–23, [http://www.r-project.org/doc/Rnews/Rnews\\_2001-3.pdf](http://www.r-project.org/doc/Rnews/Rnews_2001-3.pdf)).

In addition to registering C routines to be called by R, it can at times be useful for one package to make some of its C routines available to be called by C code in another package. An interface to support this has been provided since R 2.4.0. The interface consists of two routines declared as

```
void R_RegisterCCallable(const char *package, const char *name,
                        DL_FUNC fptr);
DL_FUNC R_GetCCallable(const char *package, const char *name);
```

A package **packA** that wants to make a C routine **myCfun** available to C code in other packages would include the call

```
R_RegisterCCallable("packA", "myCfun", myCfun);
```

in its initialization function **R\_init\_packA**. A package **packB** that wants to use this routine would retrieve the function pointer with a call of the form

```
p_myCfun = R_GetCCallable("packA", "myCfun");
```

The author of **packB** is responsible for ensuring that **p\_myCfun** has an appropriate declaration. In the future R may provide some automated tools to simplify exporting larger numbers of routines.

A package that wishes to make use of header files in other packages needs to declare them as a comma-separated list in the field **LinkingTo** in the ‘DESCRIPTION’ file. For example

```
Depends: link2, link3
LinkingTo: link2, link3
```

It must also ‘Depends’ on those packages for they have to be installed prior to this one, and loaded prior to this one (so the path to their compiled code can be found).

This then arranges that the ‘include’ directories in the installed linked-to packages are added to the include paths for C and C++ code.

A CRAN example of the use of this mechanism is package **lme4**, which links to **Matrix**.

## 5.5 Creating shared objects

Shared objects for loading into R can be created using R CMD SHLIB. This accepts as arguments a list of files which must be object files (with extension ‘.o’) or sources for C, C++, FORTRAN 77, Fortran 9x, Objective C or Objective C++ (with extensions ‘.c’, ‘.cc’ or ‘.cpp’, ‘.f’, ‘.f90’ or ‘.f95’, ‘.m’, and ‘.mm’ or ‘.M’, respectively), or commands to be passed to the linker. See *R CMD SHLIB --help* (or the R help for SHLIB) for usage information.

If compiling the source files does not work “out of the box”, you can specify additional flags by setting some of the variables **PKG\_CPPFLAGS** (for the C preprocessor, typically ‘-I’ flags), **PKG\_CFLAGS**, **PKG\_CXXFLAGS**, **PKG\_FFLAGS**, **PKG\_FCFLAGS**, and **PKG\_OBJCFLAGS** (for the C, C++, FORTRAN 77, Fortran 9x, and Objective C compilers, respectively) in the file ‘**Makevars**’ in the compilation directory (or, of course, create the object files directly from

the command line). Similarly, variable `PKG_LIBS` in `'Makevars'` can be used for additional `'-l'` and `'-L'` flags to be passed to the linker when building the shared object. (Supplying linker commands as arguments to `R CMD SHLIB` will take precedence over `PKG_LIBS` in `'Makevars'`.)

It is possible to arrange to include compiled code from other languages by setting the macro `'OBJECTS'` in file `'Makevars'`, together with suitable rules to make the objects.

Flags which are already set (for example in file `'etcR_ARCH/Makeconf'`) can be overridden by the environment variable `MAKEFLAGS` (at least for systems using a POSIX-compliant `make`), as in (Bourne shell syntax)

```
MAKEFLAGS="CFLAGS=-O3" R CMD SHLIB *.c
```

It is also possible to set such variables in personal `'Makevars'` files, which are read after the local `'Makevars'` and the system makefiles or in a site-wide `'Makevars.site'` file. See [Section “Customizing package compilation” in \*R Installation and Administration\*](#),

Note that as `R CMD SHLIB` uses `Make`, it will not remake a shared object just because the flags have changed, and if `'test.c'` and `'test.f'` both exist in the current directory

```
R CMD SHLIB test.f
```

will compile `'test.c'`!

If the `'src'` subdirectory of an add-on package contains source code with one of the extensions listed above or a file `'Makevars'` but **not** a file `Makefile`, `R CMD INSTALL` creates a shared object (for loading into R through `useDynlib()` in the `'NAMESPACE'`, or in the `.onLoad` function of the package) using the `R CMD SHLIB` mechanism. If file `'Makevars'` exists it is read first, then the system makefile and then any personal `'Makevars'` files.

If the `'src'` subdirectory of package contains a file `'Makefile'`, this is used by `R CMD INSTALL` in place of the `R CMD SHLIB` mechanism. `make` is called with makefiles `'R_HOME/etcR_ARCH/Makeconf'`, `'src/Makefile'` and any personal `'Makevars'` files (in that order). The first target found in `'src/Makefile'` is used.

It is better to make use of a `Makevars` file rather than a `Makefile`: the latter should be needed only exceptionally.

Under Windows the same commands work, but `'Makevars.win'` will be used in preference to `'Makevars'`, and only `'src/Makefile.win'` will be used by `R CMD INSTALL` with `'src/Makefile'` being ignored. For past experiences of building DLLs with a variety of compilers, see file `'README.packages'` and <http://www.stats.uwo.ca/faculty/murdoch/software/compilingDLLs/>. Under Windows you can supply an exports definitions file called `'dllname-win.def'`: otherwise all entry points in objects (but not libraries) supplied to `R CMD SHLIB` will be exported from the DLL. An example is `'stats-win.def'` for the `stats` package: a CRAN example in package **fastICA**.

If you feel tempted to read the source code and subvert these mechanisms, please resist. Far too much developer time has been wasted in chasing down errors caused by failures to follow this documentation, and even more by package authors demanding explanations as to why their packages not longer work. In particular, undocumented environment or make variables are not for use by package writers and are subject to change without notice.

## 5.6 Interfacing C++ code

Suppose we have the following hypothetical C++ library, consisting of the two files ‘X.h’ and ‘X.cpp’, and implementing the two classes X and Y which we want to use in R.

```
// X.h

class X {
public: X (); ~X ();
};

class Y {
public: Y (); ~Y ();
};
```

```
// X.cpp

#include <R.h>
#include "X.h"

static Y y;

X::X() { Rprintf("constructor X\n"); }
X::~X() { Rprintf("destructor X\n"); }
Y::Y() { Rprintf("constructor Y\n"); }
Y::~Y() { Rprintf("destructor Y\n"); }
```

To use with R, the only thing we have to do is writing a wrapper function and ensuring that the function is enclosed in

```
extern "C" {

}
```

For example,

```
// X_main.cpp:

#include "X.h"

extern "C" {

void X_main () {
    X x;
}

} // extern "C"
```

Compiling and linking should be done with the C++ compiler-linker (rather than the C compiler-linker or the linker itself); otherwise, the C++ initialization code (and hence the constructor of the static variable Y) are not called. On a properly configured system, one can simply use

```
R CMD SHLIB X.cpp X_main.cpp
```

to create the shared object, typically ‘X.so’ (the file name extension may be different on your platform). Now starting R yields

```
R version 2.14.1 Patched (2012-01-16 r58124)
Copyright (C) 2012 The R Foundation for Statistical Computing
...
Type      "q()" to quit R.

R> dyn.load(paste("X", .Platform$dynlib.ext, sep = ""))
constructor Y
R> .C("X_main")
constructor X
destructor X
list()
R> q()
Save workspace image? [y/n/c]: y
destructor Y
```

The R for Windows FAQ (‘rw-FAQ’) contains details of how to compile this example under Windows.

Earlier version of this example used C++ iostreams: this is best avoided. There is no guarantee that the output will appear in the R console, and indeed it will not on the R for Windows console. Use R code or the C entry points (see [Section 6.5 \[Printing\], page 133](#)) for all I/O if at all possible. Examples have been seen where merely loading a DLL that contained calls to C++ I/O upset R’s own C I/O (for example by resetting buffers on open files).

Most R header files can be included within C++ programs, and they should **not** be included within an `extern "C"` block (as they include C++ system headers). It may not be possible to include some R headers as they in turn include C header files that may cause conflicts—if this happens, define ‘NO\_C\_HEADERS’ before including the R headers, and

include C++ versions (such as ‘`cmath`’) of the appropriate headers yourself before the R headers.

## 5.7 Fortran I/O

We have already warned against the use of C++ iostreams not least because output is not guaranteed to appear on the R console, and this warning applies equally to Fortran (77 or 9x) output to units `*` and `6`. See [Section 6.5.1 \[Printing from FORTRAN\]](#), page 134, which describes workarounds.

In the past most Fortran compilers implemented I/O on top of the C I/O system and so the two interworked successfully. This was true of `g77`, but it is less true of `gfortran` as used in `gcc 4.y.z`. In particular, any package that makes use of Fortran I/O will when compiled on Windows interfere with C I/O: when the Fortran I/O is initialized (typically when the package is loaded) the C `stdout` and `stderr` are switched to LF line endings. (Function `La_Init` in file ‘`src/main/lapack.c`’ shows how to mitigate this.) Even worse, prior to R 2.6.2 using Fortran output when running under the Windows GUI console (`Rgui`) would hang the R session. This is now avoided by ensuring that the Fortran output is written to a file (‘`fort.6`’ in the working directory).

## 5.8 Linking to other packages

It is not in general possible to link a DLL in package **packA** to a DLL provided by package **packB** (for the security reasons mentioned in [Section 5.3 \[dyn.load and dyn.unload\]](#), page 95, and also because some platforms distinguish between shared objects and dynamic libraries), but it is on Windows.

Note that there can be tricky versioning issues here, as package **packB** could be re-installed after package **packA** — it is desirable that the API provided by package **packB** remains backwards-compatible.

### 5.8.1 Unix-alikes

It is possible to link a shared object in package **packA** to a library provided by package **packB** under limited circumstances on a Unix-alike OS. There are severe portability issues, so this is not recommended for a distributed package.

This is easiest if **packB** provides a static library ‘`packB/libs/libpackB.a`’. (This will need to be compiled with PIC flags on platforms where it matters.) Then as the code from package **packB** is incorporated when package **packA** is installed, we only need to find the static library at install time for package **packB**. The only issue is to find package **packB**, and for that we can ask R by something like

```
PKGB_PATH='echo `cat(system.file("libs", .Platform$r_arch, package="packB", mustWork=T)
| "${R_HOME}/bin/R" --vanilla --slave`
PKG_LIBS="${PKGB_PATH}/libpackB.a"
```

which will give an empty path component if sub-architectures are not in use (but that works on current platforms).

For a dynamic library ‘`packB/libs/libpackB.so`’ (‘`packB/libs/libpackB.dylib`’ on Mac OS X) we could use

```
PKGB_PATH='echo `cat(system.file("libs", .Platform$r_arch, package="packB", mustWork=T)
```

```
| "${R_HOME}/bin/R" --vanilla --slave'
PKG_LIBS=-L"${PKG_PATH}" -lpackB
```

This will work for installation, but very likely not when package **packB** is loaded, as the path to package **packB**'s 'libs' directory is not in the `ld.so`<sup>3</sup> search path. You can arrange to put it there **before** R is launched by setting (on some platforms) `LD_RUN_PATH` or `LD_LIBRARY_PATH` or adding to the `ld.so` cache (see `man ldconfig`). On platforms that support it, the path to the dynamic library can be hardcoded at install time (which assumes that the location of package **packB** will not be changed) nor the package updated to a changed API). On systems with the GNU linker (e.g. Linux) and some others (e.g. Mac OS X) this can be done by

```
PKG_PATH='echo `library(packB); cat(system.file("libs", package="packB"))` \
| "${R_HOME}/bin/R" --vanilla --slave'
PKG_LIBS=-L"${PKG_PATH}" -rpath "${PKG_PATH}" -lpackB
```

and on some other systems (e.g. Solaris with its native linker) use `-R` rather than `-rpath`.

It may be possible to figure out what is required semi-automatically from the result of `R CMD libtool --config` (look for 'hardcode').

Making headers provided by package **packB** available to the code to be compiled in package **packA** can be done by the `LinkingTo` mechanism (see [Section 5.4 \[Registering native routines\]](#), page 97).

## 5.8.2 Windows

Suppose package **packA** wants to make use of compiled code provided by **packB** in DLL 'packB/libs/exB.dll', possibly the package's DLL 'packB/libs/packB.dll'. (This can be extended to linking to more than one package in a similar way.) There are three issues to be addressed:

- Making headers provided by package **packB** available to the code to be compiled in package **packA**.

This is done by the `LinkingTo` mechanism (see [Section 5.4 \[Registering native routines\]](#), page 97).

- preparing **packA.dll** to link to 'packB/libs/exB.dll'.

This needs an entry in 'Makevars.win' of the form

```
PKG_LIBS= -L<something> -lexB
```

and one possibility is that `<something>` is the path to the installed 'packB/libs' directory. To find that we need to ask R where it is by something like

```
PKG_PATH='echo `library(packB); cat(system.file("libs", package="packB"))` \
| rterm --vanilla --slave'
PKG_LIBS= -L"${PKG_PATH}" -lexB
```

Another possibility is to use an import library, shipping with package **packA** an exports file 'exB.def'. Then 'Makevars.win' could contain

```
PKG_LIBS= -L. -lexB
```

```
all: $(SHLIB) before
```

---

<sup>3</sup> `dyld` on Mac OS X, and `DYLD_LIBRARY_PATHS` below.



```
before: libexB.dll.a
libexB.dll.a: exB.def
```

and then installing package **packA** will make and use the import library for ‘exB.dll’. (One way to prepare the exports file is to use ‘pexports.exe’.)

- loading ‘packA.dll’ which depends on ‘exB.dll’.

If exB.dll was used by package **packB** (because it is in fact ‘packB.dll’ or ‘packB.dll’ depends on it) and **packB** has been loaded before **packA**, then nothing more needs to be done as ‘exB.dll’ will already be loaded into the R executable. (This is the most common scenario).

More generally, we can use the `DLLpath` argument to `library.dynam` to ensure that exB.dll is found, for example by setting

```
library.dynam("packA", pkg, lib,
              DLLpath = system.file("libs", package="packB"))
```

Note that `DLLpath` can only set one path, and so for linking to two or more packages you would need to resort to setting `PATH`.

## 5.9 Handling R objects in C

Using C code to speed up the execution of an R function is often very fruitful. Traditionally this has been done *via* the `.C` function in R. However, if a user wants to write C code using internal R data structures, then that can be done using the `.Call` and `.External` functions. The syntax for the calling function in R in each case is similar to that of `.C`, but the two functions have different C interfaces. Generally the `.Call` interface (which is modelled on the interface of the same name in S version 4) is a little simpler to use, but `.External` is a little more general.

A call to `.Call` is very similar to `.C`, for example

```
.Call("convolve2", a, b)
```

The first argument should be a character string giving a C symbol name of code that has already been loaded into R. Up to 65 R objects can be passed as arguments. The C side of the interface is

```
#include <R.h>
#include <Rinternals.h>

SEXP convolve2(SEXP a, SEXP b)
...

```

A call to `.External` is almost identical

```
.External("convolveE", a, b)
```

but the C side of the interface is different, having only one argument

```
#include <R.h>
#include <Rinternals.h>

SEXP convolveE(SEXP args)
...

```

Here `args` is a `LISTSXP`, a Lisp-style pairlist from which the arguments can be extracted.

In each case the R objects are available for manipulation *via* a set of functions and macros defined in the header file ‘`Rinternals.h`’ or some S4-compatibility macros defined in ‘`Rdefines.h`’. See [Section 5.10 \[Interface functions `.Call` and `.External`\]](#), [page 116](#) for details on `.Call` and `.External`.

Before you decide to use `.Call` or `.External`, you should look at other alternatives. First, consider working in interpreted R code; if this is fast enough, this is normally the best option. You should also see if using `.C` is enough. If the task to be performed in C is simple enough involving only atomic vectors and requiring no call to R, `.C` suffices. The new interfaces are relatively recent additions to S and R, and a great deal of useful code has been written using just `.C` before they were available. The `.Call` and `.External` interfaces allow much more control, but they also impose much greater responsibilities so need to be used with care. Neither `.Call` nor `.External` copy their arguments: you should treat arguments you receive through these interfaces as read-only.

There are two approaches that can be taken to handling R objects from within C code. The first (historically) is to use the macros and functions that have been used to implement the core parts of R through `.Internal` calls. A public<sup>4</sup> subset of these is defined in the header file ‘`Rinternals.h`’ in the directory ‘`R_INCLUDE_DIR`’ (default ‘`R_HOME/include`’) that should be available on any R installation.

Another approach is to use R versions of the macros and functions defined for the S version 4 interface `.Call`, which are defined in the header file ‘`Rdefines.h`’. This is a somewhat simpler approach, and is to be preferred if the code is intended to be shared with S. However, it is less well documented and even less tested. Note too that some idiomatic S4 constructions with these macros (such as assigning elements of character vectors or lists) are invalid in R.

A substantial amount of R is implemented using the functions and macros described here, so the R source code provides a rich source of examples and “how to do it”: do make use of the source code for inspirational examples.

It is necessary to know something about how R objects are handled in C code. All the R objects you will deal with will be handled with the type `SEXP`<sup>5</sup>, which is a pointer to a structure with typedef `SEXP`. Think of this structure as a *variant type* that can handle all the usual types of R objects, that is vectors of various modes, functions, environments, language objects and so on. The details are given later in this section and in [Section “R Internal Structures” in \*R Internals\*](#), but for most purposes the programmer does not need to know them. Think rather of a model such as that used by Visual Basic, in which R objects are handed around in C code (as they are in interpreted R code) as the variant type, and the appropriate part is extracted for, for example, numerical calculations, only when it is needed. As in interpreted R code, much use is made of coercion to force the variant object to the right type.

### 5.9.1 Handling the effects of garbage collection

We need to know a little about the way R handles memory allocation. The memory allocated for R objects is not freed by the user; instead, the memory is from time to time *garbage*

<sup>4</sup> see [Chapter 6 \[The R API\]](#), [page 130](#): note that these are not all part of the API.

<sup>5</sup> `SEXP` is an acronym for *Simple EXpression*, common in LISP-like language syntaxes.

*collected*. That is, some or all of the allocated memory not being used is freed or marked as re-usable.

The R object types are represented by a C structure defined by a typedef `SEXP` in ‘`Rinternals.h`’. It contains several things among which are pointers to data blocks and to other `SEXPREC`s. A `SEXP` is simply a pointer to a `SEXP`.

If you create an R object in your C code, you must tell R that you are using the object by using the `PROTECT` macro on a pointer to the object. This tells R that the object is in use so it is not destroyed during garbage collection. Notice that it is the object which is protected, not the pointer variable. It is a common mistake to believe that if you invoked `PROTECT(p)` at some point then *p* is protected from then on, but that is not true once a new object is assigned to *p*.

Protecting an R object automatically protects all the R objects pointed to in the corresponding `SEXP`, for example all elements of a protected list are automatically protected.

The programmer is solely responsible for housekeeping the calls to `PROTECT`. There is a corresponding macro `UNPROTECT` that takes as argument an `int` giving the number of objects to unprotect when they are no longer needed. The protection mechanism is stack-based, so `UNPROTECT(n)` unprotects the last *n* objects which were protected. The calls to `PROTECT` and `UNPROTECT` must balance when the user’s code returns. R will warn about “stack imbalance in .Call” (or `.External`) if the housekeeping is wrong.

Here is a small example of creating an R numeric vector in C code. First we use the macros in ‘`Rinternals.h`’:

```
#include <R.h>
#include <Rinternals.h>

SEXP ab;
....
PROTECT(ab = allocVector(REALSXP, 2));
REAL(ab)[0] = 123.45;
REAL(ab)[1] = 67.89;
UNPROTECT(1);
```

and then those in ‘`Rdefines.h`’:

```
#include <R.h>
#include <Rdefines.h>

SEXP ab;
....
PROTECT(ab = NEW_NUMERIC(2));
NUMERIC_POINTER(ab)[0] = 123.45;
NUMERIC_POINTER(ab)[1] = 67.89;
UNPROTECT(1);
```

Now, the reader may ask how the R object could possibly get removed during those manipulations, as it is just our C code that is running. As it happens, we can do without the protection in this example, but in general we do not know (nor want to know) what is hiding behind the R macros and functions we use, and any of them might cause memory to be allocated, hence garbage collection and hence our object `ab` to be removed. It is usually

wise to err on the side of caution and assume that any of the R macros and functions might remove the object.

In some cases it is necessary to keep better track of whether protection is really needed. Be particularly aware of situations where a large number of objects are generated. The pointer protection stack has a fixed size (default 10,000) and can become full. It is not a good idea then to just `PROTECT` everything in sight and `UNPROTECT` several thousand objects at the end. It will almost invariably be possible to either assign the objects as part of another object (which automatically protects them) or unprotect them immediately after use.

Protection is not needed for objects which R already knows are in use. In particular, this applies to function arguments.

There is a less-used macro `UNPROTECT_PTR(s)` that unprotects the object pointed to by the `SEXP s`, even if it is not the top item on the pointer protection stack. This is rarely needed outside the parser (the R sources currently have three examples, one in `'src/main/plot3d.c'`).

Sometimes an object is changed (for example duplicated, coerced or grown) yet the current value needs to be protected. For these cases `PROTECT_WITH_INDEX` saves an index of the protection location that can be used to replace the protected value using `REPROTECT`. For example (from the internal code for `optim`)

```
PROTECT_INDEX ipx;

....
PROTECT_WITH_INDEX(s = eval(OS->R_fcall, OS->R_env), &ipx);
REPROTECT(s = coerceVector(s, REALSXP), ipx);
```

### 5.9.2 Allocating storage

For many purposes it is sufficient to allocate R objects and manipulate those. There are quite a few `allocXxx` functions defined in `'Rinternals.h'`—you may want to explore them. These allocate R objects of various types, and for the standard vector types there are equivalent `NEW_XXX` macros defined in `'Rdefines.h'`.

If storage is required for C objects during the calculations this is best allocating by calling `R_alloc`; see [Section 6.1 \[Memory allocation\]](#), page 130. All of these memory allocation routines do their own error-checking, so the programmer may assume that they will raise an error and not return if the memory cannot be allocated.

### 5.9.3 Details of R types

Users of the `'Rinternals.h'` macros will need to know how the R types are known internally: if the `'Rdefines.h'` macros are used then S4-compatible names are used.

The different R data types are represented in C by `SEXPTYPE`. Some of these are familiar from R and some are internal data types. The usual R object modes are given in the table.

| <b>SEXPTYPE</b>      | <b>R equivalent</b>                           |
|----------------------|---|
| <code>REALSXP</code> | numeric with storage mode <code>double</code> |
| <code>INTSXP</code>  | integer                                       |

|        |                              |
|--------|------------------------------|
| CPLXSP | complex                      |
| LGLSP  | logical                      |
| STRSP  | character                    |
| VECSXP | list (generic vector)        |
| LISTSP | pairlist                     |
| DOTSP  | a ‘...’ object               |
| NILSP  | NULL                         |
| SYMSXP | name/symbol                  |
| CLOSP  | function or function closure |
| ENVSP  | environment                  |

Among the important internal `SEXPTYPES` are `LANGSP`, `CHARSP`, `PROMSP`, etc. (**Note:** although it is possible to return objects of internal types, it is unsafe to do so as assumptions are made about how they are handled which may be violated at user-level evaluation.) More details are given in [Section “R Internal Structures” in \*R Internals\*](#).

Unless you are very sure about the type of the arguments, the code should check the data types. Sometimes it may also be necessary to check data types of objects created by evaluating an R expression in the C code. You can use functions like `isReal`, `isInteger` and `isString` to do type checking. See the header file ‘`Rinternals.h`’ for definitions of other such functions. All of these take a `SEXP` as argument and return 1 or 0 to indicate *TRUE* or *FALSE*. Once again there are two ways to do this, and ‘`Rdefines.h`’ has macros such as `IS_NUMERIC`.

What happens if the `SEXP` is not of the correct type? Sometimes you have no other option except to generate an error. You can use the function `error` for this. It is usually better to coerce the object to the correct type. For example, if you find that an `SEXP` is of the type `INTEGER`, but you need a `REAL` object, you can change the type by using, equivalently,

```
PROTECT(newSexp = coerceVector(oldSexp, REALSXP));
```

or

```
PROTECT(newSexp = AS_NUMERIC(oldSexp));
```

Protection is needed as a new *object* is created; the object formerly pointed to by the `SEXP` is still protected but now unused.

All the coercion functions do their own error-checking, and generate NAs with a warning or stop with an error as appropriate.

Note that these coercion functions are *not* the same as calling `as.numeric` (and so on) in R code, as they do not dispatch on the class of the object. Thus it is normally preferable to do the coercion in the calling R code.

So far we have only seen how to create and coerce R objects from C code, and how to extract the numeric data from numeric R vectors. These can suffice to take us a long way in interfacing R objects to numerical algorithms, but we may need to know a little more to create useful return objects.

### 5.9.4 Attributes

Many R objects have attributes: some of the most useful are classes and the `dim` and `dimnames` that mark objects as matrices or arrays. It can also be helpful to work with the `names` attribute of vectors.

To illustrate this, let us write code to take the outer product of two vectors (which `outer` and `%o%` already do). As usual the R code is simple

```
out <- function(x, y)
{
  storage.mode(x) <- storage.mode(y) <- "double"
  .Call("out", x, y)
}
```

where we expect `x` and `y` to be numeric vectors (possibly integer), possibly with names. This time we do the coercion in the calling R code.

C code to do the computations is

```
#include <R.h>
#include <Rinternals.h>

SEXP out(SEXP x, SEXP y)
{
  R_len_t i, j, nx = length(x), ny = length(y);
  double tmp, *rx = REAL(x), *ry = REAL(y), *rans;
  SEXP ans;

  PROTECT(ans = allocMatrix(REALSXP, nx, ny));
  rans = REAL(ans);
  for(i = 0; i < nx; i++) {
    tmp = rx[i];
    for(j = 0; j < ny; j++)
      rans[i + nx*j] = tmp * ry[j];
  }
  UNPROTECT(1);
  return(ans);
}
```

Note the way `REAL` is used: as it is a function call it can be considerably faster to store the result and index that.

However, we would like to set the `dimnames` of the result. Although `allocMatrix` provides a short cut, we will show how to set the `dim` attribute directly.

```
#include <R.h>
#include <Rinternals.h>

SEXP out(SEXP x, SEXP y)
{
  R_len_t i, j, nx = length(x), ny = length(y);
  double tmp, *rx = REAL(x), *ry = REAL(y), *rans;
  SEXP ans, dim, dimnames;
```

```

    PROTECT(ans = allocVector(REALSXP, nx*ny));
    rans = REAL(ans);
    for(i = 0; i < nx; i++) {
        tmp = rx[i];
        for(j = 0; j < ny; j++)
            rans[i + nx*j] = tmp * ry[j];
    }

    PROTECT(dim = allocVector(INTSXP, 2));
    INTEGER(dim)[0] = nx; INTEGER(dim)[1] = ny;
    setAttrib(ans, R_DimSymbol, dim);

    PROTECT(dimnames = allocVector(VECSXP, 2));
    SET_VECTOR_ELT(dimnames, 0, getAttrib(x, R_NamesSymbol));
    SET_VECTOR_ELT(dimnames, 1, getAttrib(y, R_NamesSymbol));
    setAttrib(ans, R_DimNamesSymbol, dimnames);

    UNPROTECT(3);
    return(ans);
}

```

This example introduces several new features. The `getAttrib` and `setAttrib` functions get and set individual attributes. Their second argument is a `SEXP` defining the name in the symbol table of the attribute we want; these and many such symbols are defined in the header file `'Rinternals.h'`.

There are shortcuts here too: the functions `namesgets`, `dimgets` and `dimnamesgets` are the internal versions of the default methods of `names<-`, `dim<-` and `dimnames<-` (for vectors and arrays), and there are functions such as `GetMatrixDimnames` and `GetArrayDimnames`.

What happens if we want to add an attribute that is not pre-defined? We need to add a symbol for it *via* a call to `install`. Suppose for illustration we wanted to add an attribute `"version"` with value 3.0. We could use

```

SEXP version;
PROTECT(version = allocVector(REALSXP, 1));
REAL(version)[0] = 3.0;
setAttrib(ans, install("version"), version);
UNPROTECT(1);

```

Using `install` when it is not needed is harmless and provides a simple way to retrieve the symbol from the symbol table if it is already installed. However, the lookup takes a non-trivial amount of time, so consider code such as

```

static SEXP VerSymbol = NULL;
...
if (VerSymbol == NULL) VerSymbol = install("version");

```

if it is to be done frequently.



### 5.9.5 Classes

In R the class is just the attribute named "class" so it can be handled as such, but there is a shortcut `classgets`. Suppose we want to give the return value in our example the class "mat". We can use

```
#include <R.h>
#include <Rdefines.h>

....
SEXP ans, dim, dimnames, class;
....
PROTECT(class = allocVector(STRSXP, 1));
SET_STRING_ELT(class, 0, mkChar("mat"));
classgets(ans, class);
UNPROTECT(4);
return(ans);
}
```

As the value is a character vector, we have to know how to create that from a C character array, which we do using the function `mkChar`.

### 5.9.6 Handling lists

Some care is needed with lists, as R moved early on from using LISP-like lists (now called "pairlists") to S-like generic vectors. As a result, the appropriate test for an object of mode `list` is `isNewList`, and we need `allocVector(VECSXP, n)` and *not* `allocList(n)`.

List elements can be retrieved or set by direct access to the elements of the generic vector. Suppose we have a list object

```
a <- list(f = 1, g = 2, h = 3)
```

Then we can access `a$g` as `a[[2]]` by

```
double g;
....
g = REAL(VECTOR_ELT(a, 1))[0];
```

This can rapidly become tedious, and the following function (based on one in package `stats`) is very useful:

```
/* get the list element named str, or return NULL */

SEXP getListElement(SEXP list, const char *str)
{
    SEXP elmt = R_NilValue, names = getAttrib(list, R_NamesSymbol);

    for (R_len_t i = 0; i < length(list); i++)
        if (strcmp(CHAR(STRING_ELT(names, i)), str) == 0) {
            elmt = VECTOR_ELT(list, i);
            break;
        }
    return elmt;
}
```

and enables us to say

```
double g;
g = REAL(getListElement(a, "g"))[0];
```

### 5.9.7 Handling character data

R character vectors are stored as **STRSXPs**, a vector type like **VECSXP** where every element is of type **CHARSXP**. The **CHARSXP** elements of **STRSXPs** are accessed using **STRING\_ELT** and **SET\_STRING\_ELT**.

**CHARSXPs** are read-only objects and must never be modified. In particular, the C-style string contained in a **CHARSXP** should be treated as read-only and for this reason the **CHAR** function used to access the character data of a **CHARSXP** returns **(const char \*)** (this also allows compilers to issue warnings about improper use). Since **CHARSXPs** are immutable, the same **CHARSXP** can be shared by any **STRSXP** needing an element representing the same string. R maintains a global cache of **CHARSXPs** so that there is only ever one **CHARSXP** representing a given string in memory.

You can obtain a **CHARSXP** by calling **mkChar** and providing a nul-terminated C-style string. This function will return a pre-existing **CHARSXP** if one with a matching string already exists, otherwise it will create a new one and add it to the cache before returning it to you. The variant **mkCharLen** can be used to create a **CHARSXP** from part of a buffer and will ensure null-termination.

### 5.9.8 Finding and setting variables

It will be usual that all the R objects needed in our C computations are passed as arguments to **.Call** or **.External**, but it is possible to find the values of R objects from within the C given their names. The following code is the equivalent of **get(name, envir = rho)**.

```
SEXP getvar(SEXP name, SEXP rho)
{
    SEXP ans;

    if(!isString(name) || length(name) != 1)
        error("name is not a single string");
    if(!isEnvironment(rho))
        error("rho should be an environment");
    ans = findVar(install(CHAR(STRING_ELT(name, 0))), rho);
    Rprintf("first value is %f\n", REAL(ans)[0]);
    return(R_NilValue);
}
```

The main work is done by **findVar**, but to use it we need to install **name** as a name in the symbol table. As we wanted the value for internal use, we return **NULL**.

Similar functions with syntax

```
void defineVar(SEXP symbol, SEXP value, SEXP rho)
void setVar(SEXP symbol, SEXP value, SEXP rho)
```

can be used to assign values to R variables. **defineVar** creates a new binding or changes the value of an existing binding in the specified environment frame; it is the analogue of **assign(symbol, value, envir = rho, inherits = FALSE)**, but unlike **assign**, **defineVar**

does not make a copy of the object `value`.<sup>6</sup> `setVar` searches for an existing binding for `symbol` in `rho` or its enclosing environments. If a binding is found, its value is changed to `value`. Otherwise, a new binding with the specified value is created in the global environment. This corresponds to `assign(symbol, value, envir = rho, inherits = TRUE)`.

### 5.9.9 Some convenience functions

Some operations are done so frequently that there are convenience functions to handle them. (All these are provided via the header file ‘`Rinternals.h`’.)

Suppose we wanted to pass a single logical argument `ignore_quotes`: we could use

```
int ign = asLogical(ignore_quotes);
if(ign == NA_LOGICAL) error("'ignore_quotes' must be TRUE or FALSE");
```

which will do any coercion needed (at least from a vector argument), and return `NA_LOGICAL` if the value passed was `NA` or coercion failed. There are also `asInteger`, `asReal` and `asComplex`. The function `asChar` returns a `CHARSXP`. All of these functions ignore any elements of an input vector after the first.

To return a length-one real vector we can use

```
double x;

...
return ScalarReal(x);
```

and there are versions of this for all the atomic vector types (those for a length-one character vector being `ScalarString` with argument a `CHARSXP` and `mkString` with argument `const char *`).

Some of the `isXXXX` functions differ from their apparent R-level counterparts: for example `isVector` is true for any atomic vector type (`isVectorAtomic`) and for lists and expressions (`isVectorList`) (with no check on attributes). `isMatrix` is a test of a length-2 “`dim`” attribute.

There are a series of small macros/functions to help construct pairlists and language objects (whose internal structures just differ by `SEXPTYPE`). Function `CONS(u, v)` is the basic building block: it constructs a pairlist from `u` followed by `v` (which is a pairlist or `R_NilValue`). `LCONS` is a variant that constructs a language object. Functions `list1` to `list5` construct a pairlist from one to five items, and `lang1` to `lang6` do the same for a language object (a function to call plus zero to five arguments). Functions `elt` and `lastElt` find the *i*th element and the last element of a pairlist, and `nthcdr` returns a pointer to the *n*th position in the pairlist (whose `CAR` is the *n*th item).

Functions `str2type` and `type2str` map R length-one character strings to and from `SEXPTYPE` numbers, and `type2char` maps numbers to C character strings.

#### 5.9.9.1 Semi-internal convenience functions

There is quite a collection of functions that may be used in your C code *if* you are willing to adapt to rare “API” changes. These typically contain “workhorses” of their R counterparts.

---

<sup>6</sup> You can assign a *copy* of the object in the environment frame `rho` using `defineVar(symbol, duplicate(value), rho)`.

Functions `any_duplicated` and `any_duplicated3` are fast versions of R's `any(duplicated(.))`.

Function `R_compute_identical` corresponds to R's `identical` function.

### 5.9.10 Named objects and copying

When assignments are done in R such as

```
x <- 1:10
y <- x
```

the named object is not necessarily copied, so after those two assignments `y` and `x` are bound to the same `SEXP` (the structure a `SEXP` points to). This means that any code which alters one of them has to make a copy before modifying the copy if the usual R semantics are to apply. Note that whereas `.C` and `.Fortran` do copy their arguments (unless the dangerous `dup = FALSE` is used), `.Call` and `.External` do not. So `duplicate` is commonly called on arguments to `.Call` before modifying them.

However, at least some of this copying is unneeded. In the first assignment shown, `x <- 1:10`, R first creates an object with value `1:10` and then assigns it to `x` but if `x` is modified no copy is necessary as the temporary object with value `1:10` cannot be referred to again. R distinguishes between named and unnamed objects *via* a field in a `SEXP` that can be accessed *via* the macros `NAMED` and `SET_NAMED`. This can take values

- 0           The object is not bound to any symbol
- 1           The object has been bound to exactly one symbol
- 2           The object has potentially been bound to two or more symbols, and one should act as if another variable is currently bound to this value.

Note the past tenses: R does not do full reference counting and there may currently be fewer bindings.

It is safe to modify the value of any `SEXP` for which `NAMED(foo)` is zero, and if `NAMED(foo)` is two, the value should be duplicated (*via* a call to `duplicate`) before any modification. Note that it is the responsibility of the author of the code making the modification to do the duplication, even if it is `x` whose value is being modified after `y <- x`.

The case `NAMED(foo) == 1` allows some optimization, but it can be ignored (and duplication done whenever `NAMED(foo) > 0`). (This optimization is not currently usable in user code.) It is intended for use within replacement functions. Suppose we used

```
x <- 1:10
foo(x) <- 3
```

which is computed as

```
x <- 1:10
x <- "foo<-"(x, 3)
```

Then inside `"foo<-"` the object pointing to the current value of `x` will have `NAMED(foo)` as one, and it would be safe to modify it as the only symbol bound to it is `x` and that will be rebound immediately. (Provided the remaining code in `"foo<-"` make no reference to `x`, and no one is going to attempt a direct call such as `y <- "foo<-"(x)`.)

Currently all arguments to a `.Call` call will have `NAMED` set to 2, and so users must assume that they need to be duplicated before alteration.

## 5.10 Interface functions `.Call` and `.External`

In this section we consider the details of the R/C interfaces.

These two interfaces have almost the same functionality. `.Call` is based on the interface of the same name in S version 4, and `.External` is based on `.Internal`. `.External` is more complex but allows a variable number of arguments.

### 5.10.1 Calling `.Call`

Let us convert our finite convolution example to use `.Call`, first using the ‘`Rdefines.h`’ macros. The calling function in R is

```
conv <- function(a, b) .Call("convolve2", a, b)
```

which could hardly be simpler, but as we shall see all the type checking must be transferred to the C code, which is

```
#include <R.h>
#include <Rdefines.h>

SEXP convolve2(SEXP a, SEXP b)
{
    R_len_t i, j, na, nb, nab;
    double *xa, *xb, *xab;
    SEXP ab;

    PROTECT(a = AS_NUMERIC(a));
    PROTECT(b = AS_NUMERIC(b));
    na = LENGTH(a); nb = LENGTH(b); nab = na + nb - 1;
    PROTECT(ab = NEW_NUMERIC(nab));
    xa = NUMERIC_POINTER(a); xb = NUMERIC_POINTER(b);
    xab = NUMERIC_POINTER(ab);
    for(i = 0; i < nab; i++) xab[i] = 0.0;
    for(i = 0; i < na; i++)
        for(j = 0; j < nb; j++) xab[i + j] += xa[i] * xb[j];
    UNPROTECT(3);
    return(ab);
}
```

Now for the version in ‘`Rinternals.h`’ style. Only the C code changes.

```

#include <R.h>
#include <Rinternals.h>

SEXP convolve2(SEXP a, SEXP b)
{
    R_len_t i, j, na, nb, nab;
    double *xa, *xb, *xab;
    SEXP ab;

    PROTECT(a = coerceVector(a, REALSXP));
    PROTECT(b = coerceVector(b, REALSXP));
    na = length(a); nb = length(b); nab = na + nb - 1;
    PROTECT(ab = allocVector(REALSXP, nab));
    xa = REAL(a); xb = REAL(b);
    xab = REAL(ab);
    for(i = 0; i < nab; i++) xab[i] = 0.0;
    for(i = 0; i < na; i++)
        for(j = 0; j < nb; j++) xab[i + j] += xa[i] * xb[j];
    UNPROTECT(3);
    return(ab);
}

```

This is called in exactly the same way.

### 5.10.2 Calling .External

We can use the same example to illustrate `.External`. The R code changes only by replacing `.Call` by `.External`

```
conv <- function(a, b) .External("convolveE", a, b)
```

but the main change is how the arguments are passed to the C code, this time as a single `SEXP`. The only change to the C code is how we handle the arguments.

```

#include <R.h>
#include <Rinternals.h>

SEXP convolveE(SEXP args)
{
    R_len_t i, j, na, nb, nab;
    double *xa, *xb, *xab;
    SEXP a, b, ab;

    PROTECT(a = coerceVector(CADR(args), REALSXP));
    PROTECT(b = coerceVector(CADDR(args), REALSXP));
    ...
}

```

Once again we do not need to protect the arguments, as in the R side of the interface they are objects that are already in use. The macros

```

first = CADR(args);
second = CADDR(args);
third = CADDDR(args);
fourth = CAD4R(args);

```

provide convenient ways to access the first four arguments. More generally we can use the CDR and CAR macros as in

```

args = CDR(args); a = CAR(args);
args = CDR(args); b = CAR(args);

```

which clearly allows us to extract an unlimited number of arguments (whereas .Call has a limit, albeit at 65 not a small one).

More usefully, the .External interface provides an easy way to handle calls with a variable number of arguments, as length(args) will give the number of arguments supplied (of which the first is ignored). We may need to know the names ('tags') given to the actual arguments, which we can by using the TAG macro and using something like the following example, that prints the names and the first value of its arguments if they are vector types.

```

#include <R_ext/PrtUtil.h>

SEXP showArgs(SEXP args)
{
    args = CDR(args); /* skip 'name' */
    for(int i = 0; args != R_NilValue; i++, args = CDR(args)) {
        const char *name =
            isNull(TAG(args)) ? "" : CHAR(PRINTNAME(TAG(args)));
        SEXP el = CAR(args);
        if (length(el) == 0) {
            Rprintf("[%d] '%s' R type, length 0\n", i+1, name);
            continue;
        }
        switch(TYPEOF(el)) {
        case REALSXP:
            Rprintf("[%d] '%s' %f\n", i+1, name, REAL(el)[0]);
            break;
        case LGLSXP:
        case INTSXP:
            Rprintf("[%d] '%s' %d\n", i+1, name, INTEGER(el)[0]);
            break;
        case CPLXSXP:
            {
                Rcomplex cpl = COMPLEX(el)[0];
                Rprintf("[%d] '%s' %f + %fi\n", i+1, name, cpl.r, cpl.i);
            }
            break;
        case STRSXP:
            Rprintf("[%d] '%s' %s\n", i+1, name,
                CHAR(STRING_ELT(el, 0)));
            break;
        }
    }
}

```



```

        default:
            Rprintf("[%d] '%s' R type\n", i+1, name);
        }
    }
    return(R_NilValue);
}

```

This can be called by the wrapper function

```
showArgs <- function(...) invisible(.External("showArgs", ...))
```

Note that this style of programming is convenient but not necessary, as an alternative style is

```
showArgs1 <- function(...) invisible(.Call("showArgs1", list(...)))
```

The (very similar) C code is in the scripts.

### 5.10.3 Missing and special values

One piece of error-checking the `.C` call does (unless `NAOK` is true) is to check for missing (`NA`) and IEEE special values (`Inf`, `-Inf` and `NaN`) and give an error if any are found. With the `.Call` interface these will be passed to our code. In this example the special values are no problem, as IEC60559 arithmetic will handle them correctly. In the current implementation this is also true of `NA` as it is a type of `NaN`, but it is unwise to rely on such details. Thus we will re-write the code to handle `NA`s using macros defined in `'R_exts/Arith.h'` included by `'R.h'`.

The code changes are the same in any of the versions of `convolve2` or `convolveE`:

```

...
for(i = 0; i < na; i++)
    for(j = 0; j < nb; j++)
        if(ISNA(xa[i]) || ISNA(xb[j]) || ISNA(xab[i + j]))
            xab[i + j] = NA_REAL;
        else
            xab[i + j] += xa[i] * xb[j];
...

```

Note that the `ISNA` macro, and the similar macros `ISNAN` (which checks for `NaN` or `NA`) and `R_FINITE` (which is false for `NA` and all the special values), only apply to numeric values of type `double`. Missingness of integers, logicals and character strings can be tested by equality to the constants `NA_INTEGER`, `NA_LOGICAL` and `NA_STRING`. These and `NA_REAL` can be used to set elements of R vectors to `NA`.

The constants `R_NaN`, `R_PosInf` and `R_NegInf` can be used to set doubles to the special values.

## 5.11 Evaluating R expressions from C

The main function we will use is

```
SEXP eval(SEXP expr, SEXP rho);
```

the equivalent of the interpreted R code `eval(expr, envir = rho)`, although we can also make use of `findVar`, `defineVar` and `findFun` (which restricts the search to functions).

To see how this might be applied, here is a simplified internal version of `lapply` for expressions, used as

```
a <- list(a = 1:5, b = rnorm(10), test = runif(100))
.Call("lapply", a, quote(sum(x)), new.env())
```

with C code

```
SEXP lapply(SEXP list, SEXP expr, SEXP rho)
{
    R_len_t i, n = length(list);
    SEXP ans;

    if(!isNewList(list)) error("'list' must be a list");
    if(!isEnvironment(rho)) error("'rho' should be an environment");
    PROTECT(ans = allocVector(VECSXP, n));
    for(i = 0; i < n; i++) {
        defineVar(install("x"), VECTOR_ELT(list, i), rho);
        SET_VECTOR_ELT(ans, i, eval(expr, rho));
    }
    setAttrib(ans, R_NamesSymbol, getAttrib(list, R_NamesSymbol));
    UNPROTECT(1);
    return(ans);
}
```

It would be closer to `lapply` if we could pass in a function rather than an expression. One way to do this is *via* interpreted R code as in the next example, but it is possible (if somewhat obscure) to do this in C code. The following is based on the code in ‘src/main/optimize.c’.

```
SEXP lapply2(SEXP list, SEXP fn, SEXP rho)
{
    R_len_t i, n = length(list);
    SEXP R_fcall, ans;

    if(!isNewList(list)) error("'list' must be a list");
    if(!isFunction(fn)) error("'fn' must be a function");
    if(!isEnvironment(rho)) error("'rho' should be an environment");
    PROTECT(R_fcall = lang2(fn, R_NilValue));
    PROTECT(ans = allocVector(VECSXP, n));
    for(i = 0; i < n; i++) {
        SETCADR(R_fcall, VECTOR_ELT(list, i));
        SET_VECTOR_ELT(ans, i, eval(R_fcall, rho));
    }
    setAttrib(ans, R_NamesSymbol, getAttrib(list, R_NamesSymbol));
    UNPROTECT(2);
    return(ans);
}
```

used by

```
.Call("lapply2", a, sum, new.env())
```

Function `lang2` creates an executable pairlist of two elements, but this will only be clear to those with a knowledge of a LISP-like language.

As a more comprehensive example of constructing an R call in C code and evaluating, consider the following fragment of `printAttributes` in ‘`src/main/print.c`’.

```
/* Need to construct a call to
   print(CAR(a), digits=digits)
   based on the R_print structure, then eval(call, env).
   See do_docal for the template for this sort of thing.
*/
SEXP s, t;
PROTECT(t = s = allocList(3));
SET_TYPEOF(s, LANGSXP);
SETCAR(t, install("print")); t = CDR(t);
SETCAR(t, CAR(a)); t = CDR(t);
SETCAR(t, ScalarInteger(digits));
SET_TAG(t, install("digits"));
eval(s, env);
UNPROTECT(1);
```

At this point `CAR(a)` is the R object to be printed, the current attribute. There are three steps: the call is constructed as a pairlist of length 3, the list is filled in, and the expression represented by the pairlist is evaluated.

A pairlist is quite distinct from a generic vector list, the only user-visible form of list in R. A pairlist is a linked list (with `CDR(t)` computing the next entry), with items (accessed by `CAR(t)`) and names or tags (set by `SET_TAG`). In this call there are to be three items, a symbol (pointing to the function to be called) and two argument values, the first unnamed and the second named. Setting the type to `LANGSXP` makes this a call which can be evaluated.

### 5.11.1 Zero-finding

In this section we re-work the example of Becker, Chambers & Wilks (1988, pp.~205–10) on finding a zero of a univariate function. The R code and an example are

```
zero <- function(f, guesses, tol = 1e-7) {
  f.check <- function(x) {
    x <- f(x)
    if(!is.numeric(x)) stop("Need a numeric result")
    as.double(x)
  }
  .Call("zero", body(f.check), as.double(guesses), as.double(tol),
        new.env())
}

cube1 <- function(x) (x^2 + 1) * (x - 1.5)
zero(cube1, c(0, 5))
```

where this time we do the coercion and error-checking in the R code. The C code is

```

SEXP mkans(double x)
{
    SEXP ans;
    PROTECT(ans = allocVector(REALSXP, 1));
    REAL(ans)[0] = x;
    UNPROTECT(1);
    return ans;
}

double feval(double x, SEXP f, SEXP rho)
{
    defineVar(install("x"), mkans(x), rho);
    return(REAL(eval(f, rho))[0]);
}

SEXP zero(SEXP f, SEXP guesses, SEXP stol, SEXP rho)
{
    double x0 = REAL(guesses)[0], x1 = REAL(guesses)[1],
           tol = REAL(stol)[0];
    double f0, f1, fc, xc;

    if(tol <= 0.0) error("non-positive tol value");
    f0 = feval(x0, f, rho); f1 = feval(x1, f, rho);
    if(f0 == 0.0) return mkans(x0);
    if(f1 == 0.0) return mkans(x1);
    if(f0*f1 > 0.0) error("x[0] and x[1] have the same sign");

    for(;;) {
        xc = 0.5*(x0+x1);
        if(fabs(x0-x1) < tol) return mkans(xc);
        fc = feval(xc, f, rho);
        if(fc == 0) return mkans(xc);
        if(f0*fc > 0.0) {
            x0 = xc; f0 = fc;
        } else {
            x1 = xc; f1 = fc;
        }
    }
}

```

### 5.11.2 Calculating numerical derivatives

We will use a longer example (by Saikat DebRoy) to illustrate the use of evaluation and `.External`. This calculates numerical derivatives, something that could be done as effectively in interpreted R code but may be needed as part of a larger C calculation.

An interpreted R version and an example are

```

numeric.deriv <- function(expr, theta, rho=sys.frame(sys.parent()))
{
  eps <- sqrt(.Machine$double.eps)
  ans <- eval(substitute(expr), rho)
  grad <- matrix(, length(ans), length(theta),
                 dimnames=list(NULL, theta))
  for (i in seq_along(theta)) {
    old <- get(theta[i], envir=rho)
    delta <- eps * max(1, abs(old))
    assign(theta[i], old+delta, envir=rho)
    ans1 <- eval(substitute(expr), rho)
    assign(theta[i], old, envir=rho)
    grad[, i] <- (ans1 - ans)/delta
  }
  attr(ans, "gradient") <- grad
  ans
}
omega <- 1:5; x <- 1; y <- 2
numeric.deriv(sin(omega*x*y), c("x", "y"))

```

where `expr` is an expression, `theta` a character vector of variable names and `rho` the environment to be used.

For the compiled version the call from R will be

```
.External("numeric_deriv", expr, theta, rho)
```

with example usage

```
.External("numeric_deriv", quote(sin(omega*x*y)),
          c("x", "y"), .GlobalEnv)
```

Note the need to quote the expression to stop it being evaluated in the caller.

Here is the complete C code which we will explain section by section.

```

#include <R.h> /* for DOUBLE_EPS */
#include <Rinternals.h>

SEXP numeric_deriv(SEXP args)
{
  SEXP theta, expr, rho, ans, ans1, gradient, par, dimnames;
  double tt, xx, delta, eps = sqrt(DOUBLE_EPS), *rgr, *rans;
  R_len_t start, i, j;

  expr = CADR(args);
  if(!isString(theta = CADDR(args)))
    error("theta should be of type character");
  if(!isEnvironment(rho = CADDRDR(args)))
    error("rho should be an environment");

```

```

PROTECT(ans = coerceVector(eval(expr, rho), REALSXP));
PROTECT(gradient = allocMatrix(REALSXP, LENGTH(ans), LENGTH(theta)));
rgr = REAL(gradient); rans = REAL(ans);

for(i = 0, start = 0; i < LENGTH(theta); i++, start += LENGTH(ans)) {
  PROTECT(par = findVar(install(CHAR(String_ELT(theta, i))), rho));
  tt = REAL(par)[0];
  xx = fabs(tt);
  delta = (xx < 1) ? eps : xx*eps;
  REAL(par)[0] += delta;
  PROTECT(ans1 = coerceVector(eval(expr, rho), REALSXP));
  for(j = 0; j < LENGTH(ans); j++)
    rgr[j + start] = (REAL(ans1)[j] - rans[j])/delta;
  REAL(par)[0] = tt;
  UNPROTECT(2); /* par, ans1 */
}

PROTECT(dimnames = allocVector(VECSXP, 2));
SET_VECTOR_ELT(dimnames, 1, theta);
dimnamesgets(gradient, dimnames);
setAttrib(ans, install("gradient"), gradient);
UNPROTECT(3); /* ans gradient dimnames */
return ans;
}

```

The code to handle the arguments is

```

expr = CADDR(args);
if(!isString(theta = CADDR(args)))
  error("theta should be of type character");
if(!isEnvironment(rho = CADDR(args)))
  error("rho should be an environment");

```

Note that we check for correct types of `theta` and `rho` but do not check the type of `expr`. That is because `eval` can handle many types of R objects other than `EXPRSXP`. There is no useful coercion we can do, so we stop with an error message if the arguments are not of the correct mode.

The first step in the code is to evaluate the expression in the environment `rho`, by

```
PROTECT(ans = coerceVector(eval(expr, rho), REALSXP));
```

We then allocate space for the calculated derivative by

```
PROTECT(gradient = allocMatrix(REALSXP, LENGTH(ans), LENGTH(theta)));
```

The first argument to `allocMatrix` gives the `SEXPTYPE` of the matrix: here we want it to be `REALSXP`. The other two arguments are the numbers of rows and columns.

```

for(i = 0, start = 0; i < LENGTH(theta); i++, start += LENGTH(ans)) {
  PROTECT(par = findVar(install(CHAR(String_ELT(theta, i))), rho));

```

Here, we are entering a for loop. We loop through each of the variables. In the for loop, we first create a symbol corresponding to the *i*'th element of the `STRSXP` `theta`. Here, `String_ELT(theta, i)` accesses the *i*'th element of the `STRSXP` `theta`. Macro `CHAR()` extracts the

actual character representation<sup>7</sup> of it: it returns a pointer. We then install the name and use `findVar` to find its value.

```
tt = REAL(par)[0];
xx = fabs(tt);
delta = (xx < 1) ? eps : xx*eps;
REAL(par)[0] += delta;
PROTECT(ans1 = coerceVector(eval(expr, rho), REALSXP));
```

We first extract the real value of the parameter, then calculate `delta`, the increment to be used for approximating the numerical derivative. Then we change the value stored in `par` (in environment `rho`) by `delta` and evaluate `expr` in environment `rho` again. Because we are directly dealing with original R memory locations here, R does the evaluation for the changed parameter value.

```
for(j = 0; j < LENGTH(ans); j++)
    rgr[j + start] = (REAL(ans1)[j] - rans[j])/delta;
REAL(par)[0] = tt;
UNPROTECT(2);
}
```

Now, we compute the `i`'th column of the gradient matrix. Note how it is accessed: R stores matrices by column (like FORTRAN).

```
PROTECT(dimnames = allocVector(VECSXP, 2));
SET_VECTOR_ELT(dimnames, 1, theta);
dimnamesgets(gradient, dimnames);
setAttrib(ans, install("gradient"), gradient);
UNPROTECT(3);
return ans;
}
```

First we add column names to the gradient matrix. This is done by allocating a list (a `VECSXP`) whose first element, the row names, is `NULL` (the default) and the second element, the column names, is set as `theta`. This list is then assigned as the attribute having the symbol `R_DimNamesSymbol`. Finally we set the gradient matrix as the gradient attribute of `ans`, unprotect the remaining protected locations and return the answer `ans`.

## 5.12 Parsing R code from C

Suppose an R extension want to accept an R expression from the user and evaluate it. The previous section covered evaluation, but the expression will be entered as text and needs to be parsed first. A small part of R's parse interface is declared in header file `'R_ext/Parse.h'`<sup>8</sup>.

An example of the usage can be found in the (example) Windows package `windlgs` included in the R source tree. The essential part is

<sup>7</sup> see [Section 5.15 \[Character encoding issues\]](#), page 129 for why this might not be what is required.

<sup>8</sup> This is only guaranteed to show the current interface: it is liable to change.



```

#include <R.h>
#include <Rinternals.h>
#include <R_ext/Parse.h>

SEXP menu_ttest3()
{
    char cmd[256];
    SEXP cmdSexp, cmdexpr, ans = R_NilValue;
    ParseStatus status;
    ...
    if(done == 1) {
        PROTECT(cmdSexp = allocVector(STRSXP, 1));
        SET_STRING_ELT(cmdSexp, 0, mkChar(cmd));
        cmdexpr = PROTECT(R_ParseVector(cmdSexp, -1, &status, R_NilValue));
        if (status != PARSE_OK) {
            UNPROTECT(2);
            error("invalid call %s", cmd);
        }
        /* Loop is needed here as EXPSEXP will be of length > 1 */
        for(R_len_t i = 0; i < length(cmdexpr); i++)
            ans = eval(VECTOR_ELT(cmdexpr, i), R_GlobalEnv);
        UNPROTECT(2);
    }
    return ans;
}

```

Note that a single line of text may give rise to more than one R expression.

`R_ParseVector` is essentially the code used to implement `parse(text=)` at R level. The first argument is a character vector (corresponding to `text`) and the second the maximal number of expressions to parse (corresponding to `n`). The third argument is a pointer to a variable of an enumeration type, and it is normal (as `parse` does) to regard all values other than `PARSE_OK` as an error. Other values which might be returned are `PARSE_INCOMPLETE` (an incomplete expression was found) and `PARSE_ERROR` (a syntax error), in both cases the value returned being `R_NilValue`. The fourth argument is a `srcfile` object or the R `NULL` object (as in the example above). In the former case a `srcref` attribute would be attached to the result, containing a list of `srcref` objects of the same length as the expression, to allow it to be echoed with its original formatting.

### 5.13 External pointers and weak references

The SEXPTYPES `EXTPTRSXP` and `WEAKREFSXP` can be encountered at R level, but are created in C code.

External pointer SEXPs are intended to handle references to C structures such as ‘handles’, and are used for this purpose in package **RODBC** for example. They are unusual in their copying semantics in that when an R object is copied, the external pointer object is not duplicated. (For this reason external pointers should only be used as part of an object with normal semantics, for example an attribute or an element of a list.)

An external pointer is created by

```
SEXP R_MakeExternalPtr(void *p, SEXP tag, SEXP prot);
```

where `p` is the pointer (and hence this cannot portably be a function pointer), and `tag` and `prot` are references to ordinary R objects which will remain in existence (be protected from garbage collection) for the lifetime of the external pointer object. A useful convention is to use the `tag` field for some form of type identification and the `prot` field for protecting the memory that the external pointer represents, if that memory is allocated from the R heap. Both `tag` and `prot` can be `R_NilValue`, and often are.

The elements of an external pointer can be accessed and set *via*

```
void *R_ExternalPtrAddr(SEXP s);
SEXP R_ExternalPtrTag(SEXP s);
SEXP R_ExternalPtrProtected(SEXP s);
void R_ClearExternalPtr(SEXP s);
void R_SetExternalPtrAddr(SEXP s, void *p);
void R_SetExternalPtrTag(SEXP s, SEXP tag);
void R_SetExternalPtrProtected(SEXP s, SEXP p);
```

Clearing a pointer sets its value to the C NULL pointer.

An external pointer object can have a *finalizer*, a piece of code to be run when the object is garbage collected. This can be R code or C code, and the various interfaces are, respectively.

```
void R_RegisterFinalizerEx(SEXP s, SEXP fun, Rboolean onexit);

typedef void (*R_CFinalizer_t)(SEXP);
void R_RegisterCFinalizerEx(SEXP s, R_CFinalizer_t fun, Rboolean onexit);
```

The R function indicated by `fun` should be a function of a single argument, the object to be finalized. R does not perform a garbage collection when shutting down, and the `onexit` argument of the extended forms can be used to ask that the finalizer be run during a normal shutdown of the R session. It is suggested that it is good practice to clear the pointer on finalization.

The only R level function for interacting with external pointers is `reg.finalizer` which can be used to set a finalizer.

It is probably not a good idea to allow an external pointer to be **saved** and then reloaded, but if this happens the pointer will be set to the C NULL pointer.

Weak references are used to allow the programmer to maintain information on entities without preventing the garbage collection of the entities once they become unreachable.

A weak reference contains a key and a value. The value is reachable is if it either reachable directly or *via* weak references with reachable keys. Once a value is determined to be unreachable during garbage collection, the key and value are set to `R_NilValue` and the finalizer will be run later in the garbage collection.

Weak reference objects are created by one of

```
SEXP R_MakeWeakRef(SEXP key, SEXP val, SEXP fin, Rboolean onexit);
SEXP R_MakeWeakRefC(SEXP key, SEXP val, R_CFinalizer_t fin,
                    Rboolean onexit);
```

where the R or C finalizer are specified in exactly the same way as for an external pointer object (whose finalization interface is implemented *via* weak references).

The parts can be accessed *via*

```
SEXP R_WeakRefKey(SEXP w);
SEXP R_WeakRefValue(SEXP w);
void R_RunWeakRefFinalizer(SEXP w);
```

A toy example of the use of weak references can be found at [www.stat.uiowa.edu/~luke/R/references/weakfinex.html](http://www.stat.uiowa.edu/~luke/R/references/weakfinex.html), but that is used to add finalizers to external pointers which can now be done more directly. At the time of writing no CRAN or Bioconductor package uses weak references.

### 5.13.1 An example

Package **RODBC** uses external pointers to maintain its *channels*, connections to databases. There can be several connections open at once, and the status information for each is stored in a C structure (pointed to by `this_handle`) in the code extract below) that is returned *via* an external pointer as part of the RODBC ‘channel’ (as the “`handle_ptr`” attribute). The external pointer is created by

```
SEXP ans, ptr;
PROTECT(ans = allocVector(INTSXP, 1));
ptr = R_MakeExternalPtr(thisHandle, install("RODBC_channel"), R_NilValue);
PROTECT(ptr);
R_RegisterCFinalizerEx(ptr, chanFinalizer, TRUE);
...
/* return the channel no */
INTEGER(ans)[0] = nChannels;
/* and the connection string as an attribute */
setAttrib(ans, install("connection.string"), constr);
setAttrib(ans, install("handle_ptr"), ptr);
UNPROTECT(3);
return ans;
```

Note the symbol given to identify the usage of the external pointer, and the use of the finalizer. Since the final argument when registering the finalizer is `TRUE`, the finalizer will be run at the the of the R session (unless it crashes). This is used to close and clean up the connection to the database. The finalizer code is simply

```
static void chanFinalizer(SEXP ptr)
{
    if(!R_ExternalPtrAddr(ptr)) return;
    inRODBCClose(R_ExternalPtrAddr(ptr));
    R_ClearExternalPtr(ptr); /* not really needed */
}
```

Clearing the pointer and checking for a NULL pointer avoids any possibility of attempting to close an already-closed channel.

R’s connections provide another example of using external pointers, in that case purely to be able to use a finalizer to close and destroy the connection if it is no longer is use.

## 5.14 Vector accessor functions

The vector accessors like `REAL` and `INTEGER` and `VECTOR_ELT` are *functions* when used in R extensions. (For efficiency they are macros when used in the R source code, apart from `SET_STRING_ELT` and `SET_VECTOR_ELT` which are always functions.)

The accessor functions check that they are being used on an appropriate type of `SEXP`.

If efficiency is essential, the macro versions of the accessors can be obtained by defining `'USE_RINTERNALS'` before including `'Rinternals.h'`. If you find it necessary to do so, please do test that your code compiles without `'USE_RINTERNALS'` defined, as this provides a stricter test that the accessors have been used correctly.

## 5.15 Character encoding issues

`CHARSXPs` can be marked as coming from a known encoding (Latin-1 or UTF-8). This is mainly intended for human-readable output, and most packages can just treat such `CHARSXPs` as a whole. However, if they need to be interpreted as characters or output at C level then it would normally be correct to ensure that they are converted to the encoding of the current locale: this can be done by accessing the data in the `CHARSXP` by `translateChar` rather than by `CHAR`. If re-encoding is needed this allocates memory with `R_alloc` which thus persists to the end of the `.Call/.External` call unless `vmxset` is used.

There is a similar function `translateCharUTF8` which converts to UTF-8: this has the advantage that a faithful translation is almost always possible (whereas only a few languages can be represented in the encoding of the current locale unless that is UTF-8).

There is a public interface to the encoding marked on `CHARSXPs` *via*

```
typedef enum {CE_NATIVE, CE_UTF8, CE_LATIN1, CE_SYMBOL, CE_ANY} cetype_t;
cetype_t getCharCE(SEXP);
SEXP mkCharCE(const char *, cetype_t);
```

Only `CE_UTF8` and `CE_LATIN1` are marked on `CHARSXPs` (and so `Rf_getCharCE` will only return one of the first three), and these should only be used on non-ASCII strings. Value `CE_SYMBOL` is used internally to indicate Adobe Symbol encoding. Value `CE_ANY` is used to indicate a character string that will not need re-encoding – this is used for character strings known to be in ASCII, and can also be used as an input parameter where the intention is that the string is treated as a series of bytes.

Function

```
const char *reEnc(const char *x, cetype_t ce_in, cetype_t ce_out,
                  int subst);
```

can be used to re-encode character strings: like `translateChar` it returns a string allocated by `R_alloc`. This can translate from `CE_SYMBOL` to `CE_UTF8`, but not conversely. Argument `subst` controls what to do with untranslatable characters or invalid input: this is done byte-by-byte with 1 indicates to output hex of the form `<a0>`, and 2 to replace by `.`, with any other value causing the byte to produce no output.

There is also

```
SEXP mkCharLenCE(const char *, int, cetype_t);
```

to create marked character strings of a given length.

## 6 The R API: entry points for C code

There are a large number of entry points in the R executable/DLL that can be called from C code (and some that can be called from FORTRAN code). Only those documented here are stable enough that they will only be changed with considerable notice.

The recommended procedure to use these is to include the header file ‘R.h’ in your C code by

```
#include <R.h>
```

This will include several other header files from the directory ‘R\_INCLUDE\_DIR/R\_ext’, and there are other header files there that can be included too, but many of the features they contain should be regarded as undocumented and unstable.

An alternative is to include the header file ‘S.h’, which may be useful when porting code from S. This includes rather less than ‘R.h’, and has some extra compatibility definitions (for example the `S_complex` type from S).

The defines used for compatibility with S sometimes causes conflicts (notably with Windows headers), and the known problematic defines can be removed by defining `STRICT_R_HEADERS`.

Most of these header files, including all those included by ‘R.h’, can be used from C++ code. Some others need to be included within an `extern "C"` declaration, and for clarity this is advisable for all R header files.

**Note:** Because R re-maps many of its external names to avoid clashes with user code, it is *essential* to include the appropriate header files when using these entry points.

This remapping can cause problems<sup>1</sup>, and can be eliminated by defining `R_NO_REMAP` and prepending ‘Rf\_’ to *all* the function names used from ‘Rinternals.h’ and ‘R\_ext/Error.h’.

We can classify the entry points as

|                |  |
|----------------|--|
| <i>API</i>     | Entry points which are documented in this manual and declared in an installed header file. These can be used in distributed packages and will only be changed after deprecation. |
| <i>public</i>  | Entry points declared in an installed header file that are exported on all R platforms but are not documented and subject to change without notice.                              |
| <i>private</i> | Entry points that are used when building R and exported on all R platforms but are not declared in the installed header files. Do not use these in distributed code.             |
| <i>hidden</i>  | Entry points that are where possible (Windows and some modern Unix-alike compilers/loaders when using R as a shared library) not exported.                                       |

### 6.1 Memory allocation

There are two types of memory allocation available to the C programmer, one in which R manages the clean-up and the other in which user has full control (and responsibility).

---

<sup>1</sup> Known problems are redefining `error`, `length`, `vector` and `warning`

### 6.1.1 Transient storage allocation

Here R will reclaim the memory at the end of the call to `.C`. Use

```
char *R_alloc(size_t n, int size)
```

which allocates  $n$  units of  $size$  bytes each. A typical usage (from package `stats`) is

```
x = (int *) R_alloc(nrows(merge)+2, sizeof(int));
```

(`size_t` is defined in `'stddef.h'` which the header defining `R_alloc` includes.)

There is a similar call, `S_alloc` (for compatibility with older versions of S) which zeroes the memory allocated,

```
char *S_alloc(long n, int size)
```

and

```
char *S_realloc(char *p, long new, long old, int size)
```

which changes the allocation size from  $old$  to  $new$  units, and zeroes the additional units.

For compatibility with current versions of S, header `'S.h'` (only) defines wrapper macros equivalent to

```
type* Salloc(long n, int type)
```

```
type* Srealloc(char *p, long new, long old, int type)
```

This memory is taken from the heap, and released at the end of the `.C`, `.Call` or `.External` call. Users can also manage it, by noting the current position with a call to `vmxget` and clearing memory allocated subsequently by a call to `vmxset`. This is only recommended for experts.

Note that this memory will be freed on error or user interrupt (if allowed: see [Section 6.12 \[Allowing interrupts\]](#), page 143).

Note that although  $n$  is `long`, there are limits imposed by R's internal allocation mechanism. These will only come into play on 64-bit systems, where the current limit for  $n$  is just under 16Gb.

### 6.1.2 User-controlled memory

The other form of memory allocation is an interface to `malloc`, the interface providing R error handling. This memory lasts until freed by the user and is additional to the memory allocated for the R workspace.

The interface functions are

```
type* Calloc(size_t n, type)
```

```
type* Realloc(any *p, size_t n, type)
```

```
void Free(any *p)
```

providing analogues of `calloc`, `realloc` and `free`. If there is an error during allocation it is handled by R, so if these routines return the memory has been successfully allocated or freed. `Free` will set the pointer  $p$  to `NULL`. (Some but not all versions of S do so.)

Users should arrange to `Free` this memory when no longer needed, including on error or user interrupt. This can often be done most conveniently from an `on.exit` action in the calling R function – see `pwilcox` for an example.

Do not assume that memory allocated by `Calloc/Realloc` comes from the same pool as used by `malloc`: in particular do not use `free` or `strdup` with it.

These entry points need to be prefixed by `R_` if `STRICT_R_HEADERS` has been defined.

## 6.2 Error handling

The basic error handling routines are the equivalents of `stop` and `warning` in R code, and use the same interface.

```
void error(const char * format, ...);
void warning(const char * format, ...);
```

These have the same call sequences as calls to `printf`, but in the simplest case can be called with a single character string argument giving the error message. (Don't do this if the string contains '%' or might otherwise be interpreted as a format.)

If `STRICT_R_HEADERS` is not defined there is also an S-compatibility interface which uses calls of the form

```
PROBLEM ..... ERROR
MESSAGE ..... WARN
PROBLEM ..... RECOVER(NULL_ENTRY)
MESSAGE ..... WARNING(NULL_ENTRY)
```

the last two being the forms available in all S versions. Here '.....' is a set of arguments to `printf`, so can be a string or a format string followed by arguments separated by commas.

### 6.2.1 Error handling from FORTRAN

There are two interface function provided to call `error` and `warning` from FORTRAN code, in each case with a simple character string argument. They are defined as

```
subroutine rexit(message)
subroutine rwarn(message)
```

Messages of more than 255 characters are truncated, with a warning.

## 6.3 Random number generation

The interface to R's internal random number generation routines is

```
double unif_rand();
double norm_rand();
double exp_rand();
```

giving one uniform, normal or exponential pseudo-random variate. However, before these are used, the user must call

```
GetRNGstate();
```

and after all the required variates have been generated, call

```
PutRNGstate();
```

These essentially read in (or create) `.Random.seed` and write it out after use.

File `'S.h'` defines `seed_in` and `seed_out` for S-compatibility rather than `GetRNGstate` and `PutRNGstate`. These take a `long *` argument which is ignored.

The random number generator is private to R; there is no way to select the kind of RNG or set the seed except by evaluating calls to the R functions.

The C code behind R's `rxxx` functions can be accessed by including the header file `'Rmath.h'`; See [Section 6.7.1 \[Distribution functions\]](#), [page 135](#). Those calls generate a single variate and should also be enclosed in calls to `GetRNGstate` and `PutRNGstate`.



In addition, there is an interface (defined in header ‘R\_ext/Applic.h’) to the generation of random 2-dimensional tables with given row and column totals using Patefield’s algorithm.

```
void rcont2 (int* nrow, int* ncol, int* nrowt, int* ncolt, int* ntotal, double* fact, int* jwork, int* matrix) [Function]
```

Here, *nrow* and *ncol* give the numbers *nr* and *nc* of rows and columns and *nrowt* and *ncolt* the corresponding row and column totals, respectively, *ntotal* gives the sum of the row (or columns) totals, *jwork* is a workspace of length *nc*, and on output *matrix* contains the *nr* \* *nc* generated random counts in the usual column-major order.

## 6.4 Missing and IEEE special values

A set of functions is provided to test for NA, Inf, -Inf and NaN. These functions are accessed *via* macros:

|             |                              |
|-------------|------------------------------|
| ISNA(x)     | True for R’s NA only         |
| ISNAN(x)    | True for R’s NA and IEEE NaN |
| R_FINITE(x) | False for Inf, -Inf, NA, NaN |

and *via* function `R_IsNaN` which is true for NaN but not NA.

Do use `R_FINITE` rather than `isfinite` or `finite`; the latter is often mendacious and `isfinite` is only available on a some platforms, on which `R_FINITE` is a macro expanding to `isfinite`.

Currently in C code `ISNAN` is a macro calling `isnan`. (Since this gives problems on some C++ systems, if the R headers is called from C++ code a function call is used.)

You can check for Inf or -Inf by testing equality to `R_PosInf` or `R_NegInf`, and set (but not test) an NA as `NA_REAL`.

All of the above apply to *double* variables only. For integer variables there is a variable accessed by the macro `NA_INTEGER` which can used to set or test for missingness.

## 6.5 Printing

The most useful function for printing from a C routine compiled into R is `Rprintf`. This is used in exactly the same way as `printf`, but is guaranteed to write to R’s output (which might be a GUI console rather than a file, and can be re-directed by `sink`). It is wise to write complete lines (including the “\n”) before returning to R. It is defined in ‘R\_ext/Print.h’.

The function `REprintf` is similar but writes on the error stream (`stderr`) which may or may not be different from the standard output stream.

Functions `Rvprintf` and `REvprintf` are analogues using the `vprintf` interface. Because that is a C99 interface, they are only defined by ‘R\_ext/Print.h’ in C++ code if the macro `R_USE_C99_IN_CXX` is defined when it is included.

Another circumstance when it may be important to use these functions is when using parallel computation on a cluster of computational nodes, as their output will be re-directed/logged appropriately.

### 6.5.1 Printing from FORTRAN

On many systems FORTRAN `write` and `print` statements can be used, but the output may not interleave well with that of C, and will be invisible on GUI interfaces. They are not portable and best avoided.

Three subroutines are provided to ease the output of information from FORTRAN code.

```
subroutine dblepr(label, nchar, data, ndata)
subroutine realpr(label, nchar, data, ndata)
subroutine intpr (label, nchar, data, ndata)
```

Here *label* is a character label of up to 255 characters, *nchar* is its length (which can be -1 if the whole label is to be used), and *data* is an array of length at least *ndata* of the appropriate type (double precision, real and integer respectively). These routines print the label on one line and then print *data* as if it were an R vector on subsequent line(s). They work with zero *ndata*, and so can be used to print a label alone.

## 6.6 Calling C from FORTRAN and vice versa

Naming conventions for symbols generated by FORTRAN differ by platform: it is not safe to assume that FORTRAN names appear to C with a trailing underscore. To help cover up the platform-specific differences there is a set of macros that should be used.

`F77_SUB(name)`

to define a function in C to be called from FORTRAN

`F77_NAME(name)`

to declare a FORTRAN routine in C before use

`F77_CALL(name)`

to call a FORTRAN routine from C

`F77_COMDECL(name)`

to declare a FORTRAN common block in C

`F77_COM(name)`

to access a FORTRAN common block from C

On most current platforms these are all the same, but it is unwise to rely on this. Note that names with underscores are not legal in FORTRAN 77, and are not portably handled by the above macros. (Also, all FORTRAN names for use by R are lower case, but this is not enforced by the macros.)

For example, suppose we want to call R's normal random numbers from FORTRAN. We need a C wrapper along the lines of

```
#include <R.h>

void F77_SUB(rndstart)(void) { GetRNGstate(); }
void F77_SUB(rndend)(void) { PutRNGstate(); }
double F77_SUB(normrnd)(void) { return norm_rand(); }
```

to be called from FORTRAN as in

```

subroutine testit()
double precision normrnd, x
call rndstart()
x = normrnd()
call dblepr("X was", 5, x, 1)
call rndend()
end

```

Note that this is not guaranteed to be portable, for the return conventions might not be compatible between the C and FORTRAN compilers used. (Passing values *via* arguments is safer.)

The standard packages, for example **stats**, are a rich source of further examples.

## 6.7 Numerical analysis subroutines

R contains a large number of mathematical functions for its own use, for example numerical linear algebra computations and special functions.

The header files ‘R\_ext/BLAS.h’, ‘R\_ext/Lapack.h’ and ‘R\_ext/Linpack.h’ contains declarations of the BLAS, LAPACK and LINPACK/EISPACK linear algebra functions included in R. These are expressed as calls to FORTRAN subroutines, and they will also be usable from users’ FORTRAN code. Although not part of the official API, this set of subroutines is unlikely to change (but might be supplemented).

The header file ‘Rmath.h’ lists many other functions that are available and documented in the following subsections. Many of these are C interfaces to the code behind R functions, so the R function documentation may give further details.

### 6.7.1 Distribution functions

The routines used to calculate densities, cumulative distribution functions and quantile functions for the standard statistical distributions are available as entry points.

The arguments for the entry points follow the pattern of those for the normal distribution:

```

double dnorm(double x, double mu, double sigma, int give_log);
double pnorm(double x, double mu, double sigma, int lower_tail,
             int give_log);
double qnorm(double p, double mu, double sigma, int lower_tail,
             int log_p);
double rnorm(double mu, double sigma);

```

That is, the first argument gives the position for the density and CDF and probability for the quantile function, followed by the distribution’s parameters. Argument *lower\_tail* should be TRUE (or 1) for normal use, but can be FALSE (or 0) if the probability of the upper tail is desired or specified.

Finally, *give\_log* should be non-zero if the result is required on log scale, and *log\_p* should be non-zero if *p* has been specified on log scale.

Note that you directly get the cumulative (or “integrated”) *hazard* function,  $H(t) = -\log(1 - F(t))$ , by using

```
- pdist(t, ..., /*lower_tail = */ FALSE, /* give_log = */ TRUE)
```

or shorter (and more cryptic) `- pdist(t, ..., 0, 1)`.

The random-variate generation routine `rnorm` returns one normal variate. See [Section 6.3 \[Random numbers\]](#), [page 132](#), for the protocol in using the random-variate routines.

Note that these argument sequences are (apart from the names and that `rnorm` has no *n*) mainly the same as the corresponding R functions of the same name, so the documentation of the R functions can be used. Note that the exponential and gamma distributions are parametrized by `scale` rather than `rate`.

For reference, the following table gives the basic name (to be prefixed by ‘d’, ‘p’, ‘q’ or ‘r’ apart from the exceptions noted) and distribution-specific arguments for the complete set of distributions.

|                         |                        |  |
|-------------------------|------------------------|--|
| beta                    | <code>beta</code>      | <code>a, b</code>                        |
| non-central beta        | <code>nbeta</code>     | <code>a, b, ncp</code>                   |
| binomial                | <code>binom</code>     | <code>n, p</code>                        |
| Cauchy                  | <code>cauchy</code>    | <code>location, scale</code>             |
| chi-squared             | <code>chisq</code>     | <code>df</code>                          |
| non-central chi-squared | <code>nchisq</code>    | <code>df, ncp</code>                     |
| exponential             | <code>exp</code>       | <code>scale (and <b>not</b> rate)</code> |
| F                       | <code>f</code>         | <code>n1, n2</code>                      |
| non-central F           | <code>nf</code>        | <code>n1, n2, ncp</code>                 |
| gamma                   | <code>gamma</code>     | <code>shape, scale</code>                |
| geometric               | <code>geom</code>      | <code>p</code>                           |
| hypergeometric          | <code>hyper</code>     | <code>NR, NB, n</code>                   |
| logistic                | <code>logis</code>     | <code>location, scale</code>             |
| lognormal               | <code>lnorm</code>     | <code>logmean, logsd</code>              |
| negative binomial       | <code>nbinom</code>    | <code>size, prob</code>                  |
| normal                  | <code>norm</code>      | <code>mu, sigma</code>                   |
| Poisson                 | <code>pois</code>      | <code>lambda</code>                      |
| Student’s t             | <code>t</code>         | <code>n</code>                           |
| non-central t           | <code>nt</code>        | <code>df, delta</code>                   |
| Studentized range       | <code>tukey (*)</code> | <code>rr, cc, df</code>                  |
| uniform                 | <code>unif</code>      | <code>a, b</code>                        |
| Weibull                 | <code>weibull</code>   | <code>shape, scale</code>                |
| Wilcoxon rank sum       | <code>wilcox</code>    | <code>m, n</code>                        |
| Wilcoxon signed rank    | <code>signrank</code>  | <code>n</code>                           |

Entries marked with an asterisk only have ‘p’ and ‘q’ functions available, and none of the non-central distributions have ‘r’ functions. After a call to `dwilcox`, `pwilcox` or `qwilcox` the function `wilcox_free()` should be called, and similarly for the signed rank functions.

### 6.7.2 Mathematical functions

|   |            |
|---|------------|
| <code>double gammafn (double x)</code>    | [Function] |
| <code>double lgammafn (double x)</code>   | [Function] |
| <code>double digamma (double x)</code>    | [Function] |
| <code>double trigamma (double x)</code>   | [Function] |
| <code>double tetragamma (double x)</code> | [Function] |
| <code>double pentagamma (double x)</code> | [Function] |

`double psigamma (double x, double deriv)` [Function]  
 The Gamma function, the natural logarithm of its absolute value and first four derivatives and the  $n$ -th derivative of Psi, the digamma function, which is the derivative of `lgammafn`. In other words, `digamma(x)` is the same as `(psigamma(x,0), trigamma(x) == psigamma(x,1), etc.`

`double beta (double a, double b)` [Function]  
`double lbeta (double a, double b)` [Function]  
 The (complete) Beta function and its natural logarithm.

`double choose (double n, double k)` [Function]  
`double lchoose (double n, double k)` [Function]  
 The number of combinations of  $k$  items chosen from  $n$  and the natural logarithm of its absolute value, generalized to arbitrary real  $n$ .  $k$  is rounded to the nearest integer (with a warning if needed).

`double bessel_i (double x, double nu, double expo)` [Function]  
`double bessel_j (double x, double nu)` [Function]  
`double bessel_k (double x, double nu, double expo)` [Function]  
`double bessel_y (double x, double nu)` [Function]  
 Bessel functions of types I, J, K and Y with index  $nu$ . For `bessel_i` and `bessel_k` there is the option to return  $\exp(-x) I(x; nu)$  or  $\exp(x) K(x; nu)$  if `expo` is 2. (Use `expo == 1` for unscaled values.)

### 6.7.3 Numerical Utilities

There are a few other numerical utility functions available as entry points.

`double R_pow (double x, double y)` [Function]  
`double R_pow_di (double x, int i)` [Function]  
`R_pow(x, y)` and `R_pow_di(x, i)` compute  $x^y$  and  $x^i$ , respectively using R-FINITE checks and returning the proper result (the same as R) for the cases where  $x$ ,  $y$  or  $i$  are 0 or missing or infinite or NaN.

`double log1p (double x)` [Function]  
 Computes  $\log(1 + x)$  (*log 1 plus x*), accurately even for small  $x$ , i.e.,  $|x| \ll 1$ .  
 This should be provided by your platform, in which case it is not included in ‘`Rmath.h`’, but is (probably) in ‘`math.h`’ which ‘`Rmath.h`’ includes.

`double log1pmx (double x)` [Function]  
 Computes  $\log(1 + x) - x$  (*log 1 plus x minus x*), accurately even for small  $x$ , i.e.,  $|x| \ll 1$ .

`double expm1 (double x)` [Function]  
 Computes  $\exp(x) - 1$  (*exp x minus 1*), accurately even for small  $x$ , i.e.,  $|x| \ll 1$ .  
 This should be provided by your platform, in which case it is not included in ‘`Rmath.h`’, but is (probably) in ‘`math.h`’ which ‘`Rmath.h`’ includes.

`double lgamma1p (double x)` [Function]  
 Computes  $\log(\Gamma(x + 1))$  (*log(gamma(1 plus x))*), accurately even for small  $x$ , i.e.,  $0 < x < 0.5$ .

```
double logspace_add (double logx, double logy) [Function]
double logspace_sub (double logx, double logy) [Function]
    Compute the log of a sum or difference from logs of terms, i.e., “x + y” as log
    (exp(logx) + exp(logy)) and “x - y” as log (exp(logx) - exp(logy)), without
    causing unnecessary overflows or throwing away too much accuracy.

int  imax2 (int x, int y) [Function]
int  imin2 (int x, int y) [Function]
double fmax2 (double x, double y) [Function]
double fmin2 (double x, double y) [Function]
    Return the larger (max) or smaller (min) of two integer or double numbers, respec-
    tively. Note that fmax2 and fmin2 differ from C99’s fmax and fmin when one of the
    arguments is a NaN: these versions return NaN.

double sign (double x) [Function]
    Compute the signum function, where sign(x) is 1, 0, or -1, when x is positive, 0, or
    negative, respectively, and NaN if x is a NaN.

double fsign (double x, double y) [Function]
    Performs “transfer of sign” and is defined as  $|x| * \text{sign}(y)$ .

double fprec (double x, double digits) [Function]
    Returns the value of x rounded to digits decimal digits (after the decimal point).
    This is the function used by R’s round().

double fround (double x, double digits) [Function]
    Returns the value of x rounded to digits significant decimal digits.
    This is the function used by R’s signif().

double ftrunc (double x) [Function]
    Returns the value of x truncated (to an integer value) towards zero.
```

### 6.7.4 Mathematical constants

R has a set of commonly used mathematical constants encompassing constants usually found ‘math.h’ and contains further ones that are used in statistical computations. All these are defined to (at least) 30 digits accuracy in ‘Rmath.h’. The following definitions use  $\ln(x)$  for the natural logarithm ( $\log(x)$  in R).

| Name     | Definition ( $\ln = \log$ ) | round(value, 7) |
|----------|-----------------------------|-----------------|
| M_E      | $e$                         | 2.7182818       |
| M_LOG2E  | $\log_2(e)$                 | 1.4426950       |
| M_LOG10E | $\log_{10}(e)$              | 0.4342945       |
| M_LN2    | $\ln(2)$                    | 0.6931472       |
| M_LN10   | $\ln(10)$                   | 2.3025851       |
| M_PI     | $\pi$                       | 3.1415927       |
| M_PI_2   | $\pi/2$                     | 1.5707963       |
| M_PI_4   | $\pi/4$                     | 0.7853982       |
| M_1_PI   | $1/\pi$                     | 0.3183099       |

|                |                     |           |
|----------------|---------------------|-----------|
| M_2_PI         | $2/\pi$             | 0.6366198 |
| M_2_SQRTPI     | $2/\sqrt{\pi}$      | 1.1283792 |
| M_SQRT2        | $\sqrt{2}$          | 1.4142136 |
| M_SQRT1_2      | $1/\sqrt{2}$        | 0.7071068 |
| M_SQRT_3       | $\sqrt{3}$          | 1.7320508 |
| M_SQRT_32      | $\sqrt{32}$         | 5.6568542 |
| M_LOG10_2      | $\log_{10}(2)$      | 0.3010300 |
| M_2PI          | $2\pi$              | 6.2831853 |
| M_SQRT_PI      | $\sqrt{\pi}$        | 1.7724539 |
| M_1_SQRT_2PI   | $1/\sqrt{2\pi}$     | 0.3989423 |
| M_SQRT_2dPI    | $\sqrt{2/\pi}$      | 0.7978846 |
| M_LN_SQRT_PI   | $\ln(\sqrt{\pi})$   | 0.5723649 |
| M_LN_SQRT_2PI  | $\ln(\sqrt{2\pi})$  | 0.9189385 |
| M_LN_SQRT_PId2 | $\ln(\sqrt{\pi/2})$ | 0.2257914 |

There are a set of constants (PI, DOUBLE\_EPS) (and so on) defined (unless STRICT\_R\_HEADERS is defined) in the included header ‘R\_ext/Constants.h’, mainly for compatibility with S.

Further, the included header ‘R\_ext/Boolean.h’ has constants TRUE and FALSE = 0 of type Rboolean in order to provide a way of using “logical” variables in C consistently.

## 6.8 Optimization

The C code underlying `optim` can be accessed directly. The user needs to supply a function to compute the function to be minimized, of the type

```
typedef double optimfn(int n, double *par, void *ex);
```

where the first argument is the number of parameters in the second argument. The third argument is a pointer passed down from the calling routine, normally used to carry auxiliary information.

Some of the methods also require a gradient function

```
typedef void optimgr(int n, double *par, double *gr, void *ex);
```

which passes back the gradient in the `gr` argument. No function is provided for finite-differencing, nor for approximating the Hessian at the result.

The interfaces (defined in header ‘R\_ext/Applic.h’) are

- Nelder Mead:

```
void nmmin(int n, double *xin, double *x, double *Fmin, optimfn fn,
           int *fail, double abstol, double intol, void *ex,
           double alpha, double beta, double gamma, int trace,
           int *fnccount, int maxit);
```

- BFGS:

```
void vmmin(int n, double *x, double *Fmin,
           optimfn fn, optimgr gr, int maxit, int trace,
           int *mask, double abstol, double reltol, int nREPORT,
           void *ex, int *fnccount, int *grcount, int *fail);
```

- Conjugate gradients:



```
void cgmin(int n, double *xin, double *x, double *Fmin,
           optimfn fn, optimgr gr, int *fail, double abstol,
           double intol, void *ex, int type, int trace,
           int *fnccount, int *grccount, int maxit);
```

- Limited-memory BFGS with bounds:

```
void lbfgsb(int n, int lmm, double *x, double *lower,
            double *upper, int *nbd, double *Fmin, optimfn fn,
            optimgr gr, int *fail, void *ex, double factr,
            double pgtol, int *fnccount, int *grccount,
            int maxit, char *msg, int trace, int nREPORT);
```

- Simulated annealing:

```
void samini(int n, double *x, double *Fmin, optimfn fn, int maxit,
            int tmax, double temp, int trace, void *ex);
```

Many of the arguments are common to the various methods. `n` is the number of parameters, `x` or `xin` is the starting parameters on entry and `x` the final parameters on exit, with final value returned in `Fmin`. Most of the other parameters can be found from the help page for `optim`: see the source code ‘`src/appl/lbfgsb.c`’ for the values of `nbd`, which specifies which bounds are to be used.

## 6.9 Integration

The C code underlying `integrate` can be accessed directly. The user needs to supply a *vectorizing* C function to compute the function to be integrated, of the type

```
typedef void integr_fn(double *x, int n, void *ex);
```

where `x[]` is both input and output and has length `n`, i.e., a C function, say `fn`, of type `integr_fn` must basically do `for(i in 1:n) x[i] := f(x[i], ex)`. The vectorization requirement can be used to speed up the integrand instead of calling it `n` times. Note that in the current implementation built on QUADPACK, `n` will be either 15 or 21. The `ex` argument is a pointer passed down from the calling routine, normally used to carry auxiliary information.

There are interfaces (defined in header ‘`R_ext/Applic.h`’) for definite and for indefinite integrals. ‘Indefinite’ means that at least one of the integration boundaries is not finite.

- Finite:

```
void Rdqags(integr_fn f, void *ex, double *a, double *b,
            double *epsabs, double *epsrel,
            double *result, double *abserr, int *neval, int *ier,
            int *limit, int *lenw, int *last,
            int *iwork, double *work);
```

- Indefinite:

```
void Rdqagi(integr_fn f, void *ex, double *bound, int *inf,
            double *epsabs, double *epsrel,
            double *result, double *abserr, int *neval, int *ier,
            int *limit, int *lenw, int *last,
            int *iwork, double *work);
```

Only the 3rd and 4th argument differ for the two integrators; for the definite integral, using `Rdqags`, `a` and `b` are the integration interval bounds, whereas for an indefinite integral, using `Rdqagi`, `bound` is the finite bound of the integration (if the integral is not doubly-infinite) and `inf` is a code indicating the kind of integration range,

`inf = 1` corresponds to `(bound, +Inf)`,

`inf = -1` corresponds to `(-Inf, bound)`,

`inf = 2` corresponds to `(-Inf, +Inf)`,

`f` and `ex` define the integrand function, see above; `epsabs` and `epsrel` specify the absolute and relative accuracy requested, `result`, `abserr` and `last` are the output components `value`, `abs.err` and `subdivisions` of the R function `integrate`, where `neval` gives the number of integrand function evaluations, and the error code `ier` is translated to R's `integrate()` `$message`, look at that function definition. `limit` corresponds to `integrate(..., subdivisions = *)`. It seems you should always define the two work arrays and the length of the second one as

```
lenw = 4 * limit;
iwork = (int *) R_alloc(limit, sizeof(int));
work = (double *) R_alloc(lenw, sizeof(double));
```

The comments in the source code in `'src/appl/integrate.c'` give more details, particularly about reasons for failure (`ier >= 1`).

## 6.10 Utility functions

R has a fairly comprehensive set of sort routines which are made available to users' C code. These are declared in header file `'R_ext/Utils.h'` (included by `'R.h'`) and include the following.

```
void R_isort (int* x, int n) [Function]
void R_rsort (double* x, int n) [Function]
void R_csort (Rcomplex* x, int n) [Function]
void rsort_with_index (double* x, int* index, int n) [Function]
```

The first three sort integer, real (double) and complex data respectively. (Complex numbers are sorted by the real part first then the imaginary part.) NAs are sorted last.

`rsort_with_index` sorts on `x`, and applies the same permutation to `index`. NAs are sorted last.

```
void revsort (double* x, int* index, int n) [Function]
```

Is similar to `rsort_with_index` but sorts into decreasing order, and NAs are not handled.

```
void iPsort (int* x, int n, int k) [Function]
void rPsort (double* x, int n, int k) [Function]
void cPsort (Rcomplex* x, int n, int k) [Function]
```

These all provide (very) partial sorting: they permute `x` so that `x[k]` is in the correct place with smaller values to the left, larger ones to the right.

```

void R_qsort (double *v, int i, int j) [Function]
void R_qsort_I (double *v, int *I, int i, int j) [Function]
void R_qsort_int (int *iv, int i, int j) [Function]
void R_qsort_int_I (int *iv, int *I, int i, int j) [Function]

```

These routines sort  $v[i:j]$  or  $iv[i:j]$  (using 1-indexing, i.e.,  $v[1]$  is the first element) calling the quicksort algorithm as used by R's `sort(v, method = "quick")` and documented on the help page for the R function `sort`. The `..._I()` versions also return the `sort.index()` vector in  $I$ . Note that the ordering is *not* stable, so tied values may be permuted.

Note that NAs are not handled (explicitly) and you should use different sorting functions if NAs can be present.

```

subroutine qsort4 (double precision v, integer indx, integer ii, integer jj) [Function]

```

```

subroutine qsort3 (double precision v, integer ii, integer jj) [Function]

```

The FORTRAN interface routines for sorting double precision vectors are `qsort3` and `qsort4`, equivalent to `R_qsort` and `R_qsort_I`, respectively.

```

void R_max_col (double* matrix, int* nr, int* nc, int* maxes, int* ties_meth) [Function]

```

Given the  $nr$  by  $nc$  matrix `matrix` in column-major ("FORTRAN") order, `R_max_col()` returns in `maxes[i-1]` the column number of the maximal element in the  $i$ -th row (the same as R's `max.col()` function). In the case of ties (multiple maxima), `*ties_meth` is an integer code in 1:3 determining the method: 1 = "random", 2 = "first" and 3 = "last". See R's help page `?max.col`.

```

int findInterval (double* xt, int n, double x, Rboolean rightmost_closed, Rboolean all_inside, int ilo, int* mflag) [Function]

```

Given the ordered vector `xt` of length  $n$ , return the interval or index of  $x$  in `xt[]`, typically  $\max(i; 1 \leq i \leq n \ \& \ xt[i] \leq x)$  where we use 1-indexing as in R and FORTRAN (but not C). If `rightmost_closed` is true, also returns  $n - 1$  if  $x$  equals `xt[n]`. If `all_inside` is not 0, the result is coerced to lie in  $1:(n-1)$  even when  $x$  is outside the `xt[]` range. On return, `*mflag` equals  $-1$  if  $x < xt[1]$ ,  $+1$  if  $x \geq xt[n]$ , and 0 otherwise.

The algorithm is particularly fast when `ilo` is set to the last result of `findInterval()` and  $x$  is a value of a sequence which is increasing or decreasing for subsequent calls.

There is also an `F77_CALL(interv)()` version of `findInterval()` with the same arguments, but all pointers.

The following two functions do *numerical* colorspace conversion from HSV to RGB and back. Note that all colours must be in  $[0,1]$ .

```

void hsv2rgb (double h, double s, double v, double *r, double *g, double *b) [Function]

```

```

void rgb2hsv (double r, double g, double b, double *h, double *s, double *v) [Function]

```

A system-independent interface to produce the name of a temporary file is provided as

```
char * R_tmpnam (const char *prefix, const char *tmpdir) [Function]
char * R_tmpnam2 (const char *prefix, const char *tmpdir, const char *fileext) [Function]
```

Return a pathname for a temporary file with name beginning with *prefix* and ending with *fileext* in directory *tmpdir*. A NULL prefix or extension is replaced by "". Note that the return value is **malloced** and should be **freed** when no longer needed (unlike the system call `tmpnam`).

There is also the internal function used to expand file names in several R functions, and called directly by `path.expand`.

```
const char * R_ExpandFileName (const char *fn) [Function]
```

Expand a path name *fn* by replacing a leading tilde by the user's home directory (if defined). The precise meaning is platform-specific; it will usually be taken from the environment variable `HOME` if this is defined.

## 6.11 Re-encoding

R has its own C-level interface to the encoding conversion capabilities provided by `iconv` because there are incompatibilities between the declarations in different implementations of `iconv`.

These are declared in header file '`R_ext/Riconv.h`'.

```
void *Riconv_open (const char *to, const char *from) [Function]
```

Set up a pointer to an encoding object to be used to convert between two encodings: "" indicates the current locale.

```
size_t Riconv (void *cd, const char **inbuf, size_t *inbytesleft, [Function]
               char **outbuf, size_t *outbytesleft)
```

Convert as much as possible of *inbuf* to *outbuf*. Initially the `int` variables indicate the number of bytes available in the buffers, and they are updated (and the `char` pointers are updated to point to the next free byte in the buffer). The return value is the number of characters converted, or `(size_t)-1` (beware: `size_t` is usually an unsigned type). It should be safe to assume that an error condition sets `errno` to one of `E2BIG` (the output buffer is full), `EILSEQ` (the input cannot be converted, and might be invalid in the encoding specified) or `EINVAL` (the input does not end with a complete multi-byte character).

```
int Riconv_close (void * cd) [Function]
```

Free the resources of an encoding object.

## 6.12 Allowing interrupts

No port of R can be interrupted whilst running long computations in compiled code, so programmers should make provision for the code to be interrupted at suitable points by calling from C

```
#include <R_ext/Utils.h>
```

```
void R_CheckUserInterrupt(void);
```

and from FORTRAN

```
subroutine rchkusr()
```

These check if the user has requested an interrupt, and if so branch to R's error handling functions.

Note that it is possible that the code behind one of the entry points defined here if called from your C or FORTRAN code could be interruptible or generate an error and so not return to your code.

## 6.13 Platform and version information

The header files define `USING_R`, which can be used to test if the code is indeed being used with R.

Header file `'Rconfig.h'` (included by `'R.h'`) is used to define platform-specific macros that are mainly for use in other header files. The macro `WORDS_BIGENDIAN` is defined on big-endian<sup>2</sup> systems (e.g. most OSes on Sparc and PowerPC hardware) and not on little-endian systems (such as `i686` and `x86_64` on all OSes, and Linux on Alpha and Itanium). It can be useful when manipulating binary files. The macro `SUPPORT_OPENMP` is defined on suitable systems as from R 2.13.0, and can be used in conjunction with the `SUPPORT_OPENMP_*` macros in packages that want to make use of OpenMP.

Header file `'Rversion.h'` (**not** included by `'R.h'`) defines a macro `R_VERSION` giving the version number encoded as an integer, plus a macro `R_Version` to do the encoding. This can be used to test if the version of R is late enough, or to include back-compatibility features. For protection against very old versions of R which did not have this macro, use a construction such as

```
#if defined(R_VERSION) && R_VERSION >= R_Version(1, 9, 0)
...
#endif
```

More detailed information is available in the macros `R_MAJOR`, `R_MINOR`, `R_YEAR`, `R_MONTH` and `R_DAY`: see the header file `'Rversion.h'` for their format. Note that the minor version includes the patchlevel (as in `'9.0'`).

## 6.14 Inlining C functions

The C99 keyword `inline` should be recognized by all compilers now used to build R. Portable code which might be used with earlier versions of R can be written using the macro `R_INLINE` (defined in file `'Rconfig.h'` included by `'R.h'`), as for example from package **cluster**

```
#include <R.h>

static R_INLINE int ind_2(int l, int j)
{
...
}
```

Be aware that using inlining with functions in more than one compilation unit is almost impossible to do portably, see <http://www.greenend.org.uk/rjk/2003/03/inline.html>,

<sup>2</sup> <http://en.wikipedia.org/wiki/Endianness>.

so this usage is for `static` functions as in the example. All the R configure code has checked is that `R_INLINE` can be used in a single C file with the compiler used to build R. We recommend that packages making extensive use of inlining include their own configure code.

## 6.15 Controlling visibility

Header `'R_ext/Visibility'` has some definitions for controlling the visibility of entry points. These are only effective when `'HAVE_VISIBILITY_ATTRIBUTE'` is defined – this is checked when R is configured and recorded in header `'Rconfig.h'` (included by `'R_ext/Visibility.h'`). It is generally defined on modern Unix-alikes with a recent compiler (e.g. `gcc4`), but not supported on Windows. Minimizing the visibility of symbols in a shared library will both speed up its loading (unlikely to be significant) and reduce the possibility of linking to the wrong entry points of the same name.

C/C++ entry points prefixed by `attribute_hidden` will not be visible in the shared object. There is no comparable mechanism for FORTRAN entry points, but there is a more comprehensive scheme used by, for example package **stats**. Most compilers which allow control of visibility will allow control of visibility for all symbols *via* a flag, and where known the flag is encapsulated in the macros `'C_VISIBILITY'` and `F77_VISIBILITY` for C and FORTRAN compilers. These are defined in `'etc/Makeconf'` and so available for normal compilation of package code. For example, `'src/Makevars'` could include

```
PKG_CFLAGS=$(C_VISIBILITY)
PKG_FFLAGS=$(F77_VISIBILITY)
```

This would end up with **no** visible entry points, which would be pointless. However, the effect of the flags can be overridden by using the `attribute_visible` prefix. A shared object which registers its entry points needs only for have one visible entry point, its initializer, so for example package **stats** has

```
void attribute_visible R_init_stats(DllInfo *dll)
{
    R_registerRoutines(dll, CEntries, CallEntries, FortEntries, NULL);
    R_useDynamicSymbols(dll, FALSE);
    ...
}
```

The visibility mechanism is not available on Windows, but there is an equally effective way to control which entry points are visible, by supplying a definitions file `'pkgname/src/pkgname-win.def'`: only entry points listed in that file will be visible. Again using **stats** as an example, it has

```
LIBRARY stats.dll
EXPORTS
    R_init_stats
```

## 6.16 Using these functions in your own C code

It is possible to build **Mathlib**, the R set of mathematical functions documented in `'Rmath.h'`, as a standalone library `'libRmath'` under both Unix-alikes and Windows. (This includes the functions documented in [Section 6.7 \[Numerical analysis subroutines\]](#), [page 135](#) as from that header file.)

The library is not built automatically when R is installed, but can be built in the directory ‘src/nmath/standalone’ in the R sources: see the file ‘README’ there. To use the code in your own C program include

```
#define MATHLIB_STANDALONE
#include <Rmath.h>
```

and link against ‘-lRmath’ (and perhaps ‘-lm’. There is an example file ‘test.c’.

A little care is needed to use the random-number routines. You will need to supply the uniform random number generator

```
double unif_rand(void)
```

or use the one supplied (and with a dynamic library or DLL you will have to use the one supplied, which is the Marsaglia-multicarry with an entry points

```
set_seed(unsigned int, unsigned int)
```

to set its seeds and

```
get_seed(unsigned int *, unsigned int *)
```

to read the seeds).

## 6.17 Organization of header files

The header files which R installs are in directory ‘R\_INCLUDE\_DIR’ (default ‘R\_HOME/include’). This currently includes

|                       |   |
|-----------------------|---|
| ‘R.h’                 | includes many other files   |
| ‘S.h’                 | different version for code ported from S  |
| ‘Rinternals.h’        | definitions for using R’s internal structures   |
| ‘Rdefines.h’          | macros for an S-like interface to the above   |
| ‘Rmath.h’             | standalone math library   |
| ‘Rversion.h’          | R version information   |
| ‘Rinterface.h’        | for add-on front-ends (Unix-alikes only)  |
| ‘Rembedded.h’         | for add-on front-ends   |
| ‘R_ext/Applic.h’      | optimization and integration  |
| ‘R_ext/BLAS.h’        | C definitions for BLAS routines   |
| ‘R_ext/Callbacks.h’   | C (and R function) top-level task handlers  |
| ‘R_ext/GetX11Image.h’ | X11Image interface used by package <b>trkplot</b>   |
| ‘R_ext/Lapack.h’      | C definitions for some LAPACK routines  |
| ‘R_ext/Linpack.h’     | C definitions for some LINPACK routines, not all of which are included in R                       |
| ‘R_ext/Parse.h’       | a small part of R’s parse interface   |
| ‘R_ext/RConvertors.h’ |   |
| ‘R_ext/Rdynload.h’    | needed to register compiled code in packages  |
| ‘R_ext/R-ftp-http.h’  | interface to internal method of <b>download.file</b>  |
| ‘R_ext/Riconv.h’      | interface to <b>iconv</b>   |
| ‘R_ext/RStartup.h’    | for add-on front-ends   |
| ‘R_ext/Visibility.h’  | definitions controlling visibility  |
| ‘R_ext/eventloop.h’   | for add-on front-ends and for packages that need to share in the R event loops (on all platforms) |

The following headers are included by ‘R.h’:



|                                  |  |
|----------------------------------|--|
| <code>'Rconfig.h'</code>         | configuration info that is made available  |
| <code>'R_ext/Arith.h'</code>     | handling for NAs, NaNs, Inf/-Inf   |
| <code>'R_ext/Boolean.h'</code>   | TRUE/FALSE type  |
| <code>'R_ext/Complex.h'</code>   | C typedefs for R's complex   |
| <code>'R_ext/Constants.h'</code> | constants  |
| <code>'R_ext/Error.h'</code>     | error handling   |
| <code>'R_ext/Memory.h'</code>    | memory allocation  |
| <code>'R_ext/Print.h'</code>     | Rprintf and variations.  |
| <code>'R_ext/Random.h'</code>    | random number generation   |
| <code>'R_ext/RS.h'</code>        | definitions common to <code>'R.h'</code> and <code>'S.h'</code> , including<br>F77_CALL etc. |
| <code>'R_ext/Utils.h'</code>     | sorting and other utilities  |
| <code>'R_ext/libextern.h'</code> | definitions for exports from <code>'R.dll'</code> on<br>Windows.                             |

The graphics systems are exposed in headers `'R_ext/GraphicsEngine.h'`, `'R_ext/GraphicsDevice.h'` (which it includes) and `'R_ext/QuartzDevice.h'`. Some entry points from the `stats` package are in `'R_ext/stats_package.h'` (currently related to the internals of `nls` and `nlminb`).

## 7 Generic functions and methods

R programmers will often want to add methods for existing generic functions, and may want to add new generic functions or make existing functions generic. In this chapter we give guidelines for doing so, with examples of the problems caused by not adhering to them.

This chapter only covers the ‘informal’ class system copied from S3, and not with the S4 (formal) methods of package **methods**.

The key function for methods is `NextMethod`, which dispatches the next method. It is quite typical for a method function to make a few changes to its arguments, dispatch to the next method, receive the results and modify them a little. An example is

```
t.data.frame <- function(x)
{
  x <- as.matrix(x)
  NextMethod("t")
}
```

Also consider `predict.glm`: it happens that in R for historical reasons it calls `predict.lm` directly, but in principle (and in S originally and currently) it could use `NextMethod`. (`NextMethod` seems under-used in the R sources. Do be aware that there are S/R differences in this area, and the example above works because there is a *next* method, the default method, not that a new method is selected when the class is changed.)

Any method a programmer writes may be invoked from another method by `NextMethod`, with the arguments appropriate to the previous method. Further, the programmer cannot predict which method `NextMethod` will pick (it might be one not yet dreamt of), and the end user calling the generic needs to be able to pass arguments to the next method. For this to work

*A method must have all the arguments of the generic, including ... if the generic does.*

It is a grave misunderstanding to think that a method needs only to accept the arguments it needs. The original S version of `predict.lm` did not have a `...` argument, although `predict` did. It soon became clear that `predict.glm` needed an argument `dispersion` to handle over-dispersion. As `predict.lm` had neither a `dispersion` nor a `...` argument, `NextMethod` could no longer be used. (The legacy, two direct calls to `predict.lm`, lives on in `predict.glm` in R, which is based on the workaround for S3 written by Venables & Ripley.)

Further, the user is entitled to use positional matching when calling the generic, and the arguments to a method called by `UseMethod` are those of the call to the generic. Thus

*A method must have arguments in exactly the same order as the generic.*

To see the scale of this problem, consider the generic function `scale`, defined as

```
scale <- function (x, center = TRUE, scale = TRUE)
  UseMethod("scale")
```

Suppose an unthinking package writer created methods such as

```
scale.foo <- function(x, scale = FALSE, ...) { }
```

Then for `x` of class `"foo"` the calls

```
scale(x, , TRUE)
scale(x, scale = TRUE)
```

would do most likely do different things, to the justifiable consternation of the end user.

To add a further twist, which default is used when a user calls `scale(x)` in our example? What if

```
scale.bar <- function(x, center, scale = TRUE) NextMethod("scale")
```

and `x` has class `c("bar", "foo")`? It is the default specified in the method that is used, but the default specified in the generic may be the one the user sees. This leads to the recommendation:

*If the generic specifies defaults, all methods should use the same defaults.*

An easy way to follow these recommendations is to always keep generics simple, e.g.

```
scale <- function(x, ...) UseMethod("scale")
```

Only add parameters and defaults to the generic if they make sense in all possible methods implementing it.

## 7.1 Adding new generics

When creating a new generic function, bear in mind that its argument list will be the maximal set of arguments for methods, including those written elsewhere years later. So choosing a good set of arguments may well be an important design issue, and there need to be good arguments *not* to include a `...` argument.

If a `...` argument is supplied, some thought should be given to its position in the argument sequence. Arguments which follow `...` must be named in calls to the function, and they must be named in full (partial matching is suppressed after `...`). Formal arguments before `...` can be partially matched, and so may ‘swallow’ actual arguments intended for `...`. Although it is commonplace to make the `...` argument the last one, that is not always the right choice.

Sometimes package writers want to make generic a function in the base package, and request a change in R. This may be justifiable, but making a function generic with the old definition as the default method does have a small performance cost. It is never necessary, as a package can take over a function in the base package and make it generic by something like

```
foo <- function(object, ...) UseMethod("foo")
foo.default <- function(object, ...) base::foo(object)
```

Earlier versions of this manual suggested assigning `foo.default <- base::foo`. This is **not** a good idea, as it captures the base function at the time of installation and it might be changed as R is patched or updated.

The same idea can be applied for functions in other packages with namespaces.

## 8 Linking GUIs and other front-ends to R

There are a number of ways to build front-ends to R: we take this to mean a GUI or other application that has the ability to submit commands to R and perhaps to receive results back (not necessarily in a text format). There are other routes besides those described here, for example the package **Rserve** (from CRAN, see also <http://www.rforge.net/Rserve/>) and connections to Java in ‘SJava’ (see <http://www.omegahat.org/RSJava/>) and ‘JRI’ (part of the **rJava** package on CRAN).

### 8.1 Embedding R under Unix-alikes

R can be built as a shared library<sup>1</sup> if configured with ‘--enable-R-shlib’. This shared library can be used to run R from alternative front-end programs. We will assume this has been done for the rest of this section. Also, it can be built as a static library if configured with ‘--enable-R-static-lib’, and this can be used in a very similar way.

The command-line R front-end, ‘*R\_HOME*/bin/exec/R’ is one such example, and the former GNOME (see package **gnomeGUI** on CRAN’s ‘Archive’ area) and Mac OS X consoles are others. The source for ‘*R\_HOME*/bin/exec/R’ is in file ‘src/main/Rmain.c’ and is very simple

```
int Rf_initialize_R(int ac, char **av); /* in ../unix/system.c */
void Rf_mainloop();                    /* in main.c */

extern int R_running_as_main_program;  /* in ../unix/system.c */

int main(int ac, char **av)
{
    R_running_as_main_program = 1;
    Rf_initialize_R(ac, av);
    Rf_mainloop(); /* does not return */
    return 0;
}
```

indeed, misleadingly simple. Remember that ‘*R\_HOME*/bin/exec/R’ is run from a shell script ‘*R\_HOME*/bin/R’ which sets up the environment for the executable, and this is used for

- Setting *R\_HOME* and checking it is valid, as well as the path *R\_SHARE\_DIR* and *R\_DOC\_DIR* to the installed ‘share’ and ‘doc’ directory trees. Also setting *R\_ARCH* if needed.
- Setting *LD\_LIBRARY\_PATH* to include the directories used in linking R. This is recorded as the default setting of *R\_LD\_LIBRARY\_PATH* in the shell script ‘*R\_HOME*/etc/*R\_ARCH*/ldpaths’.
- Processing some of the arguments, for example to run R under a debugger and to launch alternative front-ends to provide GUIs.

The first two of these can be achieved for your front-end by running it *via* R CMD. So, for example

---

<sup>1</sup> In the parlance of Mac OS X this is a *dynamic* library, and is the normal way to build R on that platform.

```
R CMD /usr/local/lib/R/bin/exec/R
R CMD exec/R
```

will both work in a standard R installation. (R CMD looks first for executables in ‘R\_HOME/bin’.) If you do not want to run your front-end in this way, you need to ensure that R\_HOME is set and LD\_LIBRARY\_PATH is suitable. (The latter might well be, but modern Unix/Linux systems do not normally include ‘/usr/local/lib’ (‘/usr/local/lib64’ on some architectures), and R does look there for system components.)

The other senses in which this example is too simple are that all the internal defaults are used and that control is handed over to the R main loop. There are a number of small examples<sup>2</sup> in the ‘tests/Embedding’ directory. These make use of Rf\_initEmbeddedR in ‘src/main/Rembedded.c’, and essentially use

```
#include <Rembedded.h>

int main(int ac, char **av)
{
    /* do some setup */
    Rf_initEmbeddedR(argc, argv);
    /* do some more setup */

    /* submit some code to R, which is done interactively via
       run_Rmainloop();

       A possible substitute for a pseudo-console is

       R_ReplDLLinit();
       while(R_ReplDLLdo1() > 0) {
           /* add user actions here if desired */
       }

       */
    Rf_endEmbeddedR(0);
    /* final tidying up after R is shutdown */
    return 0;
}
```

If you don’t want to pass R arguments, you can fake an argv array, for example by

```
char *argv[] = {"REmbeddedPostgres", "--silent"};
Rf_initEmbeddedR(sizeof(argv)/sizeof(argv[0]), argv);
```

However, to make a GUI we usually do want to run run\_Rmainloop after setting up various parts of R to talk to our GUI, and arranging for our GUI callbacks to be called during the R mainloop.

One issue to watch is that on some platforms Rf\_initEmbeddedR and Rf\_endEmbeddedR change the settings of the FPU (e.g. to allow errors to be trapped and to set extended precision registers).

---

<sup>2</sup> but these are not part of the automated test procedures and so little tested.

The standard code sets up a session temporary directory in the usual way, *unless* `R_TempDir` is set to a non-NULL value before `Rf_initEmbeddedR` is called. In that case the value is assumed to contain an existing writable directory (no check is done), and it is not cleaned up when R is shut down.

`Rf_initEmbeddedR` sets R to be in interactive mode: you can set `R_Interactive` (defined in ‘`Rinterface.h`’) subsequently to change this.

Note that R expects to be run with the locale category ‘`LC_NUMERIC`’ set to its default value of `C`, and so should not be embedded into an application which changes that.

### 8.1.1 Compiling against the R library

Suitable flags to compile and link against the R (shared or static) library can be found by

```
R CMD config --cppflags
R CMD config --ldflags
```

If R is installed, `pkg-config` is available and sub-architectures have not been used, alternatives for a shared R library are

```
pkg-config --cflags libR
pkg-config --libs libR
```

and for a static R library

```
pkg-config --cflags libR
pkg-config --libs --static libR
```

### 8.1.2 Setting R callbacks

For Unix-alikes there is a public header file ‘`Rinterface.h`’ that makes it possible to change the standard callbacks used by R in a documented way. This defines pointers (if `R_INTERFACE_PTRS` is defined)

```
extern void (*ptr_R_Suicide)(const char *);
extern void (*ptr_R_ShowMessage)(const char *);
extern int (*ptr_R_ReadConsole)(const char *, unsigned char *, int, int);
extern void (*ptr_R_WriteConsole)(const char *, int);
extern void (*ptr_R_WriteConsoleEx)(const char *, int, int);
extern void (*ptr_R_ResetConsole)();
extern void (*ptr_R_FlushConsole)();
extern void (*ptr_R_ClearerrConsole)();
extern void (*ptr_R_Busy)(int);
extern void (*ptr_R_CleanUp)(SA_TYPE, int, int);
extern int (*ptr_R_ShowFiles)(int, const char **, const char **,
                             const char *, Rboolean, const char *);
extern int (*ptr_R_ChooseFile)(int, char *, int);
extern int (*ptr_R_EditFile)(const char *);
extern void (*ptr_R_loadhistory)(SEXP, SEXP, SEXP, SEXP);
extern void (*ptr_R_savehistory)(SEXP, SEXP, SEXP, SEXP);
extern void (*ptr_R_addhistory)(SEXP, SEXP, SEXP, SEXP);
```

which allow standard R callbacks to be redirected to your GUI. What these do is generally documented in the file ‘`src/unix/system.txt`’.

**void R\_ShowMessage** (*char \*message*) [Function]  
 This should display the message, which may have multiple lines: it should be brought to the user's attention immediately.

**void R\_Busy** (*int which*) [Function]  
 This function invokes actions (such as change of cursor) when R embarks on an extended computation (*which=1*) and when such a state terminates (*which=0*).

**int R\_ReadConsole** (*const char \*prompt, unsigned char \*buf, int buflen, int hist*) [Function]

**void R\_WriteConsole** (*const char \*buf, int buflen*) [Function]

**void R\_WriteConsoleEx** (*const char \*buf, int buflen, int otype*) [Function]

**void R\_ResetConsole** () [Function]

**void R\_FlushConsole** () [Function]

**void R\_ClearErrConsole** () [Function]

These functions interact with a console.

**R\_ReadConsole** prints the given prompt at the console and then does a **fgets(3)**-like operation, transferring up to *buflen* characters into the buffer *buf*. The last two bytes should be set to `"\n\0"` to preserve sanity. If *hist* is non-zero, then the line should be added to any command history which is being maintained. The return value is 0 if no input is available and >0 otherwise.

**R\_WriteConsoleEx** writes the given buffer to the console, *otype* specifies the output type (regular output or warning/error). Call to **R\_WriteConsole(buf, buflen)** is equivalent to **R\_WriteConsoleEx(buf, buflen, 0)**. To ensure backward compatibility of the callbacks, **ptr\_R\_WriteConsoleEx** is used only if **ptr\_R\_WriteConsole** is set to NULL. To ensure that **stdout()** and **stderr()** connections point to the console, set the corresponding files to NULL *via*

```
R_Outputfile = NULL;
R_Consolefile = NULL;
```

**R\_ResetConsole** is called when the system is reset after an error. **R\_FlushConsole** is called to flush any pending output to the system console. **R\_ClearerrConsole** clears any errors associated with reading from the console.

**int R\_ShowFiles** (*int nfile, const char \*\*file, const char \*\*headers, const char \*wtitle, Rboolean del, const char \*pager*) [Function]

This function is used to display the contents of files.

**int R\_ChooseFile** (*int new, char \*buf, int len*) [Function]

Choose a file and return its name in *buf* of length *len*. Return value is 0 for success, > 0 otherwise.

**int R\_EditFile** (*const char \*buf*) [Function]

Send a file to an editor window.

**SEXP R\_loadhistory** (*SEXP, SEXP, SEXP, SEXP*); [Function]

**SEXP R\_savehistory** (*SEXP, SEXP, SEXP, SEXP*); [Function]

**SEXP R\_addhistory** (*SEXP, SEXP, SEXP, SEXP*); [Function]

.Internal functions for **loadhistory**, **savehistory** and **timestamp**: these are called after checking the number of arguments.



If the console has no history mechanism these can be as simple as

```
SEXP R_loadhistory (SEXP call, SEXP op, SEXP args, SEXP env)
{
    errorcall(call, "loadhistory is not implemented");
    return R_NilValue;
}
SEXP R_savehistory (SEXP call, SEXP op , SEXP args, SEXP env)
{
    errorcall(call, "savehistory is not implemented");
    return R_NilValue;
}
SEXP R_addhistory (SEXP call, SEXP op , SEXP args, SEXP env)
{
    return R_NilValue;
}
```

The `R_addhistory` function should return silently if no history mechanism is present, as a user may be calling `timestamp` purely to write the time stamp to the console.

`void R_Suicide (const char *message)` [Function]

This should abort R as rapidly as possible, displaying the message. A possible implementation is

```
void R_Suicide (const char *message)
{
    char pp[1024];
    snprintf(pp, 1024, "Fatal error: %s\n", s);
    R_ShowMessage(pp);
    R_CleanUp(SA_SUICIDE, 2, 0);
}
```

`void R_CleanUp (SA_TYPE saveact, int status, int RunLast)` [Function]

This function invokes any actions which occur at system termination. It needs to be quite complex:

```
#include <Rinterface.h>
#include <Rembedded.h>    /* for Rf_KillAllDevices */

void R_CleanUp (SA_TYPE saveact, int status, int RunLast)
{
    if(saveact == SA_DEFAULT) saveact = SaveAction;
    if(saveact == SA_SAVEASK) {
        /* ask what to do and set saveact */
    }
    switch (saveact) {
    case SA_SAVE:
        if(runLast) R_dot_Last();
        if(R_DirtyImage) R_SaveGlobalEnv();
        /* save the console history in R_HistoryFile */
        break;
    }
```

```

        case SA_NOSAVE:
            if(runLast) R_dot_Last();
            break;
        case SA_SUICIDE:
        default:
            break;
    }

    R_RunExitFinalizers();
    /* clean up after the editor e.g. CleanEd() */

    R_CleanTempDir();

    /* close all the graphics devices */
    if(saveact != SA_SUICIDE) Rf_KillAllDevices();
    fpu_setup(FALSE);

    exit(status);
}

```

### 8.1.3 Registering symbols

An application embedding R needs a different way of registering symbols because it is not a dynamic library loaded by R as would be the case with a package. Therefore R reserves a special `DllInfo` entry for the embedding application such that it can register symbols to be used with `.C`, `.Call` etc. This entry can be obtained by calling `getEmbeddingDllInfo`, so a typical use is

```

DllInfo *info = R_getEmbeddingDllInfo();
R_registerRoutines(info, cMethods, callMethods, NULL, NULL);

```

The native routines defined by `cMethod` and `callMethods` should be present in the embedding application. See [Section 5.4 \[Registering native routines\]](#), page 97 for details on registering symbols in general.

### 8.1.4 Meshing event loops

One of the most difficult issues in interfacing R to a front-end is the handling of event loops, at least if a single thread is used. R uses events and timers for

- Running X11 windows such as the graphics device and data editor, and interacting with them (e.g., using `locator()`).
- Supporting Tcl/Tk events for the **tcltk** package (for at least the X11 version of Tk).
- Preparing input.
- Timing operations, for example for profiling R code and `Sys.sleep()`.
- Interrupts, where permitted.

Specifically, the Unix-alike command-line version of R runs separate event loops for

- Preparing input at the console command-line, in file `'src/unix/sys-unix.c'`.
- Waiting for a response from a socket in the internal functions underlying FTP and HTTP transfers in `download.file()` and for direct socket access, in files `'src/`

```
modules/internet/nanoftp.c', 'src/modules/internet/nanohttp.c' and 'src/
modules/internet/Rsock.c'
```

- Mouse and window events when displaying the X11-based dataentry window, in file 'src/modules/X11/dataentry.c'. This is regarded as *modal*, and no other events are serviced whilst it is active.

There is a protocol for adding event handlers to the first two types of event loops, using types and functions declared in the header 'R\_ext/eventloop.h' and described in comments in file 'src/unix/sys-std.c'. It is possible to add (or remove) an input handler for events on a particular file descriptor, or to set a polling interval (*via* `R_wait_usec`) and a function to be called periodically *via* `R_PolledEvents`: the polling mechanism is used by the **cltk** package.

An alternative front-end needs both to make provision for other R events whilst waiting for input, and to ensure that it is not frozen out during events of the second type. This is not handled very well in the existing examples. The GNOME front-end can run a own handler for polled events by setting

```
extern int (*R_timeout_handler)();
extern long R_timeout_val;

if (R_timeout_handler && R_timeout_val)
    gtk_timeout_add(R_timeout_val, R_timeout_handler, NULL);
gtk_main ();
```

whilst it is waiting for console input. This obviously handles events for Gtk windows (such as the graphics device in the **gtkDevice** package), but not X11 events (such as the `X11()` device) or for other event handlers that might have been registered with R. It does not attempt to keep itself alive whilst R is waiting on sockets. The ability to add a polled handler as `R_timeout_handler` is used by the **cltk** package.

### 8.1.5 Threading issues

Embedded R is designed to be run in the main thread, and all the testing is done in that context. There is a potential issue with the stack-checking mechanism where threads are involved. This uses two variables declared in 'Rinterface.h' (if `CSTACK_DEFNS` is defined) as

```
extern uintptr_t R_CStackLimit; /* C stack limit */
extern uintptr_t R_CStackStart; /* Initial stack address */
```

Note that `uintptr_t` is a C99 type for which a substitute is defined in R, so your code needs to define `HAVE_UINTPTR_T` appropriately.

These will be set<sup>3</sup> when `Rf_initialize_R` is called, to values appropriate to the main thread. Stack-checking can be disabled by setting `R_CStackLimit = (uintptr_t)-1`, but it is better to if possible set appropriate values. (What these are and how to determine them are OS-specific, and the stack size limit may differ for secondary threads. If you have a choice of stack size, at least 8Mb is recommended.)

---

<sup>3</sup> at least on platforms where the values are available, that is having `getrlimit` and on Linux or having `sysctl` supporting `KERN_USRSTACK`, including FreeBSD and Mac OS X.

You may also want to consider how signals are handled: R sets signal handlers for several signals, including `SIGINT`, `SIGSEGV`, `SIGPIPE`, `SIGUSR1` and `SIGUSR2`, but these can all be suppressed by setting the variable `R_SignalHandlers` (declared in `'Rinterface.h'`) to 0.

## 8.2 Embedding R under Windows

All Windows interfaces to R call entry points in the DLL `'R.dll'`, directly or indirectly. Simpler applications may find it easier to use the indirect route *via* (D)COM.

### 8.2.1 Using (D)COM

(D)COM is a standard Windows mechanism used for communication between Windows applications. One application (here R) is run as COM server which offers services to clients, here the front-end calling application. The services are described in a 'Type Library' and are (more or less) language-independent, so the calling application can be written in C or C++ or Visual Basic or Perl or Python and so on. The 'D' in (D)COM refers to 'distributed', as the client and server can be running on different machines.

The basic R distribution is not a (D)COM server, but two addons are currently available that interface directly with R and provide a (D)COM server:

- There is a (D)COM server called `StatConnector` written by Thomas Baier available via <http://cran.r-project.org/other-software.html> or <http://sunsite.univie.ac.at/rcom/>, which works with package `rscproxy` to support transfer of data to and from R and remote execution of R commands, as well as embedding of an R graphics window. The `rcom` package on CRAN provides a (D)COM server in a running R session.
- Another (D)COM server, `RDCOMServer`, is available from <http://www.omegahat.org/>. Its philosophy is discussed in <http://www.omegahat.org/RDCOMServer/Docs/Paradigm.html> and is very different from the purpose of this section.

### 8.2.2 Calling R.dll directly

The R DLL is mainly written in C and has `_cdecl` entry points. Calling it directly will be tricky except from C code (or C++ with a little care).

There is a version of the Unix-alike interface calling

```
int Rf_initEmbeddedR(int ac, char **av);
void Rf_endEmbeddedR(int fatal);
```

which is an entry point in `'R.dll'`. Examples of its use (and a suitable `'Makefile.win'`) can be found in the `'tests/Embedding'` directory of the sources. You may need to ensure that `'R_HOME/bin'` is in your `PATH` so the R DLLs are found.

Examples of calling `'R.dll'` directly are provided in the directory `'src/gnuwin32/front-ends'`, including a simple command-line front end `'rtest.c'` whose code is

```
#define Win32
#include <windows.h>
#include <stdio.h>
#include <Rversion.h>
#define LibExtern __declspec(dllimport) extern
#include <Rembedded.h>
#include <R_ext/RStartup.h>
/* for askok and askyesnocancel */
```

```

#include <graphapp.h>

/* for signal-handling code */
#include <psignal.h>

/* simple input, simple output */

/* This version blocks all events: a real one needs to call ProcessEvents
   frequently. See rterm.c and ../system.c for one approach using
   a separate thread for input.
*/
int myReadConsole(const char *prompt, char *buf, int len, int addtohistory)
{
    fputs(prompt, stdout);
    fflush(stdout);
    if(fgets(buf, len, stdin)) return 1; else return 0;
}

void myWriteConsole(const char *buf, int len)
{
    printf("%s", buf);
}

void myCallBack(void)
{
    /* called during i/o, eval, graphics in ProcessEvents */
}

void myBusy(int which)
{
    /* set a busy cursor ... if which = 1, unset if which = 0 */
}

static void my_onintr(int sig) { UserBreak = 1; }

int main (int argc, char **argv)
{
    structRstart rp;
    Rstart Rp = &rp;
    char Rversion[25], *RHome;

    sprintf(Rversion, "%s.%s", R_MAJOR, R_MINOR);
    if(strcmp(getDLLVersion(), Rversion) != 0) {
        fprintf(stderr, "Error: R.DLL version does not match\n");
        exit(1);
    }

    R_setStartTime();
    R_DefParams(Rp);
    if((RHome = get_R_HOME()) == NULL) {
        fprintf(stderr, "R_HOME must be set in the environment or Registry\n");
        exit(1);
    }
    Rp->rhome = RHome;
    Rp->home = getRUser();
    Rp->CharacterMode = LinkDLL;
    Rp->ReadConsole = myReadConsole;
    Rp->WriteConsole = myWriteConsole;
}

```

```

Rp->CallBack = myCallBack;
Rp->ShowMessage = askok;
Rp->YesNoCancel = askyesnocancel;
Rp->Busy = myBusy;

Rp->R_Quiet = TRUE;          /* Default is FALSE */
Rp->R_Interactive = FALSE; /* Default is TRUE */
Rp->RestoreAction = SA_RESTORE;
Rp->SaveAction = SA_NOSAVE;
R_SetParams(Rp);
R_set_command_line_arguments(argc, argv);

FlushConsoleInputBuffer(GetStdHandle(STD_INPUT_HANDLE));

signal(SIGBREAK, my_onintr);
GA_initapp(0, 0);
readconsolecfg();
setup_Rmainloop();
#ifdef SIMPLE_CASE
    run_Rmainloop();
#else
    R_ReplDLLinit();
    while(R_ReplDLLdo1() > 0) {
/* add user actions here if desired */
    }
/* only get here on EOF (not q()) */
#endif
    Rf_endEmbeddedR(0);
    return 0;
}

```

The ideas are

- Check that the front-end and the linked ‘R.dll’ match – other front-ends may allow a looser match.
- Find and set the R home directory and the user’s home directory. The former may be available from the Windows Registry: it will be in `HKEY_LOCAL_MACHINE\Software\R-core\R\InstallPath` from an administrative install and `HKEY_CURRENT_USER\Software\R-core\R\InstallPath` otherwise, if selected during installation (as it is by default).
- Define startup conditions and callbacks *via* the `Rstart` structure. `R_DefParams` sets the defaults, and `R_SetParams` sets updated values.
- Record the command-line arguments used by `R_set_command_line_arguments` for use by the R function `commandArgs()`.
- Set up the signal handler and the basic user interface.
- Run the main R loop, possibly with our actions intermeshed.
- Arrange to clean up.

An underlying theme is the need to keep the GUI ‘alive’, and this has not been done in this example. The R callback `R_ProcessEvents` needs to be called frequently to ensure that Windows events in R windows are handled expeditiously. Conversely, R needs to allow the GUI code (which is running in the same process) to update itself as needed – two ways are provided to allow this:

- `R_ProcessEvents` calls the callback registered by `Rp->callback`. A version of this is used to run package Tcl/Tk for **tcltk** under Windows, for the code is

```
void R_ProcessEvents(void)
{
    while (peekevent()) doevent(); /* Windows events for GraphApp */
    if (UserBreak) { UserBreak = FALSE; onintr(); }
    R_CallBackHook();
    if(R_tcldo) R_tcldo();
}
```

- The mainloop can be split up to allow the calling application to take some action after each line of input has been dealt with: see the alternative code below `#ifdef SIMPLE_CASE`.

It may be that no R GraphApp windows need to be considered, although these include pagers, the `windows()` graphics device, the R data and script editors and various popups such as `choose.file()` and `select.list()`. It would be possible to replace all of these, but it seems easier to allow GraphApp to handle most of them.

It is possible to run R in a GUI in a single thread (as ‘RGui.exe’ shows) but it will normally be easier<sup>4</sup> to use multiple threads.

Note that R’s own front ends use a stack size of 10Mb, whereas MinGW executables default to 2Mb, and Visual C++ ones to 1Mb. The latter stack sizes are too small for a number of R applications, so general-purpose front-ends should use a larger stack size.

### 8.2.3 Finding R\_HOME

Both applications which embed R and those which use a `system` call to invoke R (as `Rscript.exe`, `Rterm.exe` or `R.exe`) need to be able to find the R ‘bin’ directory. The simplest way to do so is to ask the user to set an environment variable `R_HOME` and use that, but naive users may be flummoxed as to how to do so or what value to use.

The R for Windows installers have for a long time allowed the value of `R_HOME` to be recorded in the Windows Registry: this is optional but selected by default. *Where* it is recorded has changed over the years to allow for multiple versions of R to be installed at once, and to allow 32- and 64-bit versions of R to be installed on the same machine.

The basic Registry location is `Software\R-core\R`. For an administrative install this is under `HKEY_LOCAL_MACHINE` and on a 64-bit OS `HKEY_LOCAL_MACHINE\Software\R-core\R` is by default redirected for a 32-bit application, so a 32-bit application will see the information for the last 32-bit install, and a 64-bit application that for the last 64-bit install. For a personal install, the information is under `HKEY_CURRENT_USER\Software\R-core\R` which is seen by both 32-bit and 64-bit applications and so records the last install of either architecture. To circumvent this, there are locations `Software\R-core\R32` and `Software\R-core\R64` which always refer to one architecture.

When R is installed and recording is not disabled then two string values are written at that location for keys `InstallPath` and `Current Version`, and these keys are removed when R is uninstalled. To allow information about other installed versions to be retained, there is

<sup>4</sup> An attempt to use only threads in the late 1990s failed to work correctly under Windows 95, the predominant version of Windows at that time.



also a key named something like 2.11.0 or 2.11.0 patched or 2.12.0 Pre-release with a value for `InstallPath`.

So a comprehensive algorithm to search to `R_HOME` is something like

- Decide which of personal or administrative installs should have precedence. There are arguments both ways: we find that with roaming profiles that `HKEY_CURRENT_USER\Software` often gets reverted to an earlier version. Do the following for one or both of `HKEY_CURRENT_USER` and `HKEY_LOCAL_MACHINE`.
- If the desired architecture is known, look in `Software\R-core\R32` or `Software\R-core\R64`, and if that does not exist or the architecture is immaterial, in `Software\R-core\R`.
- If key `InstallPath` exists then this is `R_HOME` (recorded using backslashes). If it does not, look for version-specific keys like 2.11.0 alpha, pick the latest (which is of itself a complicated algorithm as 2.11.0 patched > 2.11.0 > 2.11.0 alpha > 2.8.1) and use its value for `InstallPath`.

Prior to R 2.12.0 '`R.dll`' and the various front-end executables are in '`R_HOME\bin`', but they are now in '`R_HOME\bin\i386`' or '`R_HOME\bin\x64`'. So you need to arrange to look first in the architecture-specific subdirectory and then in '`R_HOME\bin`'.

## Function and variable index

|                     |          |  |
|---------------------|----------|--|
| <b>*</b>            |          |  |
| *Riconv_open .....  | 143      |  |
| <b>.</b>            |          |  |
| .C .....            | 93       |  |
| .Call .....         | 105, 116 |  |
| .External .....     | 105, 117 |  |
| .Fortran .....      | 93       |  |
| .Last.lib .....     | 40       |  |
| .onAttach .....     | 40       |  |
| .onLoad .....       | 40       |  |
| .onUnload .....     | 40       |  |
| .Random.seed .....  | 132      |  |
| <b>\</b>            |          |  |
| \acronym .....      | 65       |  |
| \alias .....        | 57       |  |
| \arguments .....    | 59       |  |
| \author .....       | 60       |  |
| \bold .....         | 64       |  |
| \cite .....         | 65       |  |
| \code .....         | 64       |  |
| \command .....      | 65       |  |
| \concept .....      | 69       |  |
| \cr .....           | 63       |  |
| \deqn .....         | 67       |  |
| \describe .....     | 66       |  |
| \description .....  | 57       |  |
| \details .....      | 59       |  |
| \dfn .....          | 65       |  |
| \dontrun .....      | 60       |  |
| \dontshow .....     | 60       |  |
| \dots .....         | 68       |  |
| \dQuote .....       | 64       |  |
| \email .....        | 65       |  |
| \emph .....         | 64       |  |
| \enc .....          | 69       |  |
| \enumerate .....    | 65       |  |
| \env .....          | 65       |  |
| \eqn .....          | 67       |  |
| \examples .....     | 60       |  |
| \figure .....       | 68       |  |
| \file .....         | 65       |  |
| \format .....       | 62       |  |
| \href .....         | 65       |  |
| \if .....           | 70       |  |
| \ifelse .....       | 70       |  |
| \itemize .....      | 65       |  |
| \kbd .....          | 64       |  |
| \keyword .....      | 60       |  |
| \ldots .....        | 68       |  |
| \link .....         | 66       |  |
| \method .....       | 58       |  |
| \name .....         | 56       |  |
| \newcommand .....   | 71       |  |
| \note .....         | 59       |  |
| \option .....       | 65       |  |
| \out .....          | 70       |  |
| \pkg .....          | 64       |  |
| \preformatted ..... | 64       |  |
| \R .....            | 68       |  |
| \RdOpts .....       | 70       |  |
| \references .....   | 59       |  |
| \renewcommand ..... | 71       |  |
| \S3method .....     | 59       |  |
| \samp .....         | 64       |  |
| \section .....      | 63       |  |
| \seealso .....      | 60       |  |
| \Sexpr .....        | 70       |  |
| \source .....       | 62       |  |
| \sQuote .....       | 64       |  |
| \strong .....       | 64       |  |
| \tabular .....      | 66       |  |
| \title .....        | 57       |  |
| \url .....          | 65       |  |
| \usage .....        | 57       |  |
| \value .....        | 59       |  |
| \var .....          | 65       |  |
| \verb .....         | 64       |  |
| <b>B</b>            |          |  |
| bessel_i .....      | 137      |  |
| bessel_j .....      | 137      |  |
| bessel_k .....      | 137      |  |
| bessel_y .....      | 137      |  |
| beta .....          | 137      |  |
| BLAS_LIBS .....     | 18       |  |
| browser .....       | 81       |  |
| <b>C</b>            |          |  |
| Calloc .....        | 131      |  |
| CAR .....           | 118      |  |
| CDR .....           | 118      |  |
| cgmin .....         | 139      |  |
| choose .....        | 137      |  |
| CITATION .....      | 12, 52   |  |
| cPsort .....        | 141      |  |
| <b>D</b>            |          |  |
| debug .....         | 85       |  |
| debugger .....      | 84       |  |
| defineVar .....     | 113      |  |
| digamma .....       | 136      |  |
| dump.frames .....   | 83       |  |

duplicate..... 115  
 dyn.load..... 95  
 dyn.unload..... 95

## E

exp\_rand..... 132  
 expm1..... 137  
 export..... 39  
 exportClasses..... 44  
 exportClassPattern..... 44  
 exportMethods..... 44  
 exportPattern..... 39, 44

## F

FALSE..... 139  
 findInterval..... 142  
 findVar..... 113  
 FLIBS..... 18  
 fmax2..... 138  
 fmin2..... 138  
 fprec..... 138  
 Free..... 131  
 fround..... 138  
 fsign..... 138  
 ftrunc..... 138

## G

gammafn..... 136  
 gctorture..... 86  
 getAttrib..... 111  
 getCharCE..... 129  
 GetRNGstate..... 132

## H

hsv2rgb..... 142

## I

imax2..... 138  
 imin2..... 138  
 import..... 39  
 importClassesFrom..... 44  
 importFrom..... 39  
 importMethodsFrom..... 44  
 install..... 111  
 iPsort..... 141  
 ISNA..... 119, 133  
 ISNAN..... 119, 133

## L

LAPACK\_LIBS..... 18  
 lbeta..... 137

lbfgsb..... 140  
 lchoose..... 137  
 lgamma1p..... 137  
 lgammafn..... 136  
 library.dynam..... 10, 95  
 log1p..... 137  
 log1pmax..... 137  
 logspace\_add..... 138  
 logspace\_sub..... 138

## M

M\_E..... 138  
 M\_PI..... 138  
 mkChar..... 113  
 mkCharCE..... 129  
 mkCharLen..... 113  
 mkCharLenCE..... 129

## N

NA\_REAL..... 133  
 nmmin..... 139  
 norm\_rand..... 132

## O

OBJECTS..... 19, 100

## P

pentagamma..... 136  
 PKG\_CFLAGS..... 99  
 PKG\_CPPFLAGS..... 99  
 PKG\_CXXFLAGS..... 99  
 PKG\_FCFLAGS..... 99  
 PKG\_FFLAGS..... 99  
 PKG\_LIBS..... 100  
 PKG\_OBJCFLAGS..... 99  
 prompt..... 56  
 PROTECT..... 106  
 PROTECT\_WITH\_INDEX..... 108  
 psigamma..... 136  
 PutRNGstate..... 132

## Q

qsort3..... 142  
 qsort4..... 142

## R

R CMD build..... 29  
 R CMD check..... 26  
 R CMD config..... 16  
 R CMD Rd2pdf..... 73  
 R CMD Rdconv..... 73  
 R CMD SHLIB..... 99

R\_CMD\_Stangle..... 73  
 R\_CMD\_Sweave..... 73  
 R\_addhistory..... 153  
 R\_alloc..... 131  
 R\_Busy..... 153  
 R\_ChooseFile..... 153  
 R\_CleanUp..... 154  
 R\_ClearErrConsole..... 153  
 R\_csort..... 141  
 R\_EditFile..... 153  
 R\_ExpandFileName..... 143  
 R\_FINITE..... 133  
 R\_FlushConsole..... 153  
 R\_GetCCallable..... 99  
 R\_INLINE..... 144  
 R\_IsNaN..... 133  
 R\_isort..... 141  
 R\_LIBRARY\_DIR..... 17  
 R\_loadhistory..... 153  
 R\_max\_col..... 142  
 R\_NegInf..... 133  
 R\_PACKAGE\_DIR..... 17  
 R\_PACKAGE\_NAME..... 17  
 R\_ParseVector..... 126  
 R\_PosInf..... 133  
 R\_pow..... 137  
 R\_pow\_di..... 137  
 R\_qsort..... 142  
 R\_qsort\_I..... 142  
 R\_qsort\_int..... 142  
 R\_qsort\_int\_I..... 142  
 R\_ReadConsole..... 153  
 R\_RegisterCCallable..... 99  
 R\_registerRoutines..... 97  
 R\_ResetConsole..... 153  
 R\_rsort..... 141  
 R\_savehistory..... 153  
 R\_ShowFiles..... 153  
 R\_ShowMessage..... 153  
 R\_Suicide..... 154  
 R\_tmpnam..... 143  
 R\_tmpnam2..... 143  
 R\_Version..... 144  
 R\_WriteConsole..... 153  
 R\_WriteConsoleEx..... 153  
 rcont2..... 133  
 Rdqagi..... 140  
 Rdqags..... 140  
 Realloc..... 131  
 recover..... 85  
 reEnc..... 129  
 REprintf..... 133  
 REPROTECT..... 108  
 REvprintf..... 133

revsort..... 141  
 rgb2hsv..... 142  
 Riconv..... 143  
 Riconv\_close..... 143  
 Rprintf..... 133  
 Rprof..... 74, 77  
 Rprofmem..... 77  
 rPsort..... 141  
 rsort\_with\_index..... 141  
 Rvprintf..... 133

## S

S\_alloc..... 131  
 S\_realloc..... 131  
 S3method..... 40  
 SAFE\_FFLAGS..... 19  
 samin..... 140  
 seed\_in..... 132  
 seed\_out..... 132  
 setAttrib..... 111  
 setVar..... 113  
 sign..... 138  
 summaryRprof..... 77  
 system..... 93  
 system.time..... 93  
 system2..... 93

## T

tetragamma..... 136  
 trace..... 86  
 traceback..... 82  
 tracemem..... 77  
 translateChar..... 129  
 translateCharUTF8..... 129  
 trigamma..... 136  
 TRUE..... 139

## U

undebg..... 86  
 unif\_rand..... 132  
 UNPROTECT..... 106  
 UNPROTECT\_PTR..... 108  
 untracemem..... 77  
 useDynLib..... 40

## V

vmaxget..... 131  
 vmaxset..... 131  
 vmmmin..... 139

# Concept index

|   |        |  |          |
|---|--------|--|----------|
| •   |        | Error handling from C.....                 | 132      |
| .install_extras file.....                 | 33     | Error handling from FORTRAN.....           | 132      |
| .Rbuildignore file.....                   | 30     | Evaluating R expressions from C.....       | 119      |
| .Rinstignore file.....                    | 12     | external pointer.....                      | 126      |
| \   |        | <b>F</b>                                   |          |
| \linkS4class.....                         | 66     | Figures in documentation.....              | 68       |
| <b>A</b>                                  |        | finalizer.....                             | 127      |
| Allocating storage.....                   | 108    | Finding variables.....                     | 113      |
| Attributes.....                           | 109    | <b>G</b>                                   |          |
| <b>B</b>                                  |        | Gamma function.....                        | 136      |
| Bessel functions.....                     | 137    | Garbage collection.....                    | 106      |
| Beta function.....                        | 137    | Generic functions.....                     | 148      |
| Building binary packages.....             | 31     | <b>H</b>                                   |          |
| Building source packages.....             | 29     | handling character data.....               | 113      |
| <b>C</b>                                  |        | Handling lists.....                        | 112      |
| C++ code, interfacing.....                | 101    | Handling R objects in C.....               | 105      |
| Calling C from FORTRAN and vice versa.... | 134    | <b>I</b>                                   |          |
| Checking packages.....                    | 26     | IEEE special values.....                   | 119, 133 |
| citation.....                             | 12, 52 | INDEX file.....                            | 9        |
| Classes.....                              | 112    | Indices.....                               | 69       |
| cleanup file.....                         | 3      | Inspecting R objects when debugging.....   | 90       |
| conditionals.....                         | 70     | integration.....                           | 140      |
| configure file.....                       | 3      | Interfaces to compiled code.....           | 93, 116  |
| Copying objects.....                      | 115    | Interfacing C++ code.....                  | 101      |
| CRAN.....                                 | 34     | Interrupts.....                            | 143      |
| CRAN submission.....                      | 34     | <b>L</b>                                   |          |
| Creating packages.....                    | 2      | LICENCE file.....                          | 3        |
| Creating shared objects.....              | 99     | LICENSE file.....                          | 3        |
| Cross-references in documentation.....    | 66     | Lists and tables in documentation.....     | 65       |
| cumulative hazard.....                    | 135    | <b>M</b>                                   |          |
| <b>D</b>                                  |        | Marking text in documentation.....         | 63       |
| Debugging.....                            | 89     | Mathematics in documentation.....          | 67       |
| DESCRIPTION file.....                     | 4      | Memory allocation from C.....              | 130      |
| Details of R types.....                   | 108    | Memory use.....                            | 76       |
| Distribution functions from C.....        | 135    | Method functions.....                      | 148      |
| Documentation, writing.....               | 55     | Missing values.....                        | 119, 133 |
| Dynamic loading.....                      | 95     | <b>N</b>                                   |          |
| dynamic pages.....                        | 70     | namespaces.....                            | 38       |
| <b>E</b>                                  |        | Numerical analysis subroutines from C..... | 135      |
| Editing Rd files.....                     | 73     |  |          |
| encoding.....                             | 72     |  |          |

Numerical derivatives..... 122

## O

OpenMP ..... 21, 144  
Operating system access..... 93  
optimization..... 139

## P

Package builder..... 29  
Package bundles..... 13  
Package structure..... 3  
Package subdirectories..... 10  
Packages ..... 2  
Parsing R code from C ..... 125  
Platform-specific documentation ..... 69  
Printing from C ..... 133  
Printing from FORTRAN ..... 134  
Processing Rd format..... 73  
Profiling ..... 74, 76, 78

## R

Random numbers in C..... 132, 136  
Random numbers in FORTRAN..... 134  
Registering native routines..... 97

## S

Setting variables..... 113  
Sort functions from C ..... 141  
Submitting to CRAN..... 34  
Sweave..... 32

## T

tarballs..... 29  
Tidying R code..... 74

## U

user-defined macros..... 71

## V

Version information from C..... 144  
vignettes ..... 32  
Visibility ..... 145

## W

weak reference..... 127

## Z

Zero-finding ..... 121